


Enhancements to HPCToolkit for Analysis of CPU and GPU-accelerated Applications

Laksono Adhianto and John Mellor-Crummey
Rice University

Scalable Tools Workshop 2025


HPC Toolkit Workflow



HPCToolkit New Features and Enhancements

- **CPU profiling**
 - Support for Intel Top-down analysis
- **GPU profiling**
 - Support for AMD, Intel and NVIDIA GPUs
- Access to remote databases
 - Have been used in several hackathon
 - Ongoing work to improve the performance
- Presentation
 - Updated Hpcviewer GUI with new look

Support for Top-Down Model Analysis (TMA)



PERF_METRICS MSR

- Intel PERF_METRICS MSR: a special counter to provide percentages of slots for four TMA level 1 and four TMA level 2 metrics
- Four additional TMA level 2 metrics (Core Bound, Fetch Bandwidth, Machine Clears and Light Operations) can be derived from the metrics
- **Pros:** No need to configure many counters
- **Cons:** Limited granularity, not a precise event, update rate ~1-10ms (observed, not official rate), **not be suitable for fine-grain measurement**




Figure 21-40. PERF_METRICS MSR Definition for 12th Generation Intel® Core™ Processor P-core

$$\text{Core_Bound} = \text{Backend_Bound} - \text{Memory_Bound}$$

$$\text{Fetch_Bandwidth} = \text{Frontend_Bound} - \text{Fetch_Latency}$$

$$\text{Machine_Clears} = \text{Bad_Speculation} - \text{Branch-Mispredict}$$

$$\text{Light_Operations} = \text{Retiring} - \text{Heavy_Operation}$$

Tools Supporting Top-Down Analysis

Tool	Mode	Built-in top-down levels supported
Caliper	Counting & Sampling	Level 1-3
HPCToolkit	Sampling	Level 1-4 (mostly)
Likwid	Counting	Level 1 (at least)
Linux Perf stat	Counting	All levels, default level 1-2
PAPI	Counting	Level 1-2
Score-P	Counting	Level 1-2
VTune	Sampling	Level 1-4 and some level 5 & 6

Issues Top-Down Analysis in Sampling Mode

- Issues with libpfm4
 - Bug in translating from top-down events to perf_event configuration
 - Fixed in the main branch, but not in the release
 - New TOPDOWN_M pseudo event to use PERF_METRICS MSR
- Issue with Intel perfmon JSON file
 - Incorrect specification of some top-down events: Serializing_Operation, AMX_Busy, and Nop_Instructions
 - Fixed in the main branch
- Issue with Linux v5
 - Unable to group top-down events in sampling mode with perf_events
 - Grouping top-down events works fine in counting mode
 - Fixed in Linux v6

Issue with Linux v5 (case with perf tool)

```
$ perf stat -e '{slots,topdown-bad-spec}' /bin/ls  
... (success)
```

```
$ perf record -e '{slots,topdown-bad-spec}' /bin/ls  
Error: The sys_perf_event_open() syscall returned with 22 (Invalid argument) for event (topdown-bad-spec).
```


```
$ perf record -e '{slots,topdown-bad-spec}:S' /bin/ls  
Error: The sys_perf_event_open() syscall returned with 22 (Invalid argument) for event (topdown-bad-spec).
```

```
$ perf record -e '{slots,cycles,topdown-bad-spec}:S' /bin/ls  
... (success)
```


Top-Down Analysis Implementation

- Measurements
 - TMA level 1 and 2: use `PERF_METRICS` MSR
 - TMA level 3 and 4: use groups of hardware counters as specified in Intel JSON file
 - Avoid using all TMA metrics in JSON file, select only critical ones
 - Test usability of hardware counters before using them
 - CPU_CLK_UNHALTED.THREAD **VS** CPU_CLK_UNHALTED.THREAD_P
 - Avoid using deprecated counters
 - OFFCORE_REQUESTS_OUTSTANDING.ALL_DATA_RD **VS** OFFCORE_REQUESTS_OUTSTANDING.DATA_RD
 - Profiled with frequency-based sampling + multiplexing (time sharing)


Implementation Challenges



- So many events, so few registers
 - Even worse if SMT is enabled → even less registers
 - Multiplexing (time sharing) interpolates the samples → inaccuracy → uncertainty
- Accuracy issues
 - Some counters may overcount: FP scalars & FP vectors due to FMA instruction
 - Some counters may overlap: Branch resteers with MS switches, L3 hit latency with Contested accesses, ...
- Precision issues
 - PERF_METRICS MSR is not designed for fine-grain measurement, can we mix with precise events?
- Presentation: how to present effectively to users?
 - Calling-context tree + Top-down Metrics + source codes



Topdown metrics of [I] CalcElemFBHourglassForce



Top-down view Bottom-up view Flat view

Scope

	Total-Slots (I)	Total-Slots (E)
Experiment Aggregate Metrics	3.00e+11	100.0%
<thread root>	1.59e+11	53.1%
» gomp_thread_start [libgomp.so.1.0.0]	1.59e+11	53.1%
loop at team.c: 139	1.59e+11	53.1%
130 » gomp_team_barrier_wait_end [libgomp.so.1.0.0]	2.61e+10	8.7%
129 » CalcFBHourglassForceForElems(Domain&, double*, double*, double*, double*, double*, double*, double*, double, int, i...	2.51e+10	8.4%
loop at lulesh.cc: 824	2.51e+10	8.4%
923 » [I] CalcElemFBHourglassForce	1.11e+10	3.7%
lulesh.cc: 726	1.95e+09	0.6%
lulesh.cc: 724	1.29e+09	0.4%
lulesh.cc: 702	1.26e+09	0.4%
lulesh.cc: 720	9.09e+08	0.3%

GPU Support in HPCToolkit

Unreleased

	ROCM	Level0	CUDA	OpenCL
Profiling GPU Operations	✓	✓	✓	✓
Tracing of GPU operations	✓	✓	✓	✓
PC sampling	✓	✓	✓	
HW counters for kernel launches	✓			
Page migration; scratch space mgmt	✓			
Binary instrumentation		✓		


Key Additions to AMD Rocprofiler-sdk API This Year

- Initialization: `rocprofiler_configure`/`rocprofiler_force_configure`
 - Integration with HPCToolkit was surprisingly subtle
 - Either HPCToolkit initialization or `rocprofiler_configure` callback may occur first
 - HPCToolkit initialization of rocprofiler-sdk triggers `rocprofiler_force_configure`
 - `rocprofiler_configure` triggers HPCToolkit initialization
 - Requires a “rendevous” so that all HPCToolkit initialization occurs within `prepare_measurement_subsystem` and all rocprofiler-sdk initialization occurs in the scope of a callback
 - PC sampling configuration
 - Informs a tool which GPUs will be used by a process: enables selective configuration

New Design for HPCToolkit's GPU Monitoring Substrate

- New design supports multiple monitoring threads
- GPU monitoring
 - CUDA: single monitoring thread
 - Level 0 and OpenCL: unspecified threads
 - AMD rocprofiler-sdk supports multiple monitoring threads
 - HW counter reporting
 - PC sample reporting
 - Activity API for reporting GPU operations

Runtime Processing of GPU Measurement Data



Vendor GPU Monitoring Concerns

- Intel Level0 reports distinct GPU binaries per MPI rank!
- Intel's PTI View API is only half complete
 - Provides a completion callback that delivers a sequence of “activity records” for GPU ops
 - Lacking several key capabilities
 - initialization
 - intercept launch of GPU operations (for correlating them with their invocation context)
- AMD GPU OpenMP support is awkward
 - Had to use special ROCm API for monitoring rather than OMPT interface
- NVIDIA Activity API record for a kernel provides only string name
 - Lacks the precise attribution to function objects present in PC samples
 - Requires awkward recording of strings rather than addresses!

Tool Challenges and Approaches

- Tool code in the application namespace: unwanted interactions
 - Dangerous to load a tool's C++ library in application namespace
 - Application and tool may be linked with different C++ libraries
 - Application symbols interfere with tools
 - Some applications use libunwind that conflict with our tools, others define mmap
 - Some versions of bash (e.g. RHEL 8) define getenv, which interferes with tool startup
 - HPCToolkit, rocprofiler-sdk, PAPI, and libpfm all expect libc getenv
- Namespaces in Linux: dlmopen (glibc 2.3.4), LD_AUDIT (glibc 2.4)
- HPCToolkit
 - Uses multiple namespaces to avoid conflicts
 - Uses LD_AUDIT to monitor dynamic library loading and symbol binding

Software Infrastructure Woes

- Thread-local variables aren't not async signal safe
 - <https://sourceware.org/glibc/wiki/TLSandSignals>
- Pthread keys don't support multiple namespaces!
 - https://sourceware.org/bugzilla/show_bug.cgi?id=24776
- Worse: dynamic linker prior to glibc 2.34 uses pthread key
 - Support for multiple namespaces is broken on Aurora (glibc 2.31)

Pthread Keys don't Support Multiple Namespaces!

```
typedef int (*pthread_key_create_t)(pthread_key_t *, void (*)(void*));
int main() {
    pthread_key_t k1, k2;

    int rc = pthread_key_create(&k1, NULL);                                // create key in the default namespace

    void *h = dlmopen(LM_ID_NEWLM, "libpthread.so.0", RTLD_LAZY); // open libpthread in a new namespace
    assert(h != NULL);
    // find pthread_key_create in the new namespace
    pthread_key_create_t fn = (pthread_key_create_t)dlsym(h, "pthread_key_create");
    assert(fn != NULL);

    rc = fn(&k2, NULL);                                              // create a key in the new namespace
    assert(rc == 0);

    assert(k2 != k1);                                                 // the key is same in both namespaces

    return 0;
}
```


assert fails on ALL Linux systems

Summary

- Support for Top-down analysis
 - HPCToolkit provides metrics in different views and different scopes for procedures, loops or lines
 - Current support: Intel Sapphire Rapids or newer
 - Not supported: E-Core in hybrid architectures
 - Can be extended to other platforms: AMD and ARM and GPUs
 - Ongoing work: handle uncertainty
- Support for new GPU monitoring substrate, rocprofiler-sdk will soon be released
- Need continued engagement with Intel to get a usable PTI-View interface
- Re-engage Linux glibc team about improving support for namespaces, dynamic linking, and LD_AUDIT

Backup Slides

Access to Remote Databases



Multi-producer, Single-consumer, Wait-free Queue

```
1  struct Element
2      Element *next, Record r
3
4  struct Queue
5      Element *head, Element *tail
6
7  procedure init(Queue q)
8      q.head = NULL; q.tail = NULL
9
10 procedure enqueue(Queue q, Element e)
11     atomic_store(&e.next, NULL)
12     previous_tail = atomic_exchange(&q.tail, e)
13     if previous_tail is NULL then
14         atomic_store(&q.head, e)
15     else
16         atomic_store(&previous_tail.next, e)
17
```

```
18 function dequeue(Queue q)
19     first = atomic_load(&q.head)
20     if first is NULL then return NULL
21     successor = atomic_load(&first->next)
22     if successor is not NULL then
23         atomic_store(&q->head, successor)
24         return first
25     else if atomic_compare_exchange(
26             &q.tail, &first, NULL)
27         expected_head = first
28         atomic_compare_exchange(&q.head,
29             &expected_head, NULL)
30         return first
31     else
32         return NULL
```

Note: sacrifice linearizability for wait-freedom