



Understanding GPU Kernels with Instrumentation

Yumeng Liu
Advisor: John Mellor-Crummey
Rice University

Scalable Tools Workshop 2025

Performance Analysis on a GPU

- Current tools can provide
 - Traces of GPU functions (kernel launch, memory copies, etc.)
 - High-level metrics of GPU kernels (GPU utilization, bandwidth utilization, etc.)
 - Detailed metrics per instruction (instruction counts, etc.)
 - Traces of metrics
- A missing feature
 - Traces of GPU **device** function calls **inside** kernels over kernel execution time
 - Why? Behavior of complex kernels may differ greatly across threads

Complex GPU-Accelerated Application: Quicksilver

- A proxy application for a dynamic Monte Carlo particle transport code *Mercury*, **consisting of a single GPU kernel with thousands of lines of code that invokes other GPU device functions**
- High-level structure of the GPU kernel
 - Each GPU thread works on one particle
 - Each particle keeps evolving until a termination condition is met
 - **Work for particles differ between threads**

```
cycle_tracking() {  
    for all particles {  
        do {  
            compute distance to census  
            compute distance to facet  
            compute distance to reaction  
            do segment with shortest distance  
            increment tallies  
        } until census, absorbed, escaped  
    }  
}
```

Potential load imbalance, optimization opportunity

Our Goals

- Trace GPU device function calls within kernels to understand thread behavior
- Explore potential optimizations

Outline

- Our trace tool
- Potential optimization opportunities of Quicksilver found with the trace tool
- Exploration of an optimization opportunity
- Work stealing algorithm among GPU threads
- Summary

Outline

- Our trace tool
- Potential optimization opportunities of Quicksilver found with the trace tool
- Exploration of an optimization opportunity
- Work stealing algorithm among GPU threads
- Summary

Instrumentation of GPU Binaries

- Use **NVBit** for dynamic binary instrumentation on NVIDIA GPUs
- Add instrumentation to trace function invocations on each thread
 - Instrumentation points
 - Before a call (CALL)
 - Upon function entry
 - Before function exit (RET, EXIT)

Instrumentation of GPU Binaries

- At each instrumentation point, emit a record containing
 - Warp ID + Mask => Thread ID - **who**
 - Timestamp - **when**
 - Address (**Special: EXIT will be 0**) - **where**

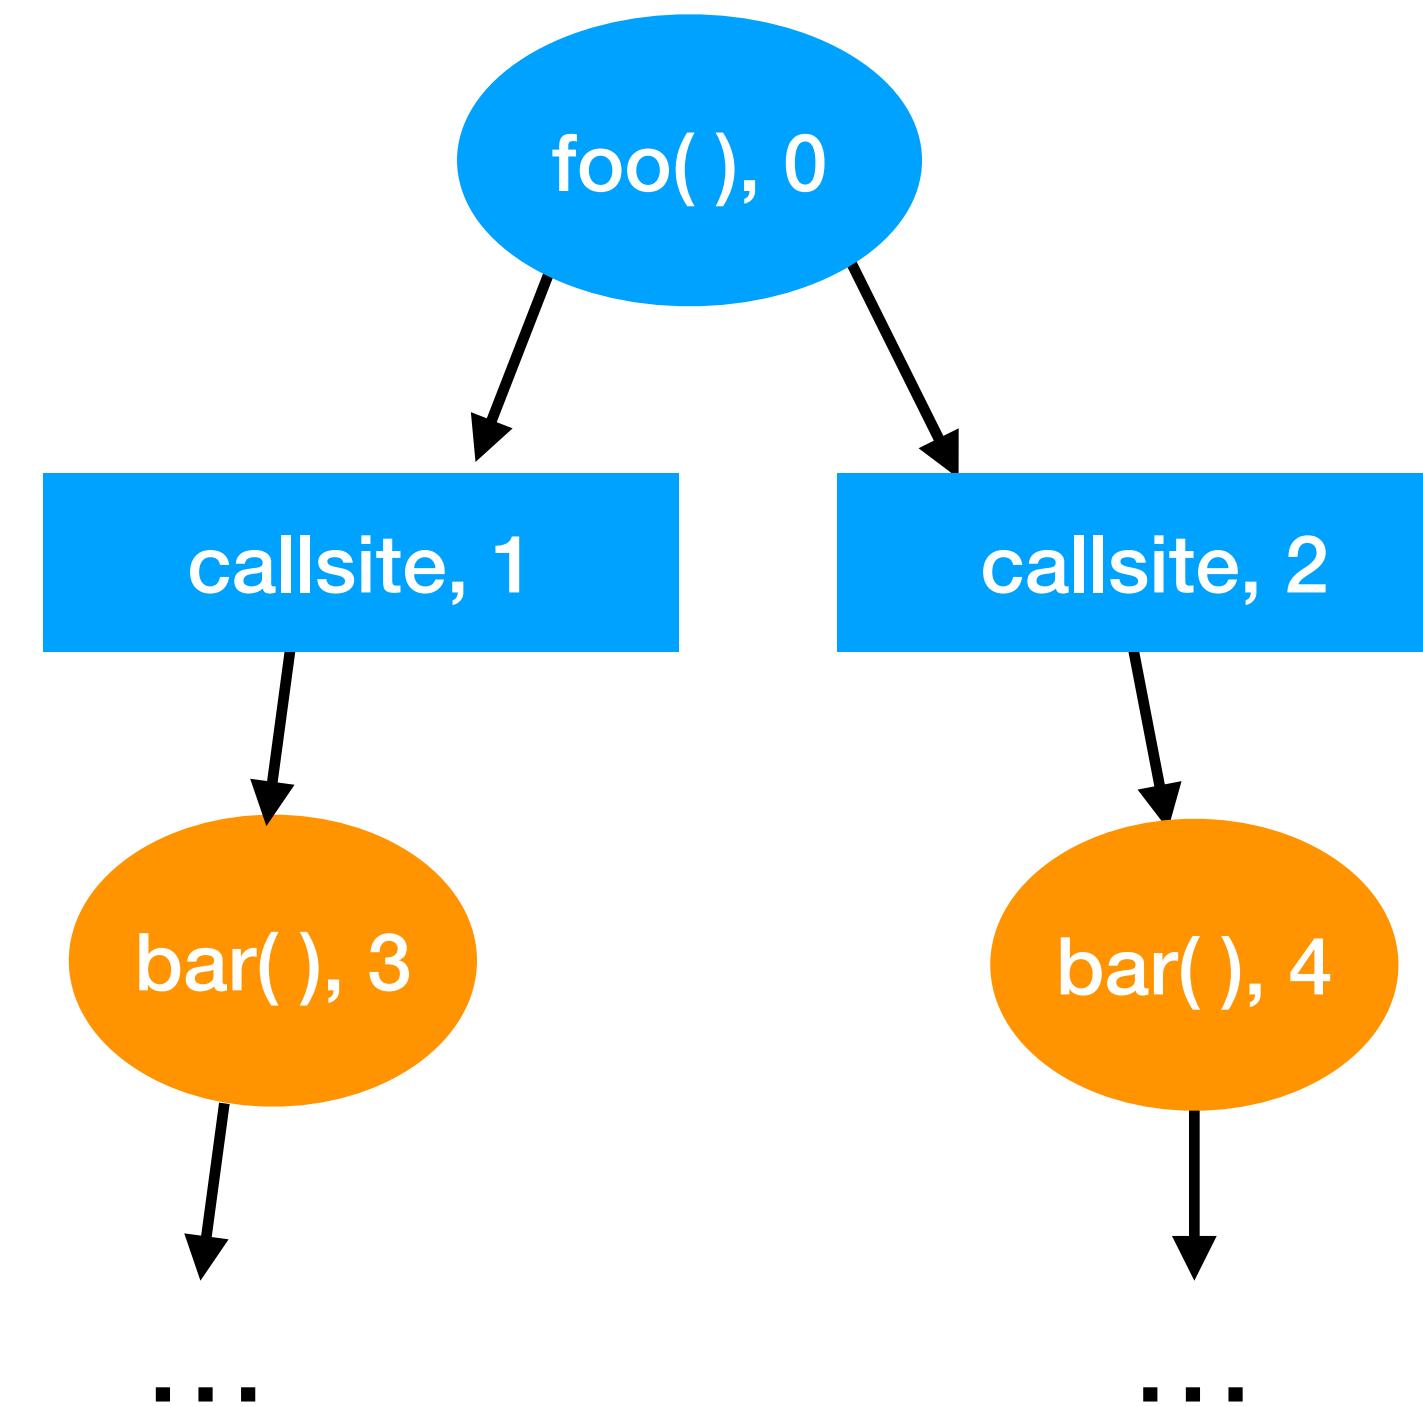
Moving Trace Records to CPU

- Use ChannelDev/ChannelHost objects from NVBit to move records back to CPU
 - One buffer on GPU, one buffer on CPU, flush the buffer when it's full

Processing Trace Stream on CPU

- Input: a stream of records containing **who** (thread id), **where** (function), **when** (timestamp)
- Output: **Calling Context Tree (CCT)** and trace, using HPCToolkit measurement format

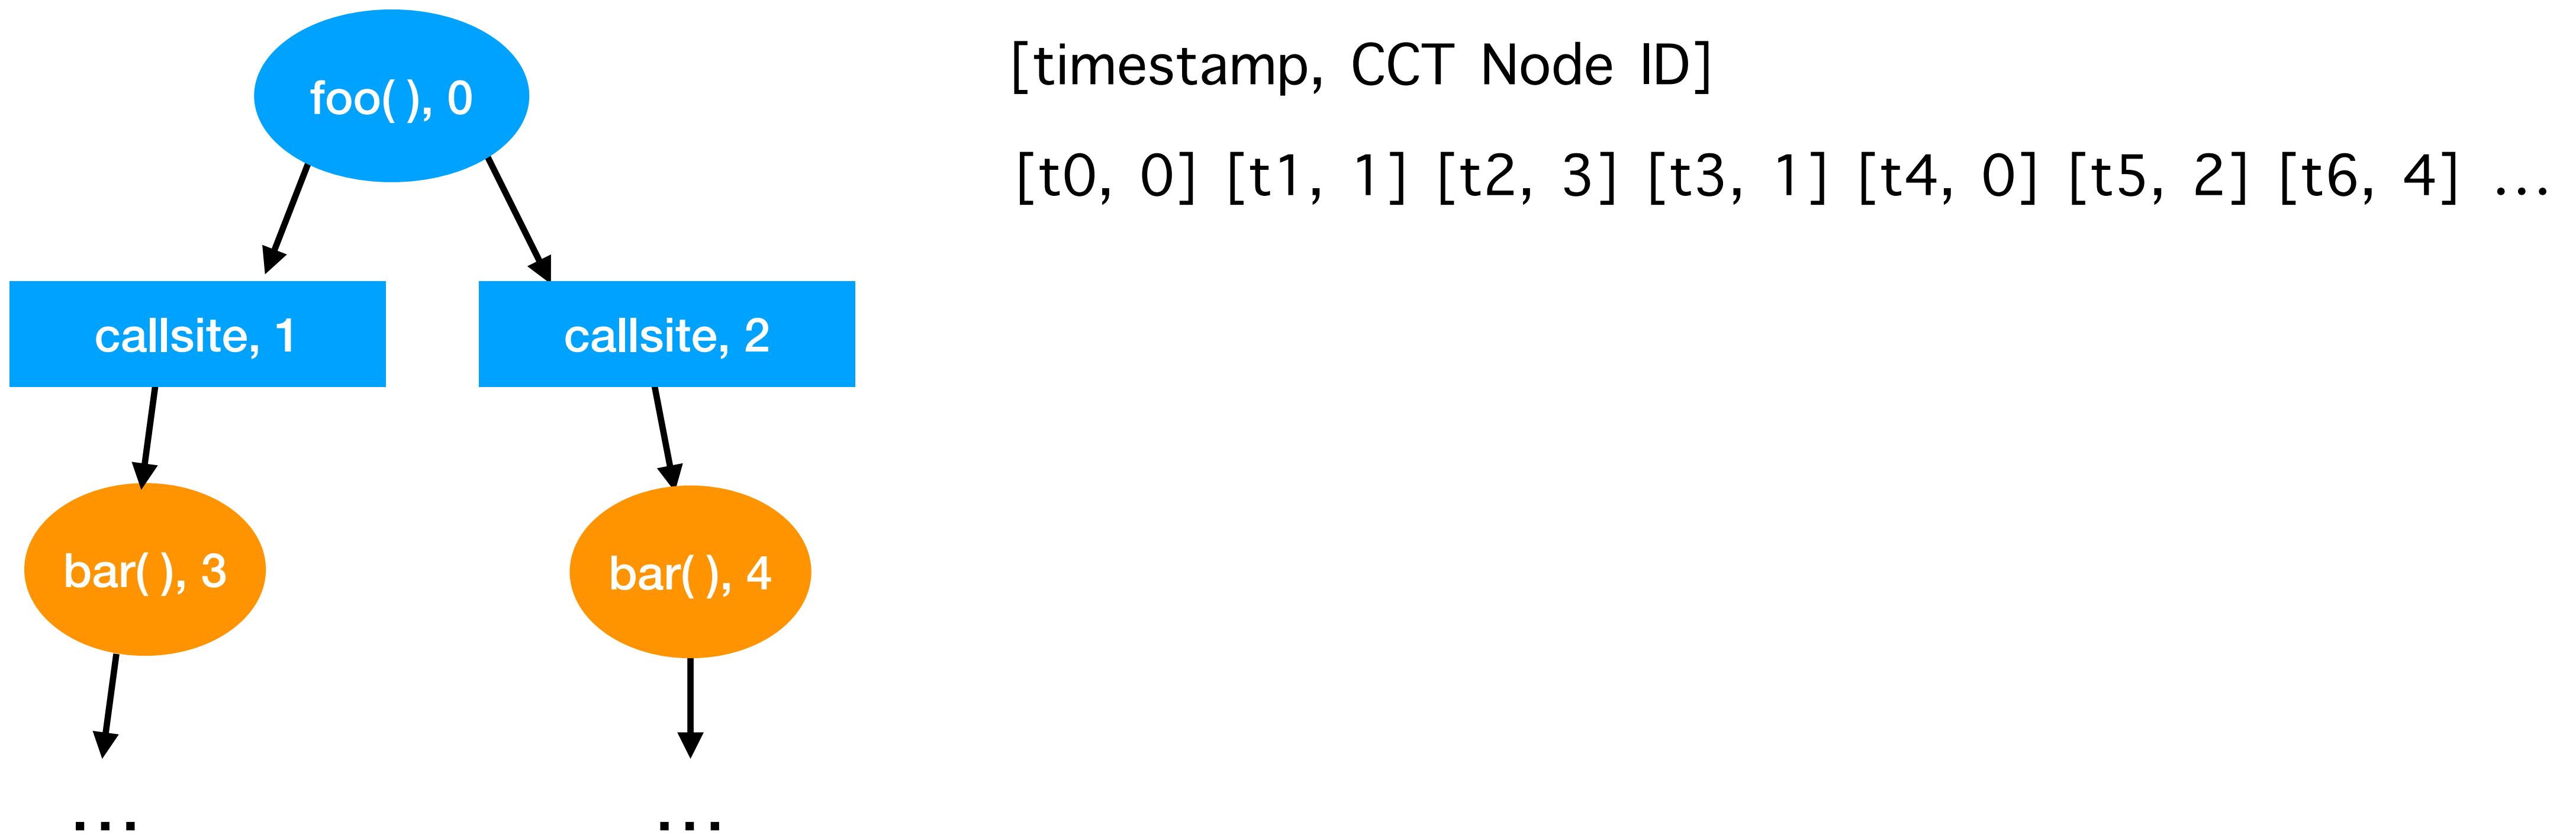
```
void foo() {  
    // ...  
    bar();  
    // ...  
    bar();  
    // ...  
}
```



Callsite information
differentiates multiple
calls by a function to the
same callee

Processing Trace Stream on CPU

- Input: a stream of records containing **who** (thread id), **where** (function), **when** (timestamp)
- Output: Calling Context Tree (CCT) and **trace**, using HPCToolkit measurement format



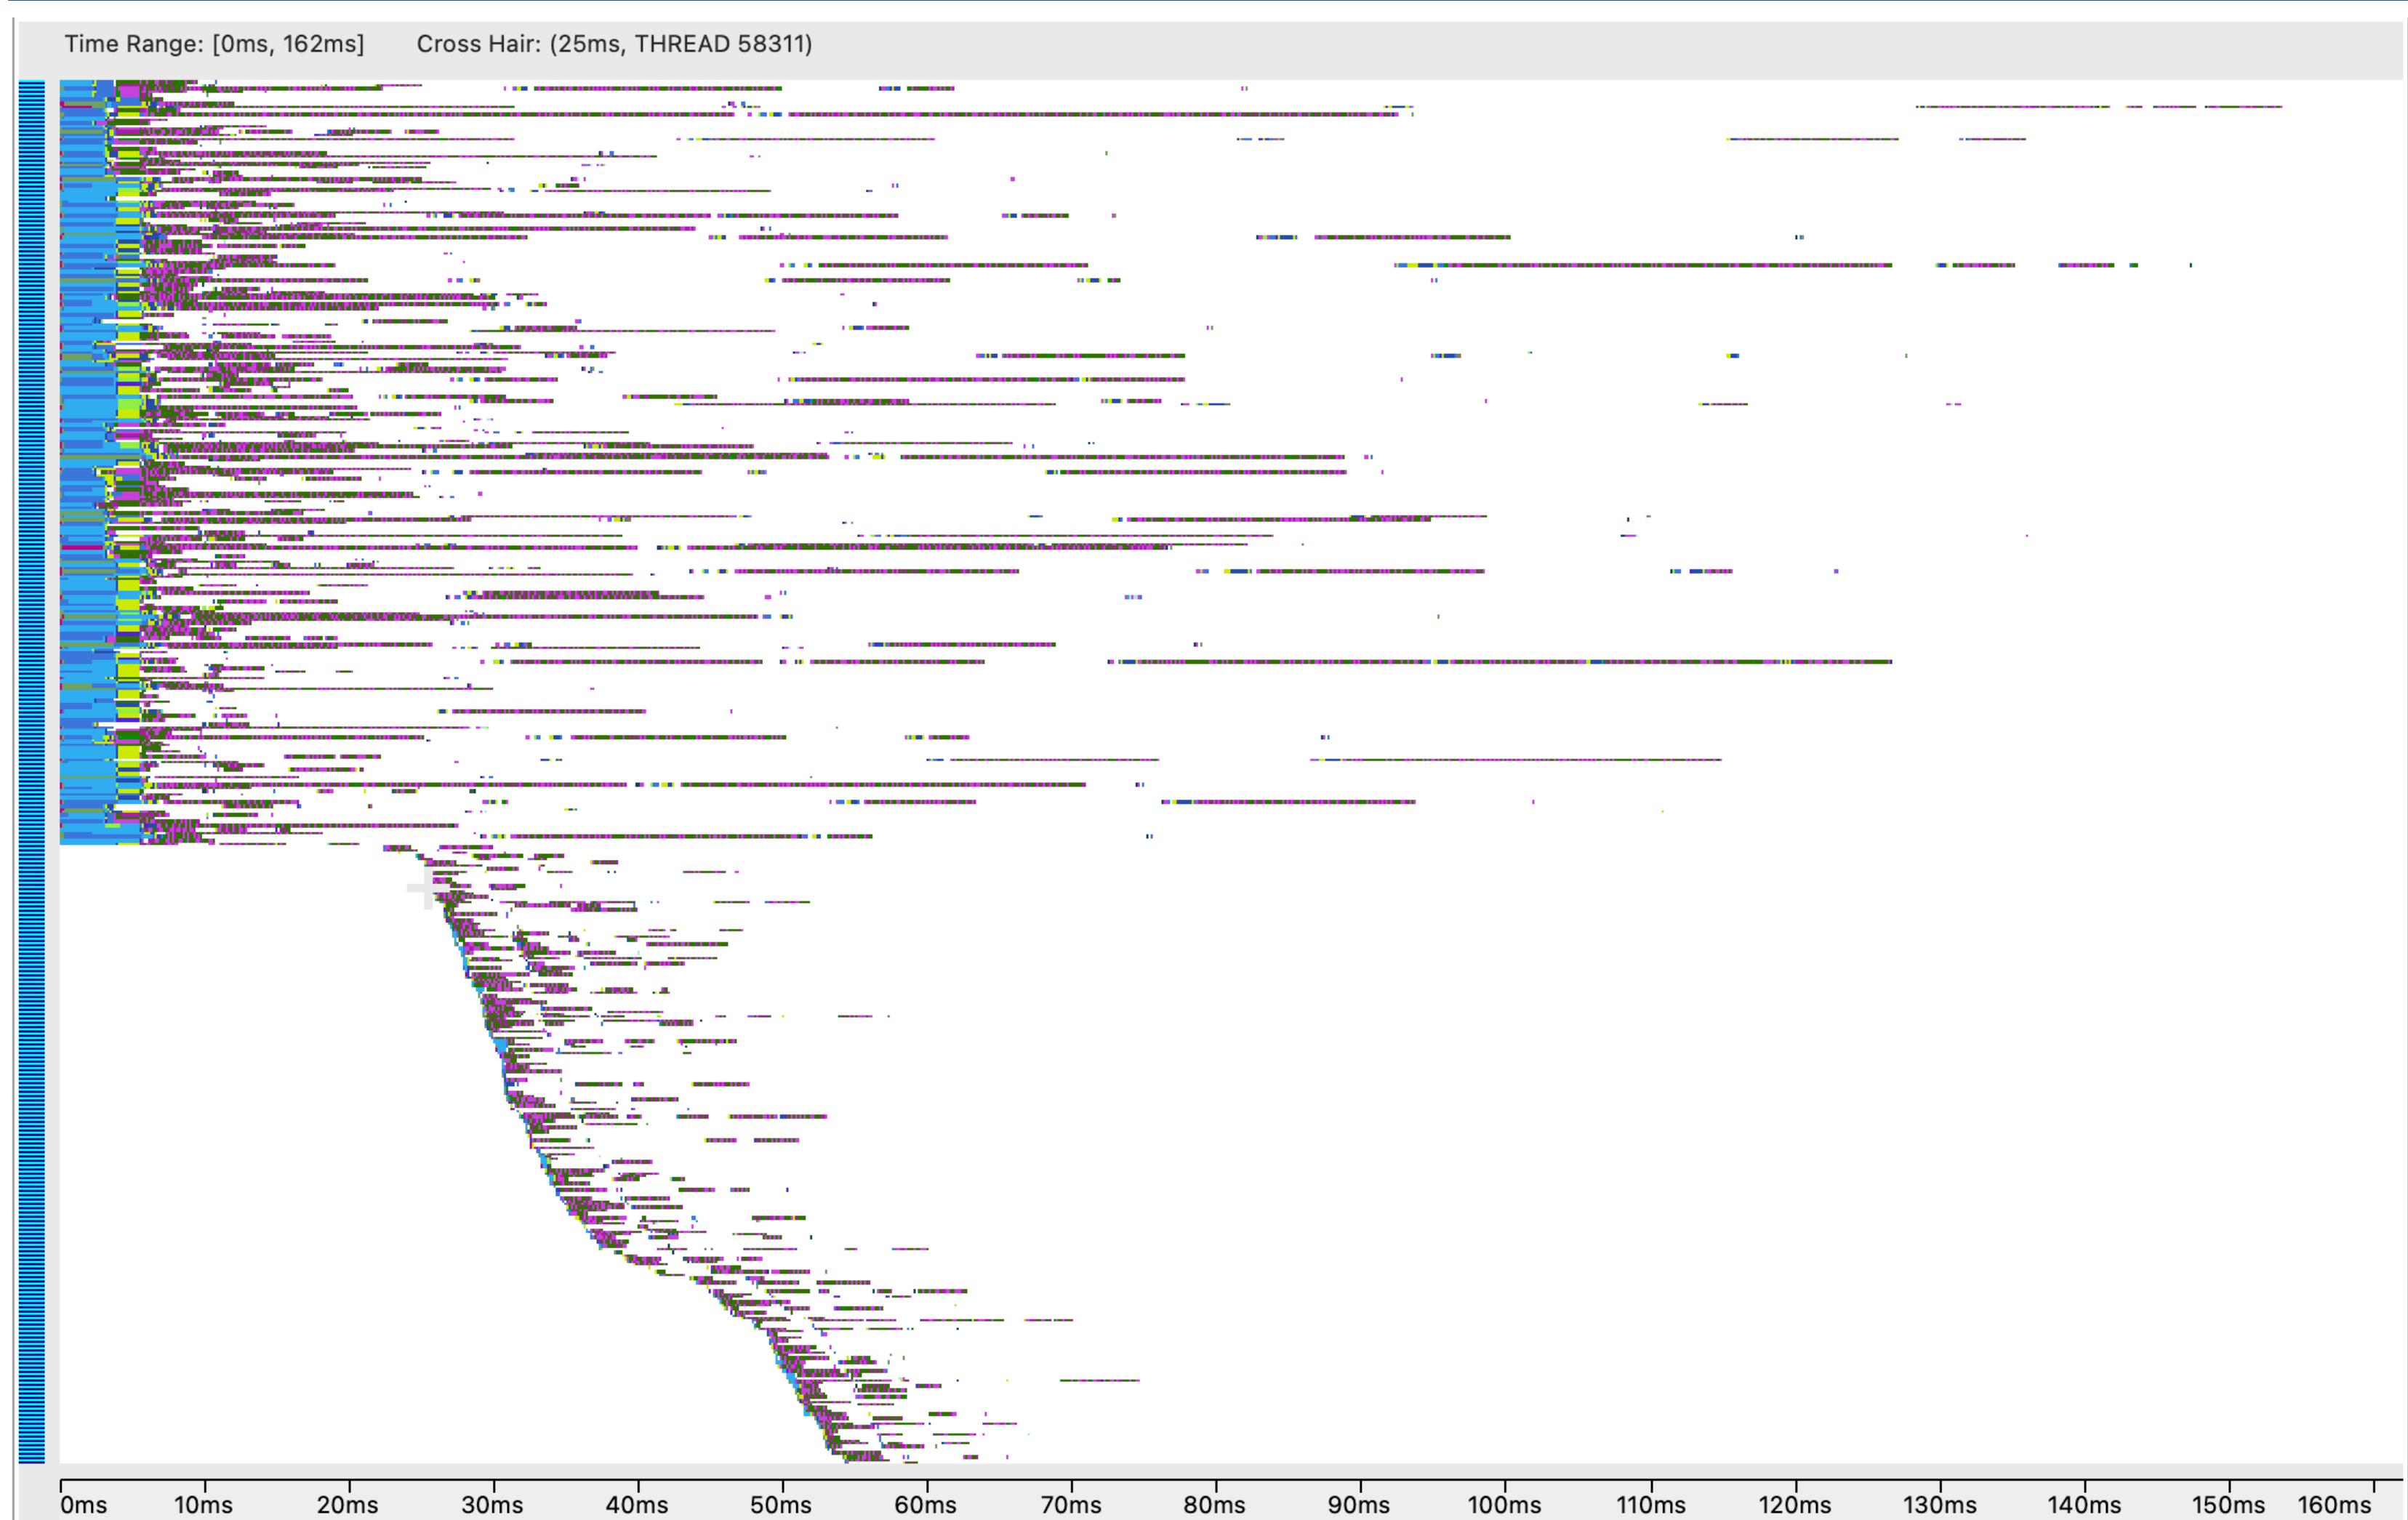
Post-mortem Analysis

- Analyze the GPU binary with *hpcstruct*
- Interpret the trace data with *hpcprof*
- View the traces in *hpcviewer*

Outline

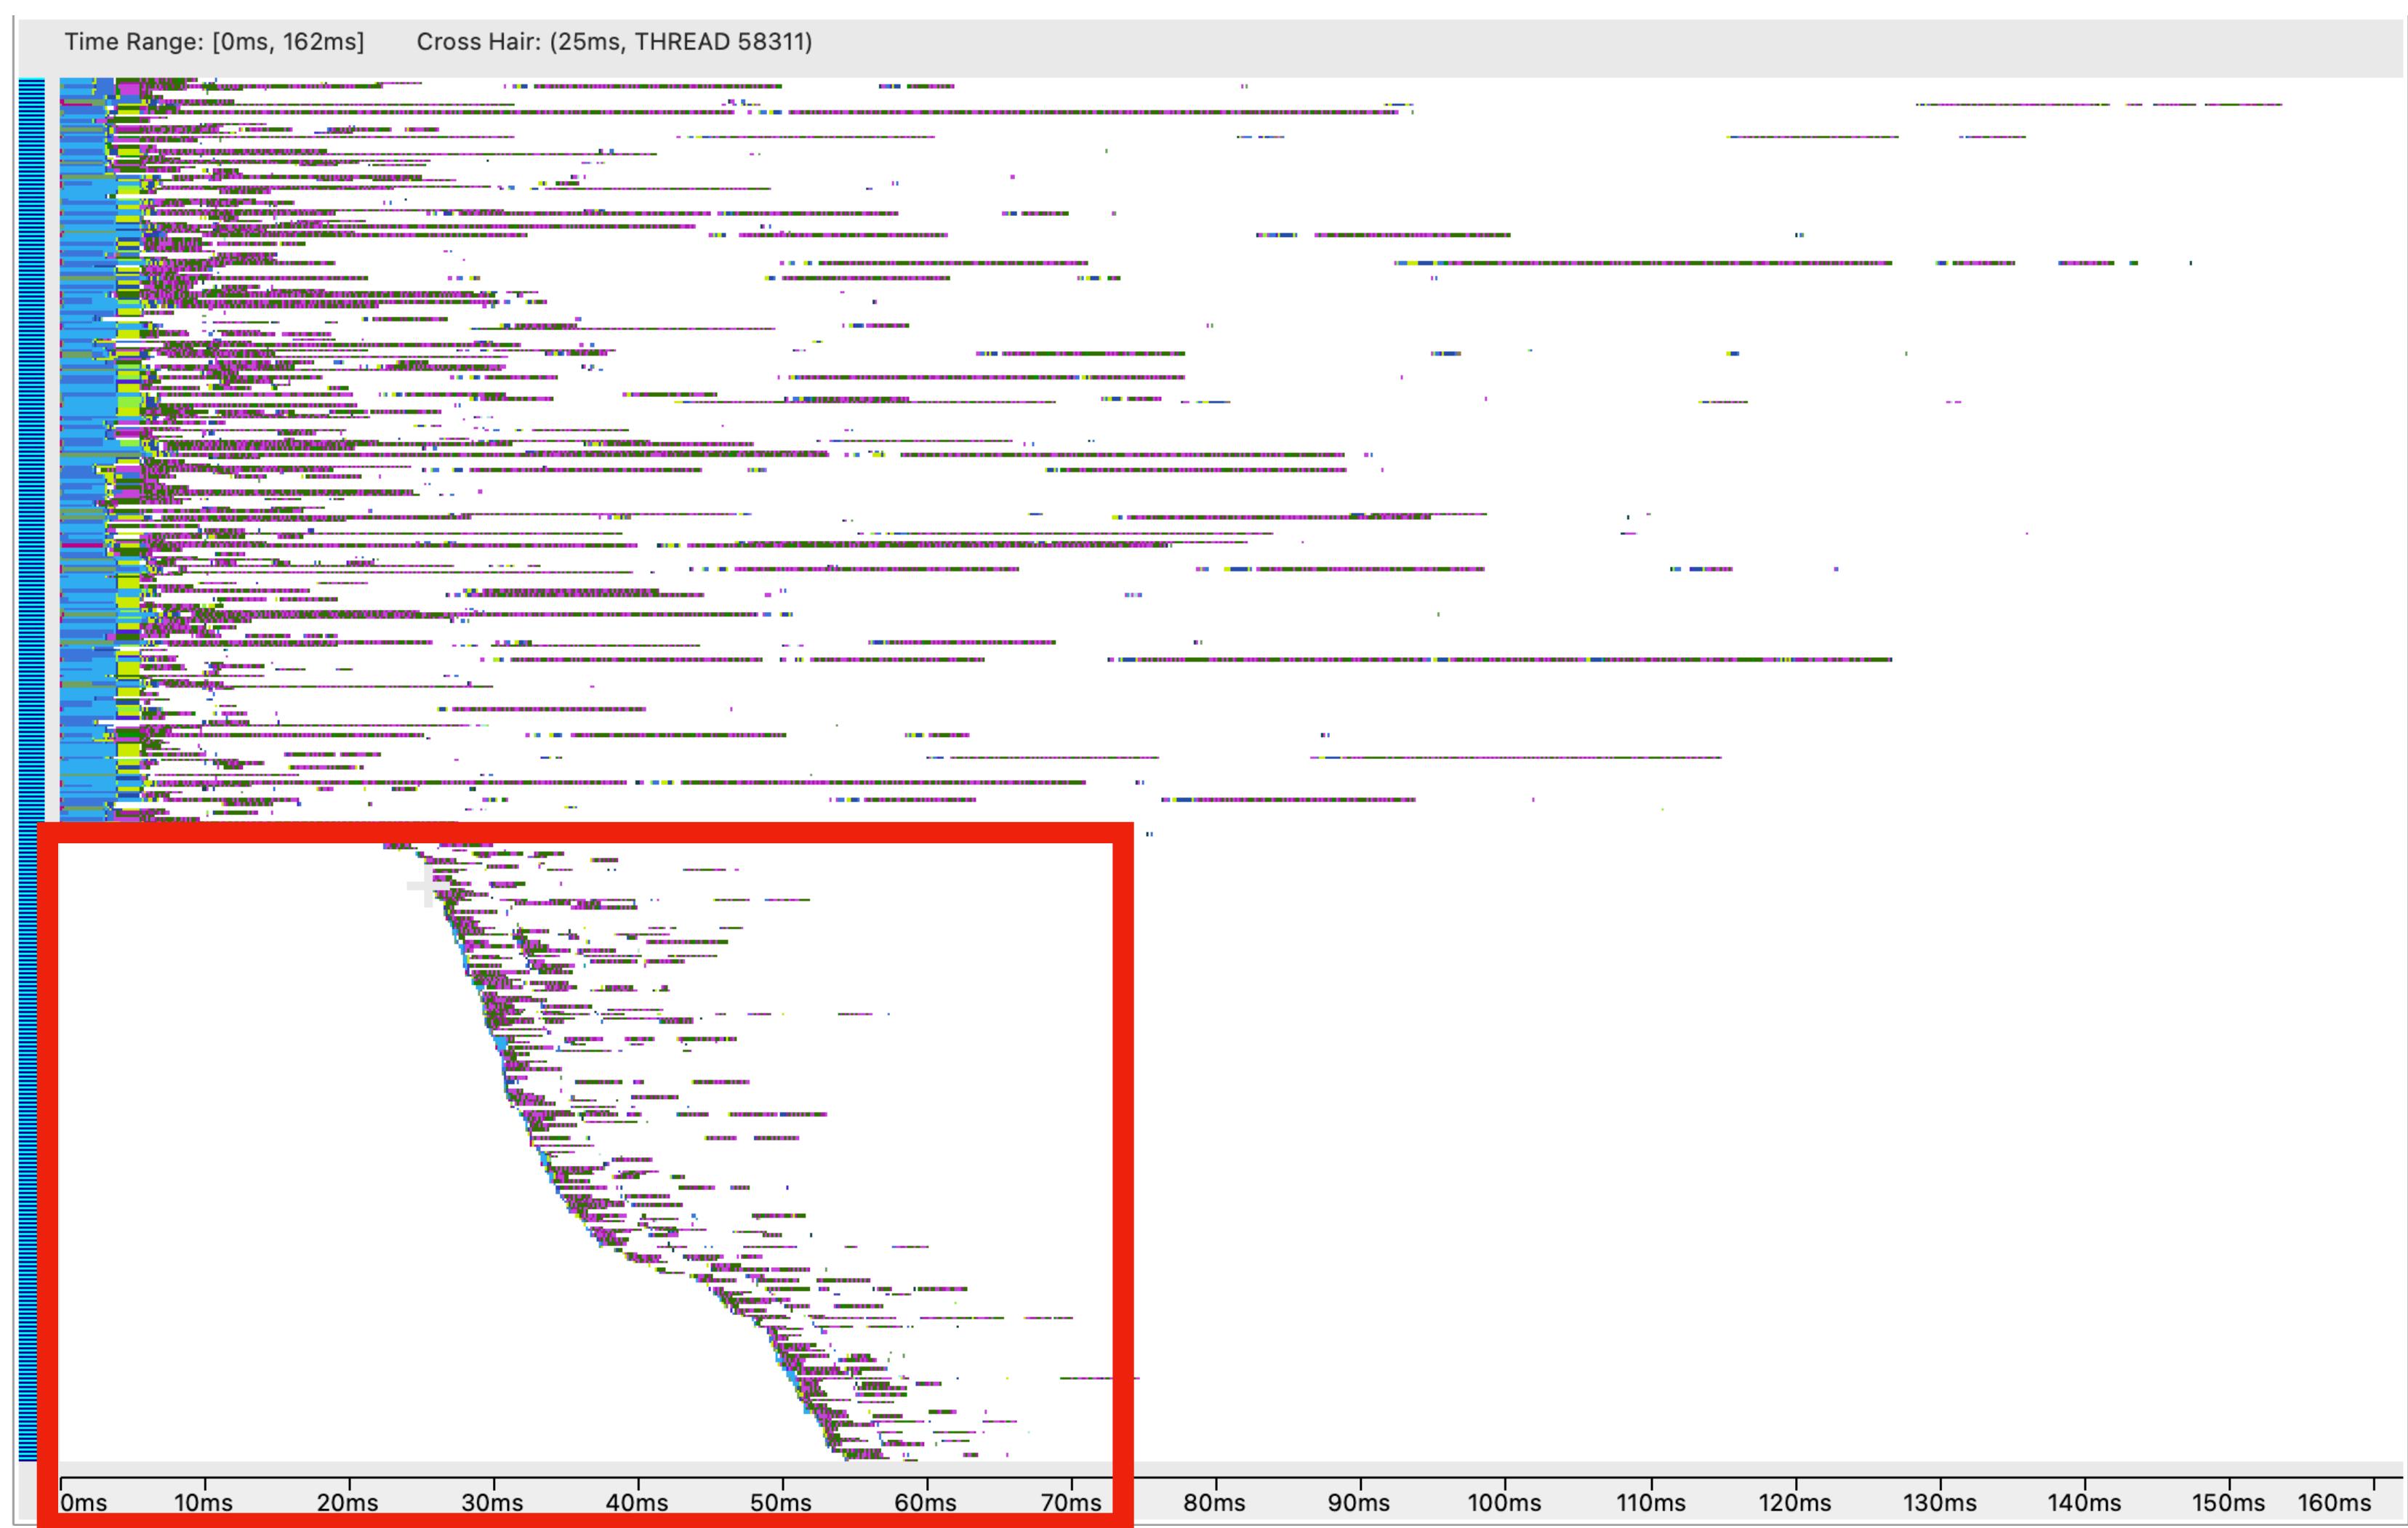
- Our trace tool
- **Potential optimization opportunities of Quicksilver found with the trace tool**
- Exploration of an optimization opportunity
- Work stealing algorithm among GPU threads
- Summary

Quicksilver in Our Tool



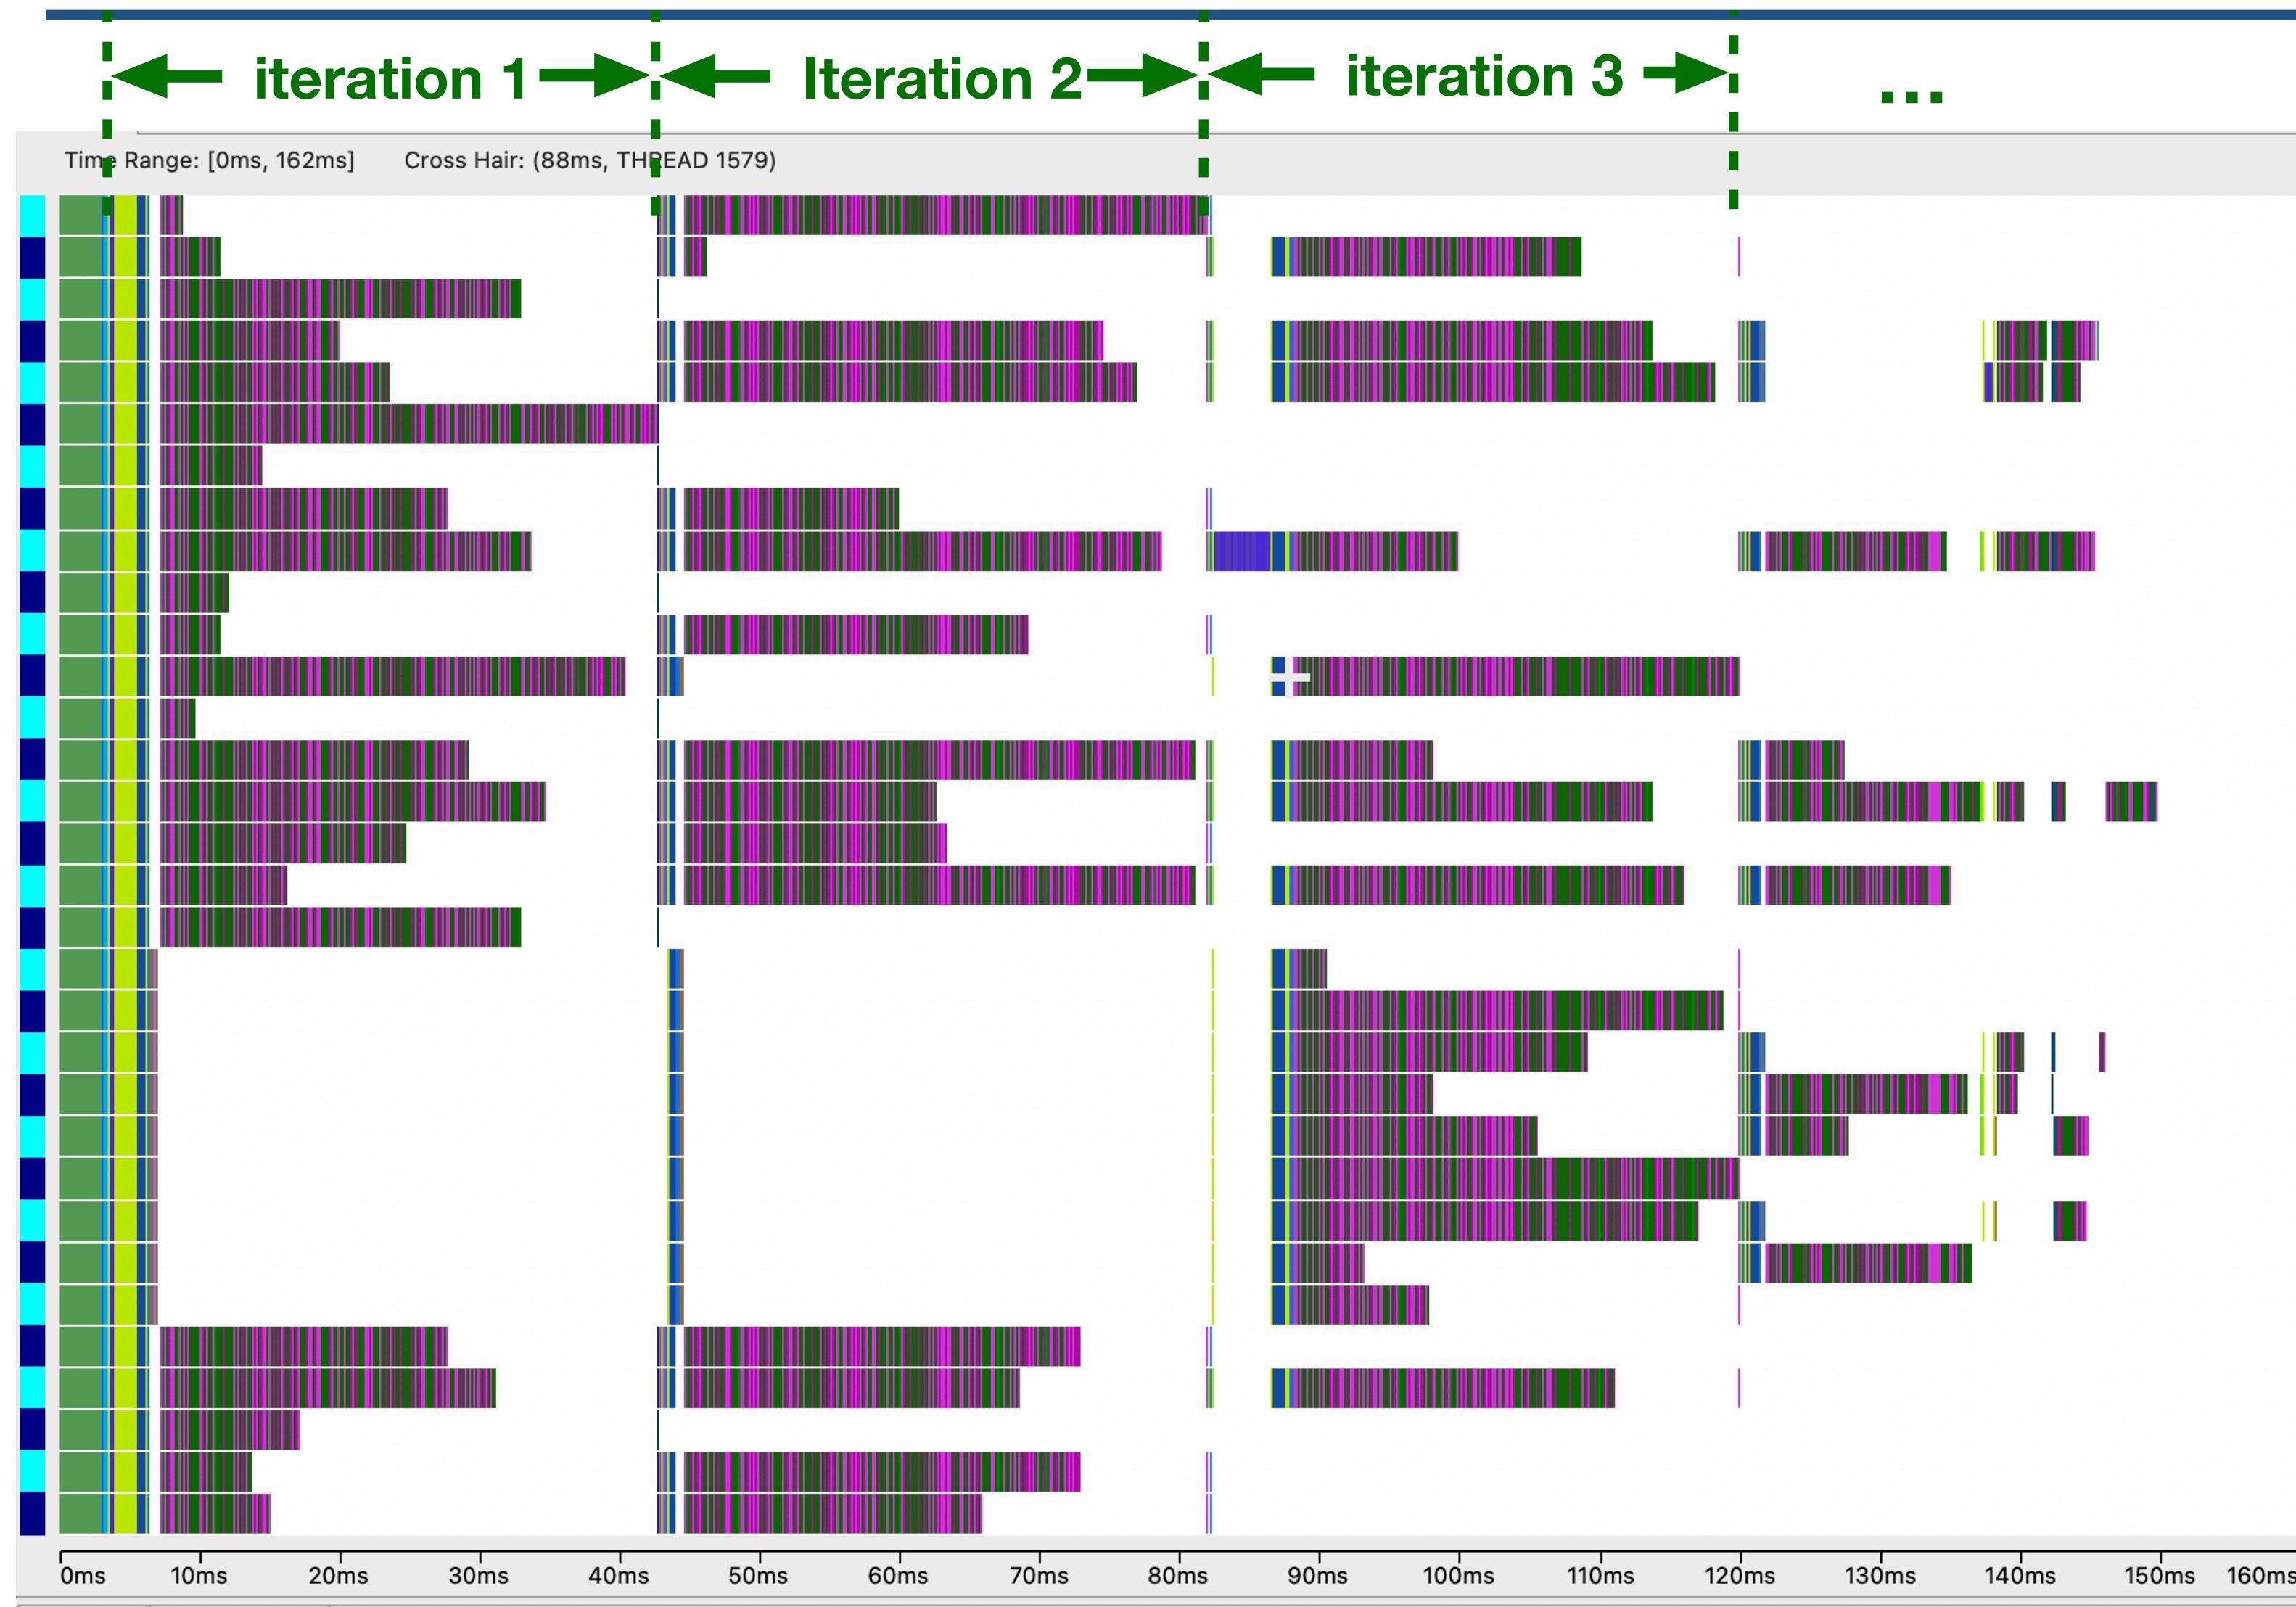
- 100,000 particles
- 100,096 threads
- First impression
 - A lot of idleness
 - Obvious load imbalance

Quicksilver in Our Tool



- 100,000 particles
- 100,096 threads
- First impression
 - A lot of idleness
 - Obvious load imbalance
- Roughly half threads start late

Quicksilver in Our Tool



```
CycleTrackingFunction(){
```

```
do{
```

```
    calculate an outcome  
    switch(outcome){  
        case 1  
        case 2  
        ...  
    }  
}
```

```
}while(keepTracking);
```

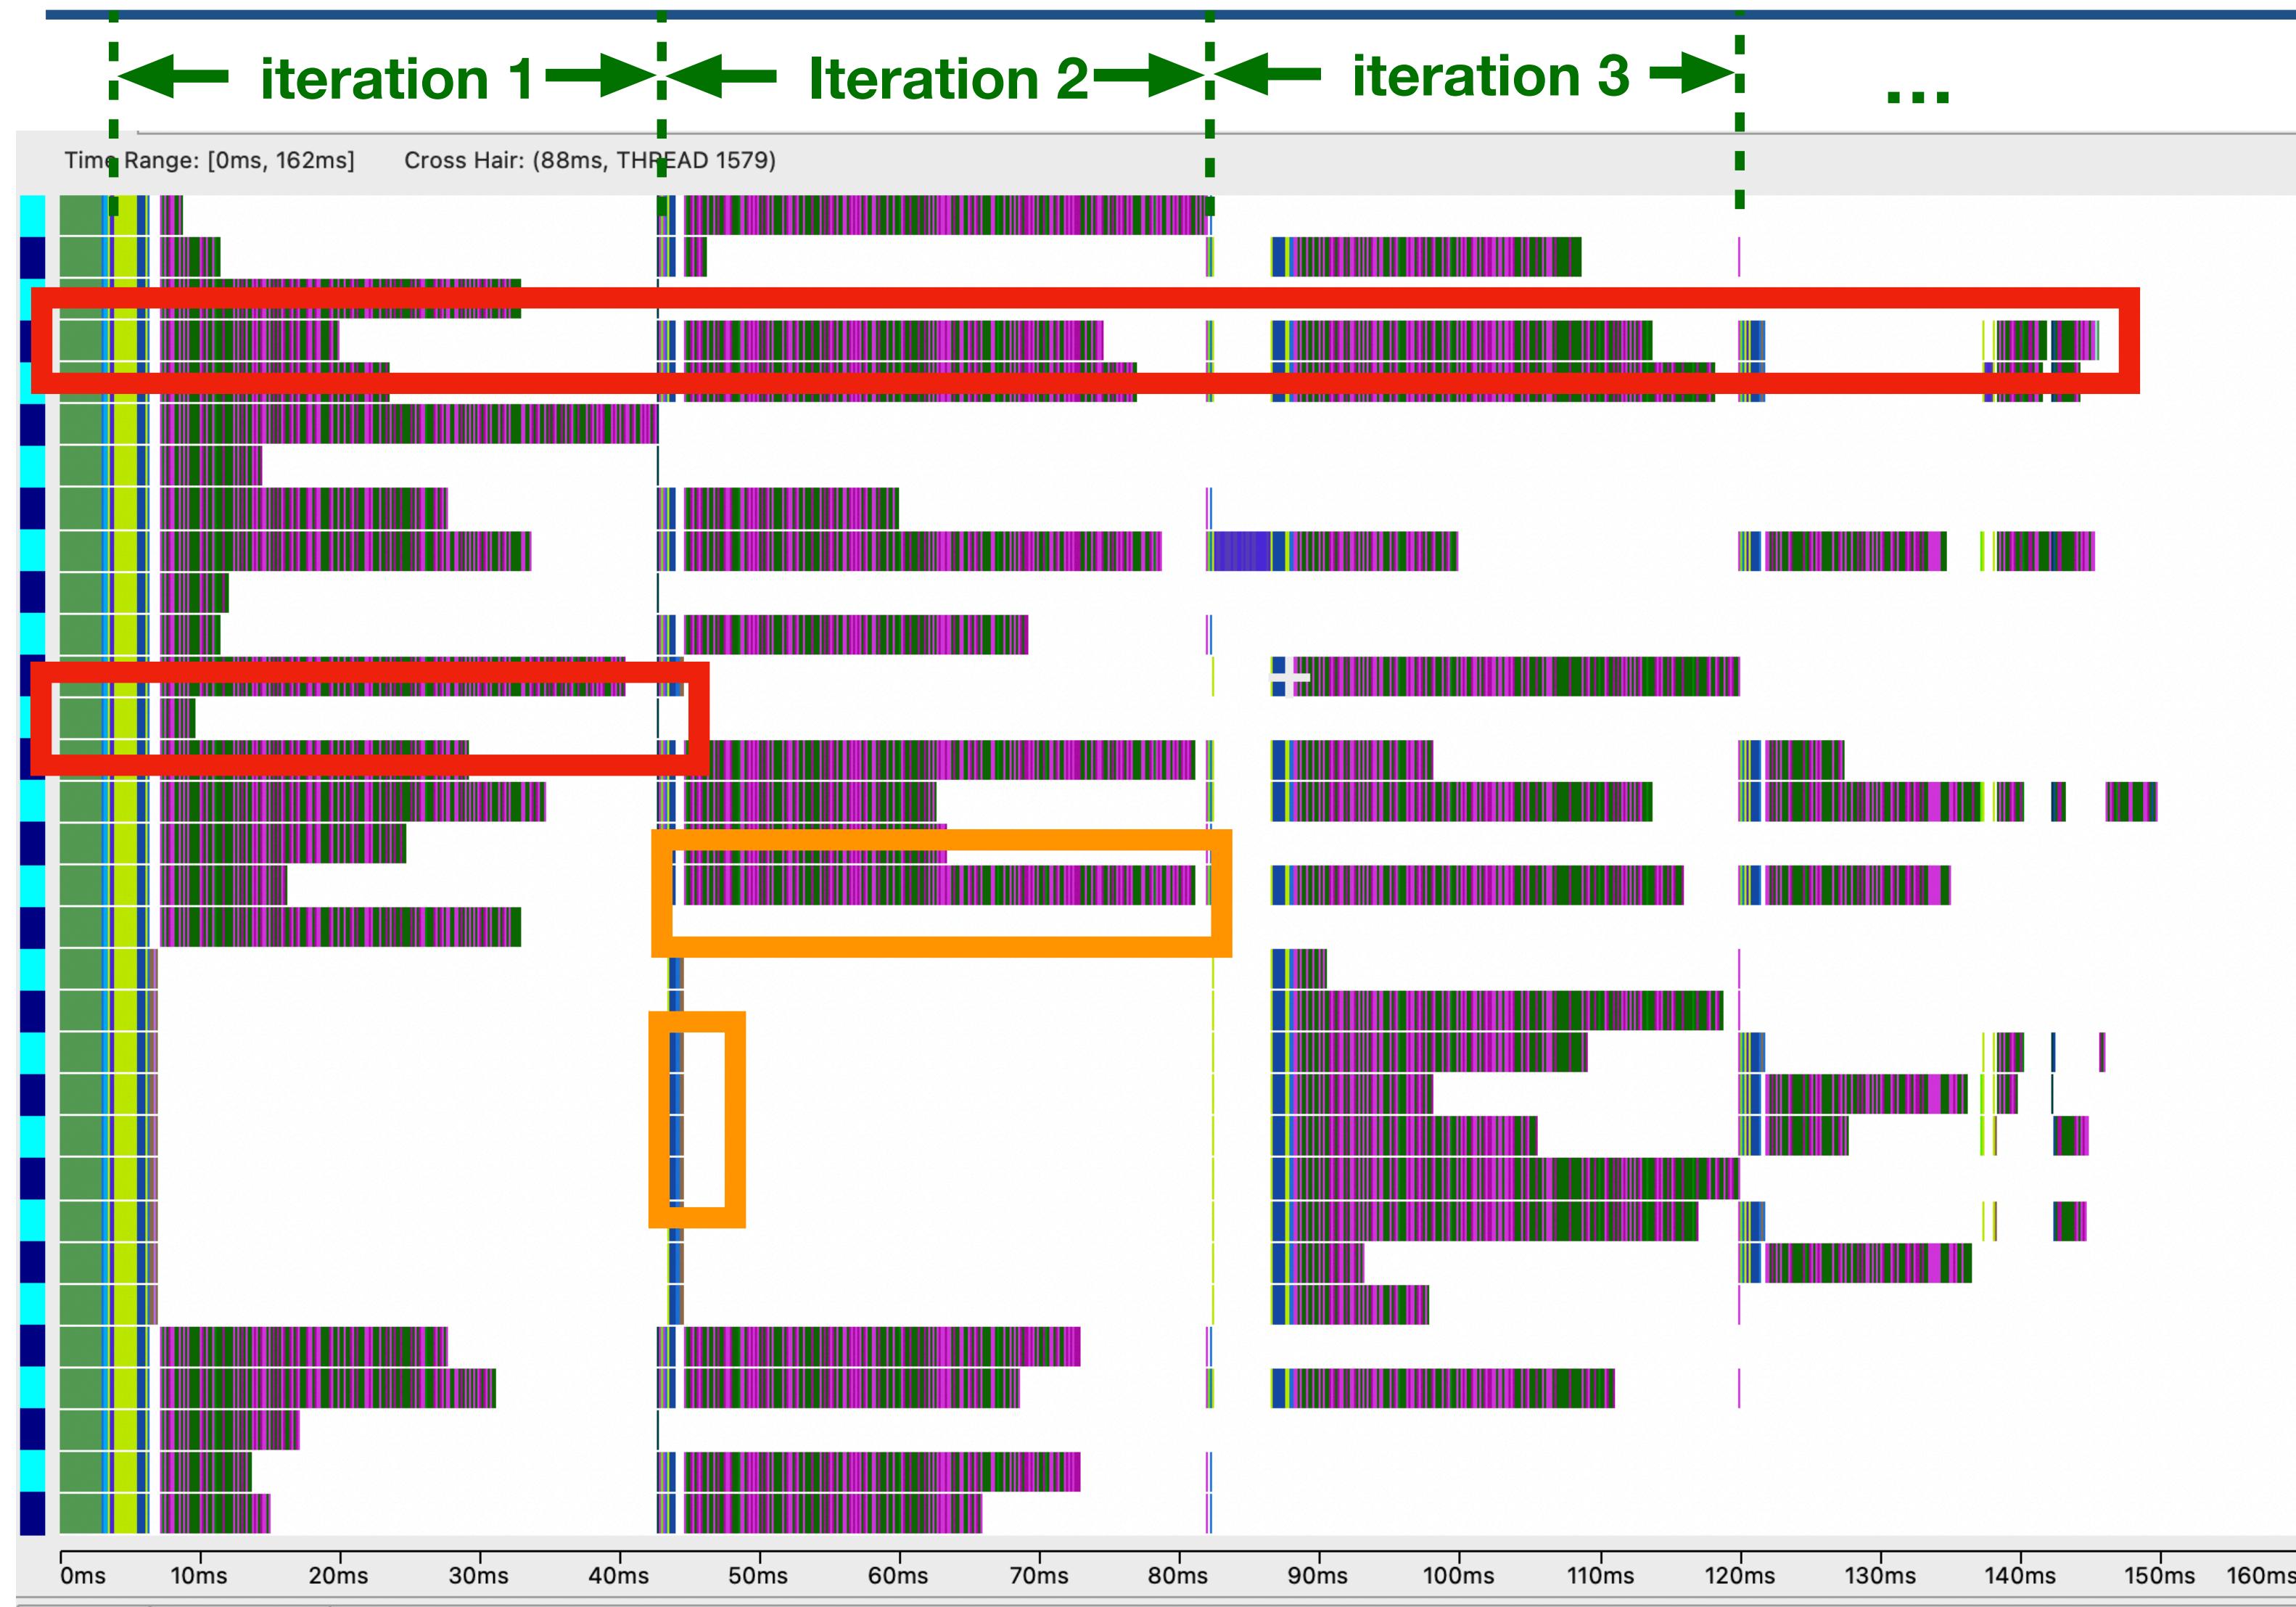
```
}
```

- Each case

- Some computation

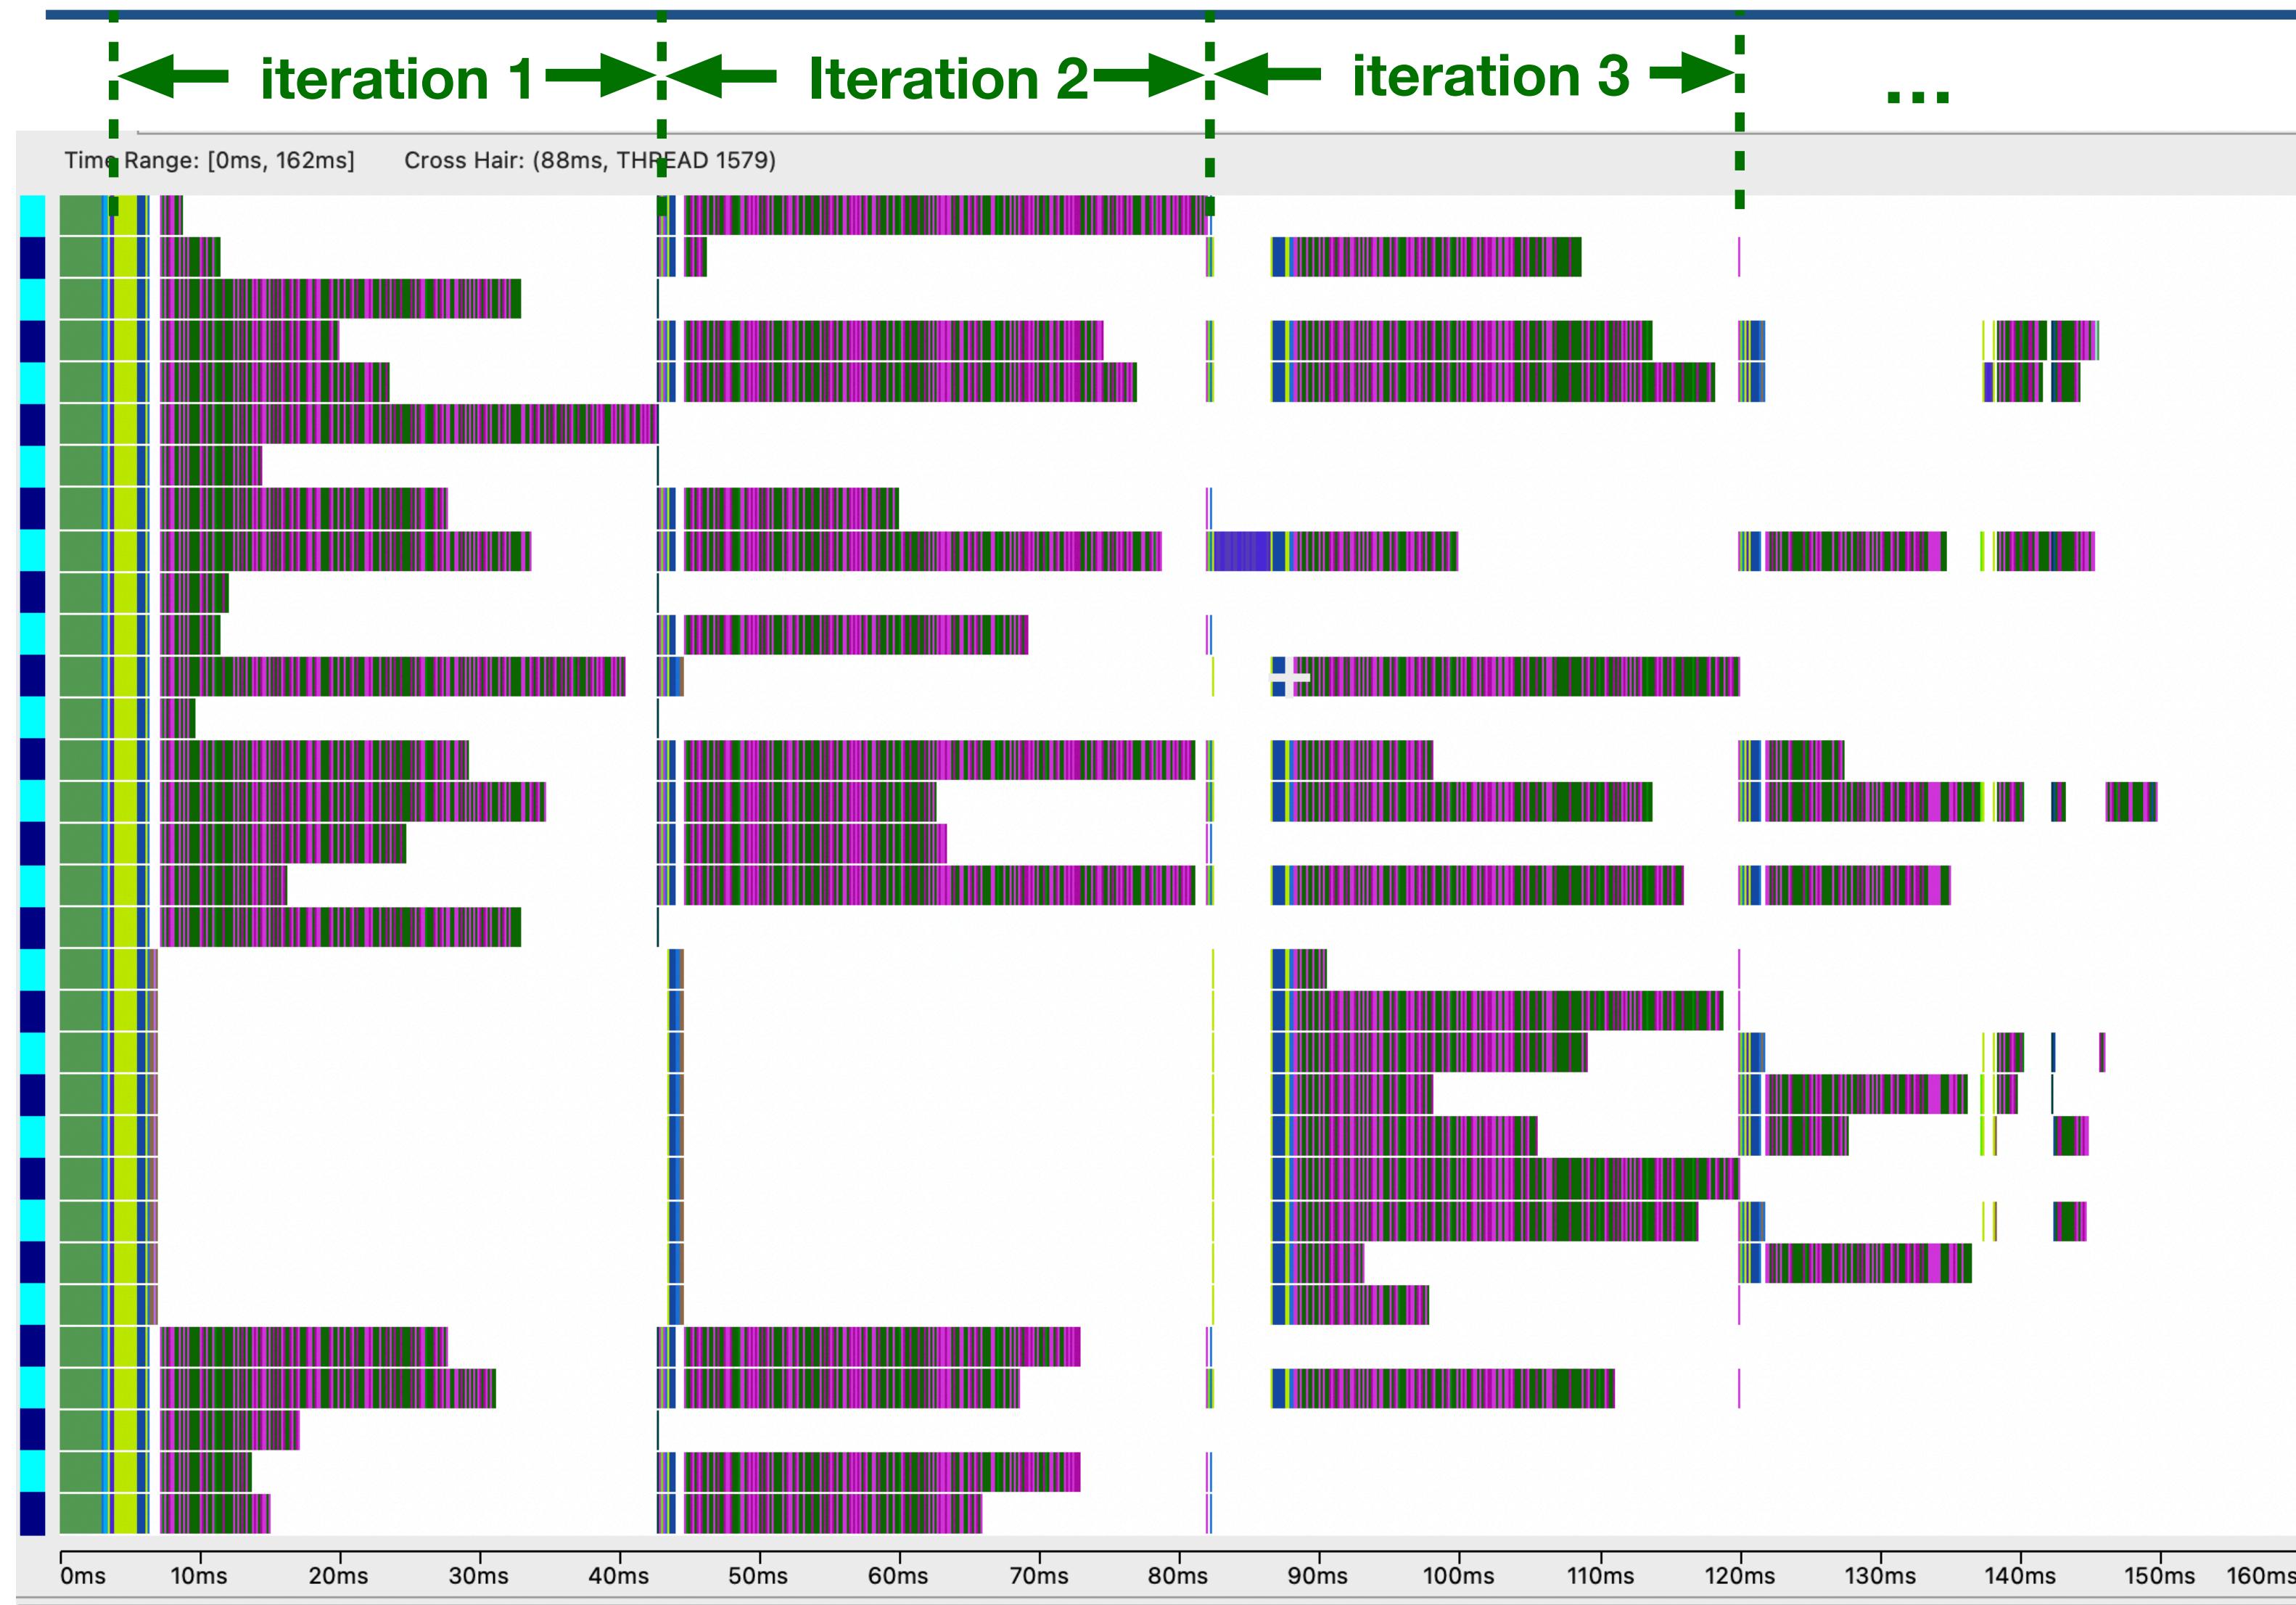
- Decide keepTracking

Quicksilver in Our Tool



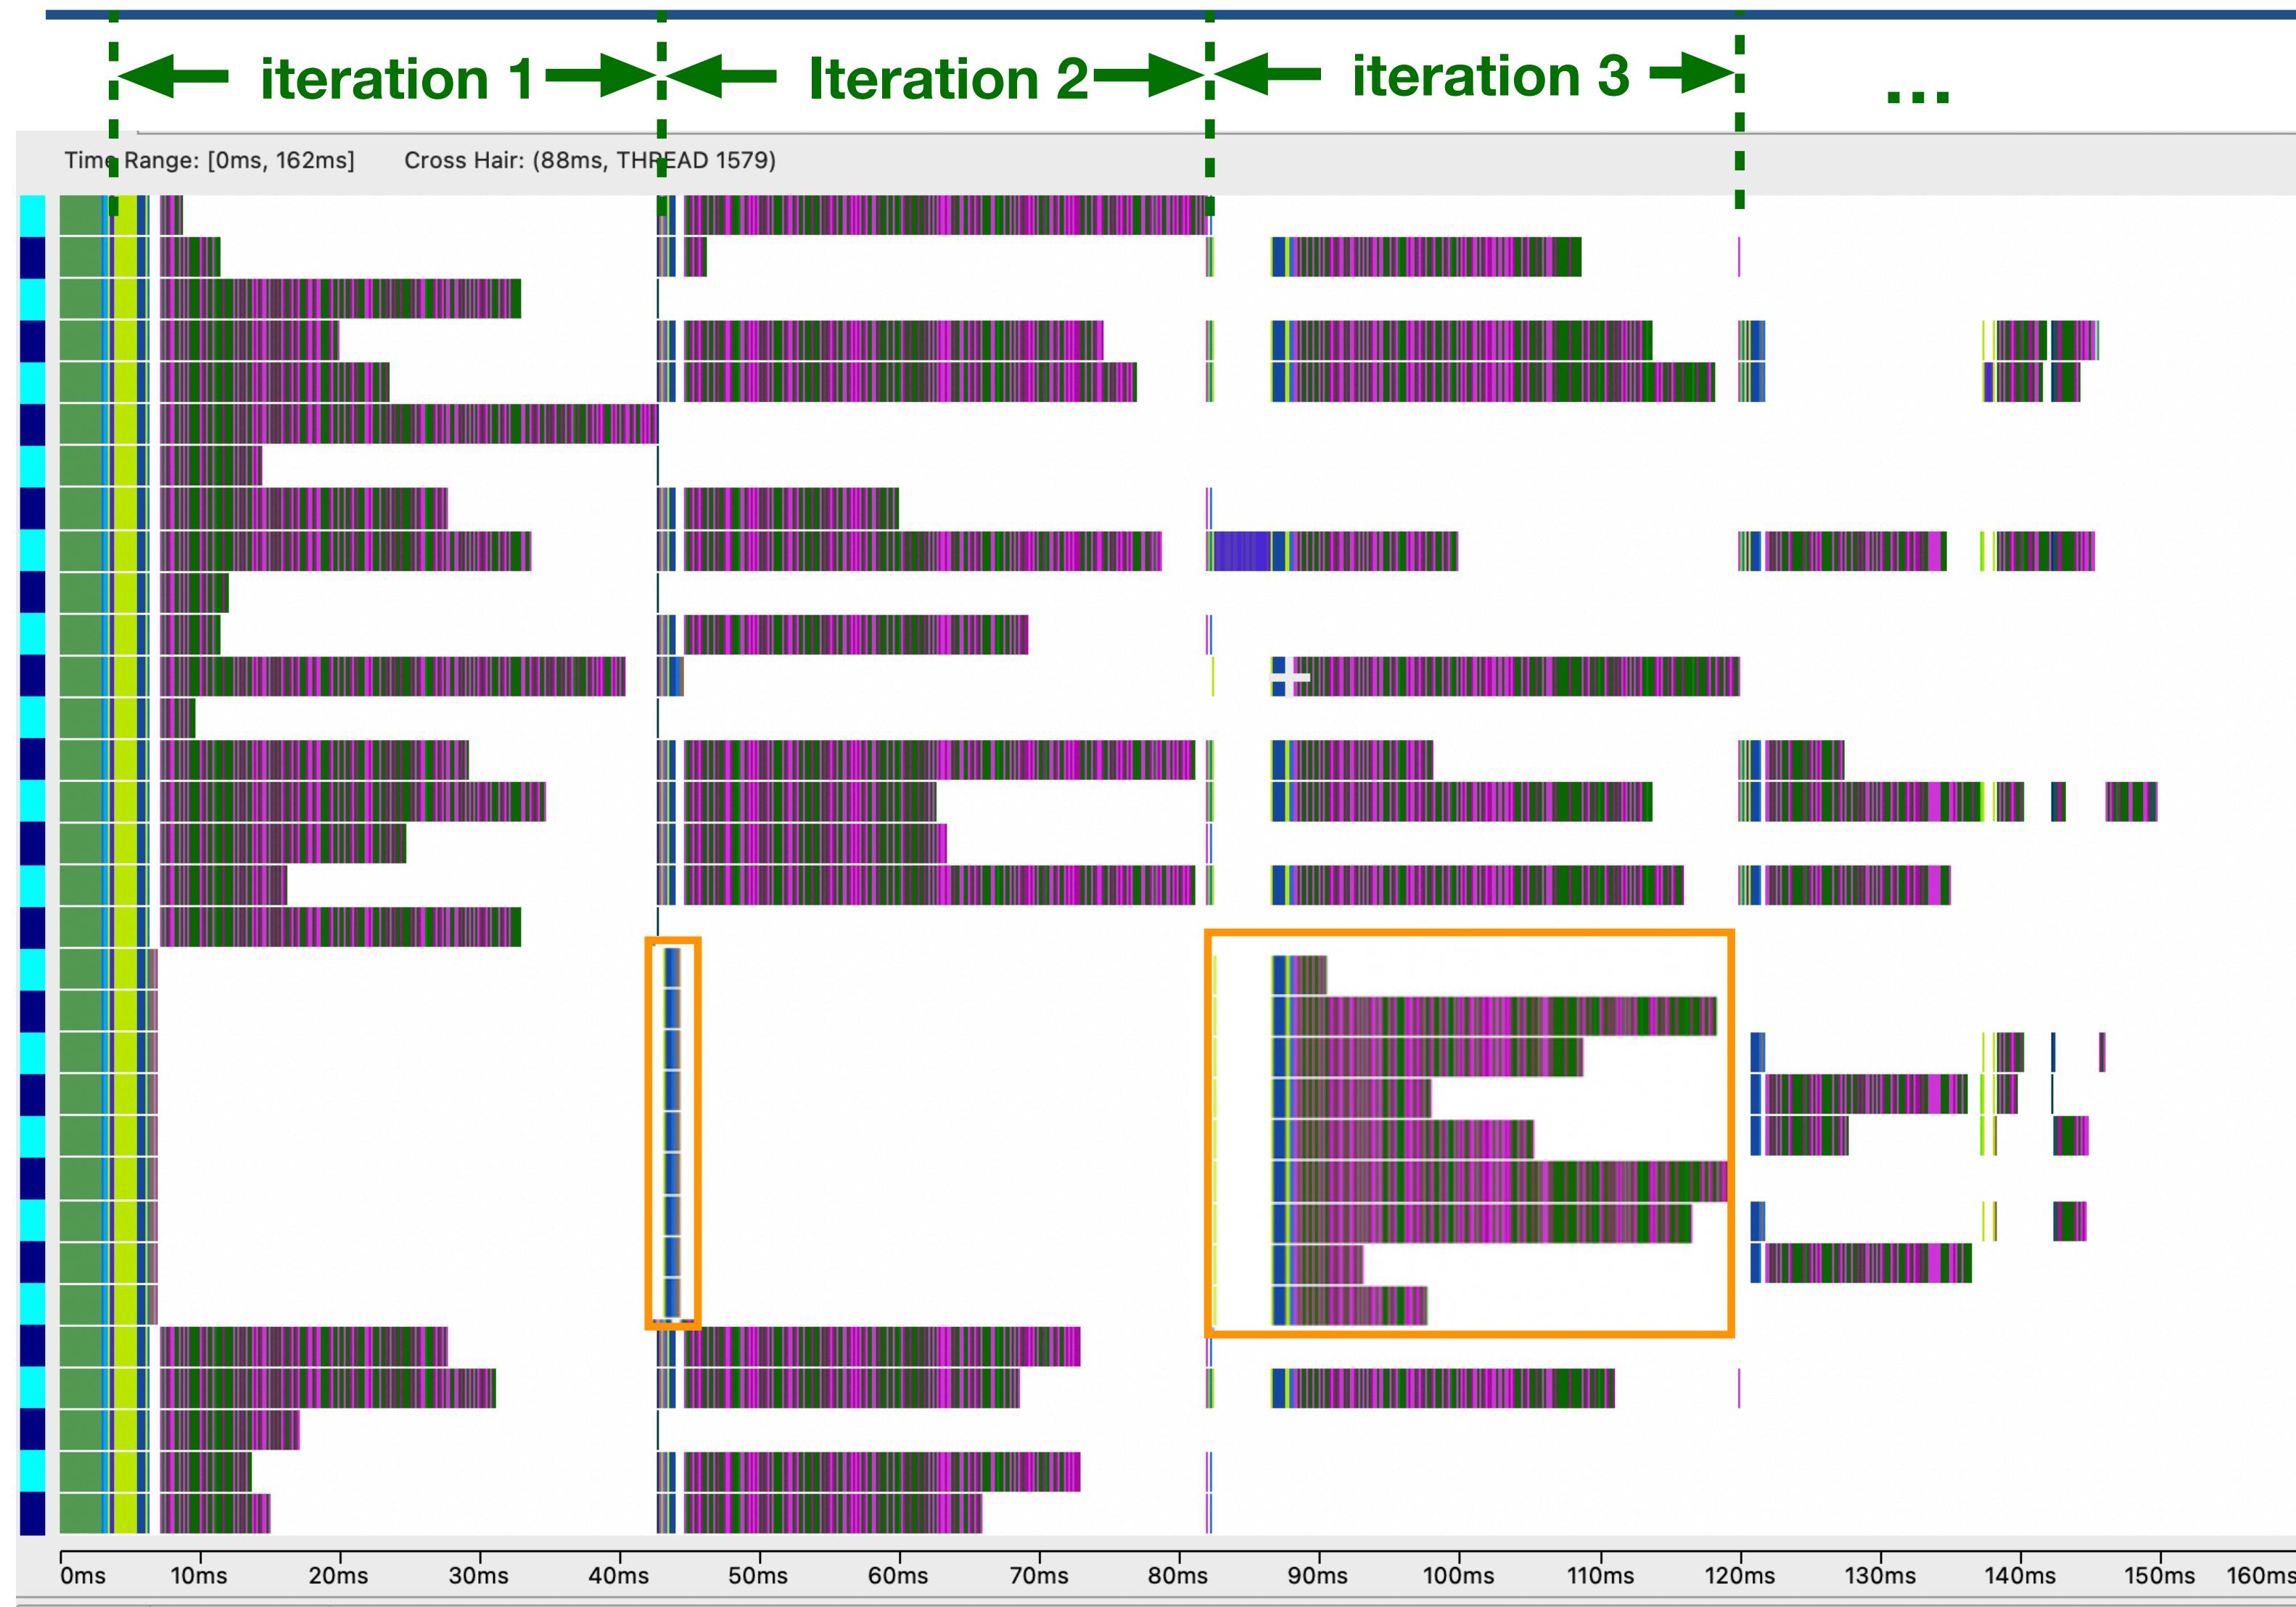
- Two imbalance
 - Imbalance **between** tasks
 - Imbalance between cases **within** each iteration

Quicksilver in Our Tool



- Two imbalance
 - Imbalance **between** tasks
 - Imbalance between cases **within** each iteration
 - Two optimization opportunities
 - Compact tasks among threads

Quicksilver in Our Tool



- Two imbalance
 - Imbalance **between** tasks
 - Imbalance between cases **within** each iteration
- Two optimization opportunities
 - Compact tasks among threads
 - Each thread executes short cases until the next case is the long one, and all the threads work on the long case together

Outline

- Our trace tool
- Potential optimization opportunities of Quicksilver found with the trace tool
- **Exploration of an optimization opportunity**
- Work stealing algorithm among GPU threads
- Summary

Exploration with Mini Proxy App

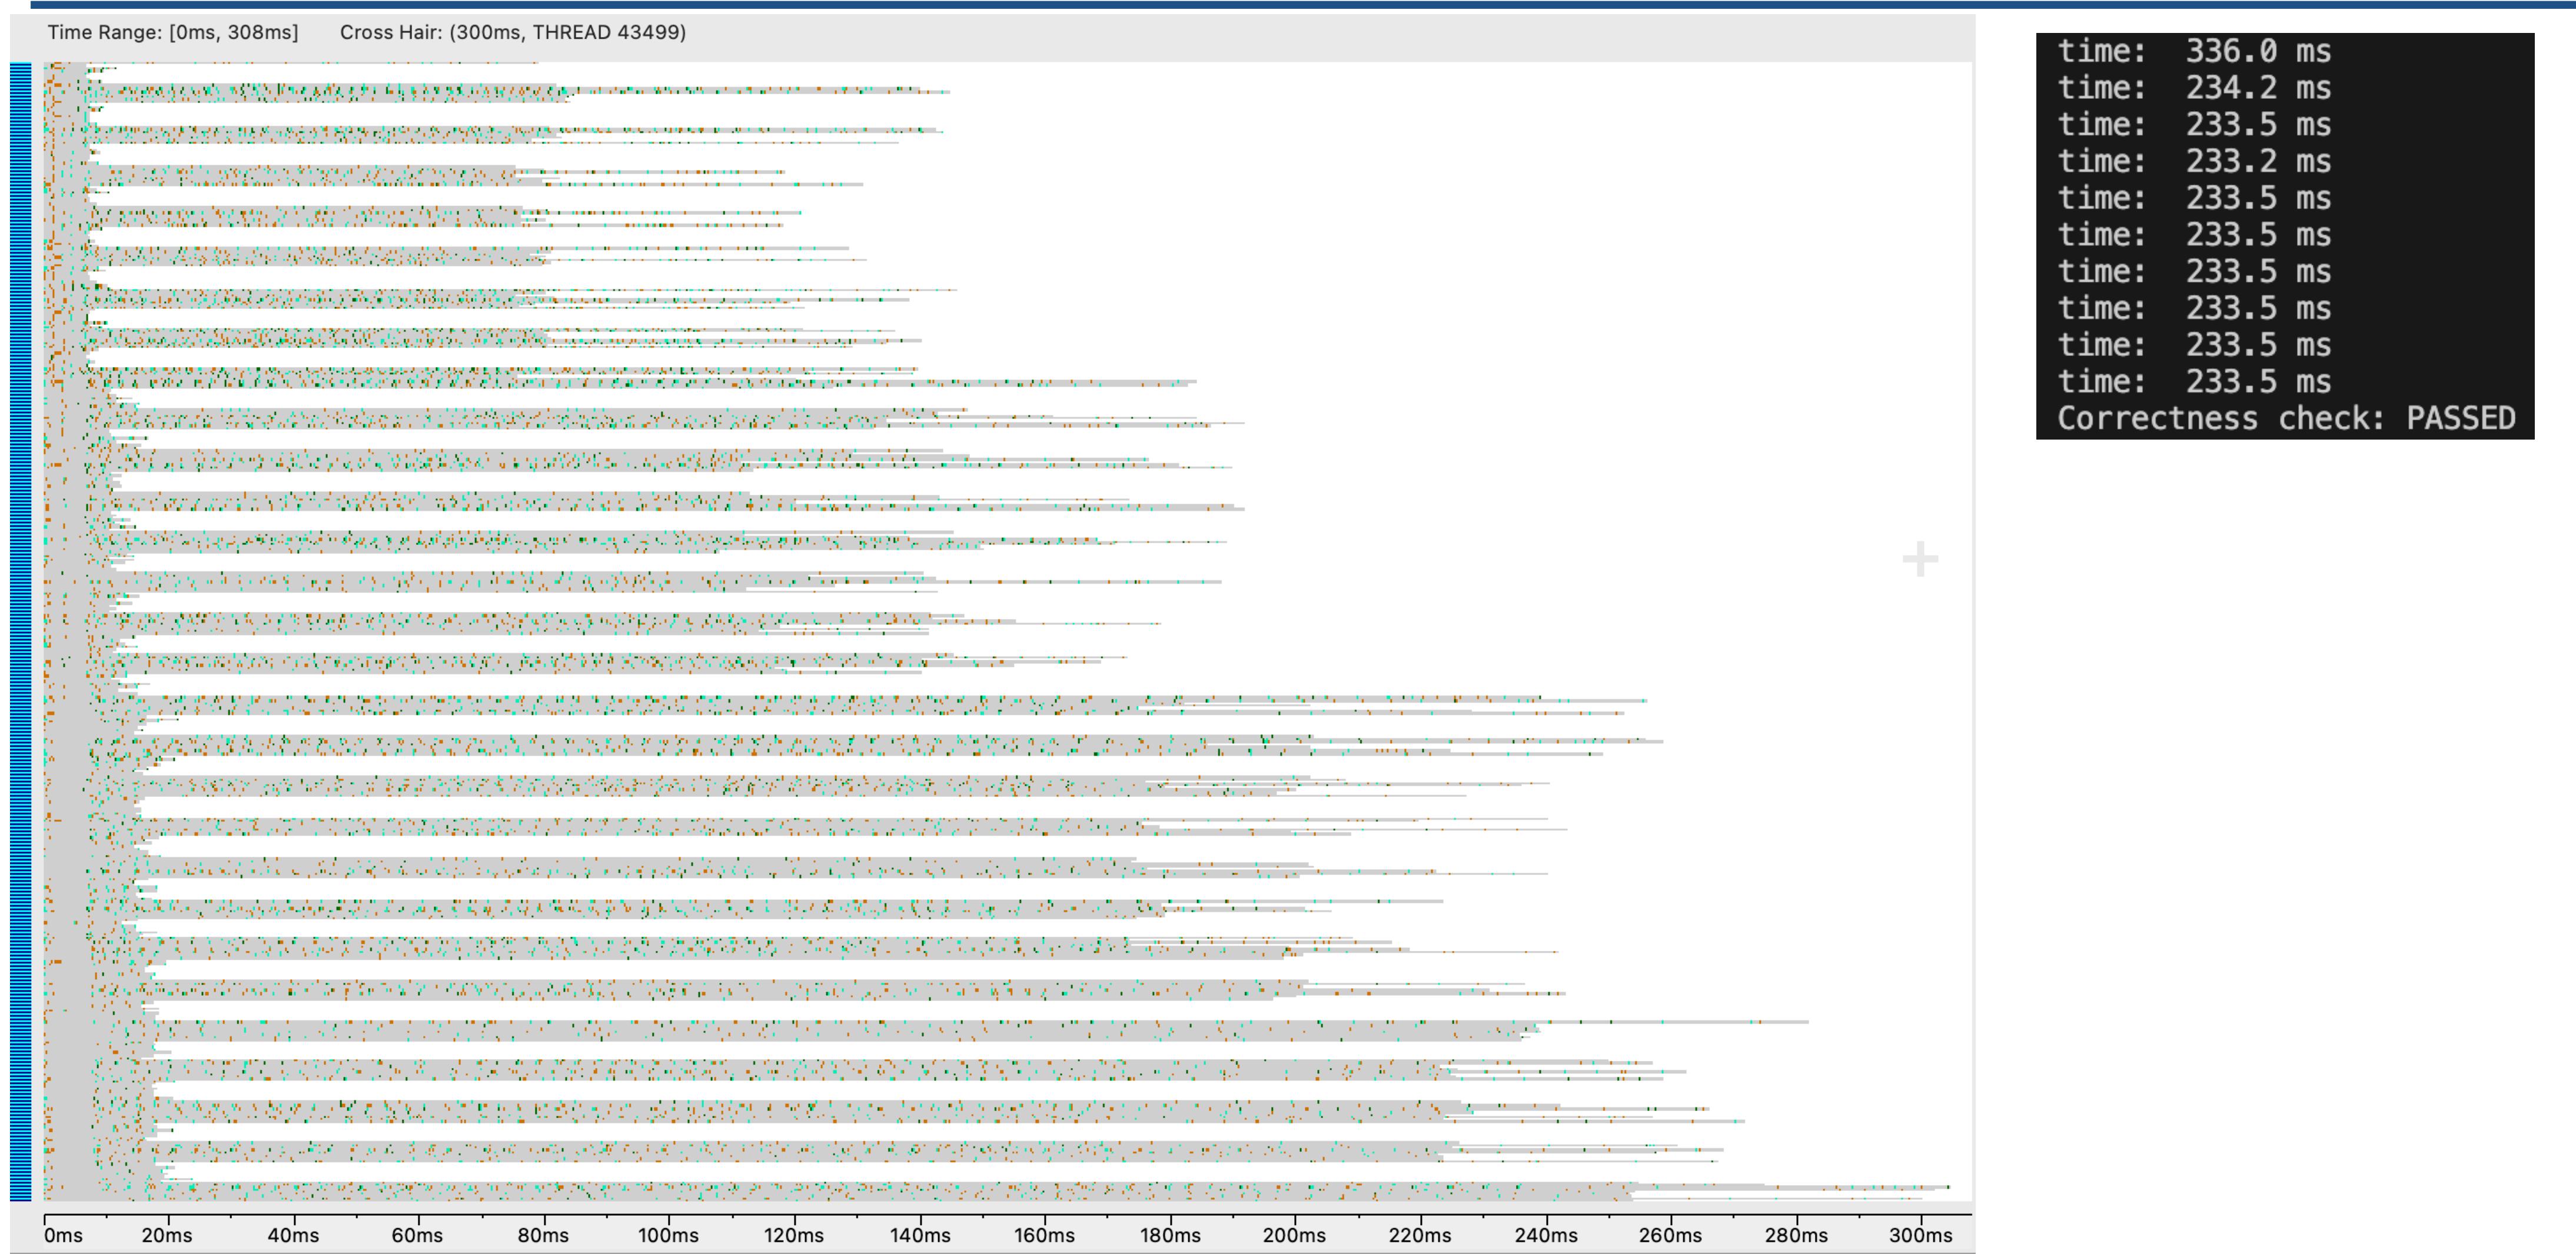
- N threads and N tasks, each thread works for one task, $N = 100,096$
- Each task can be either SHORT ($5 * 100,000$ additions) or LONG ($100 * 100,000$ additions)
- Terms:
 - Even threads: threads with even id
 - Odd threads: threads with odd id
 - First half threads: threads with id in $[0, N/2]$
 - Second half threads: threads with id in $[N/2, N]$

Mini Proxy App - Potential Benefits

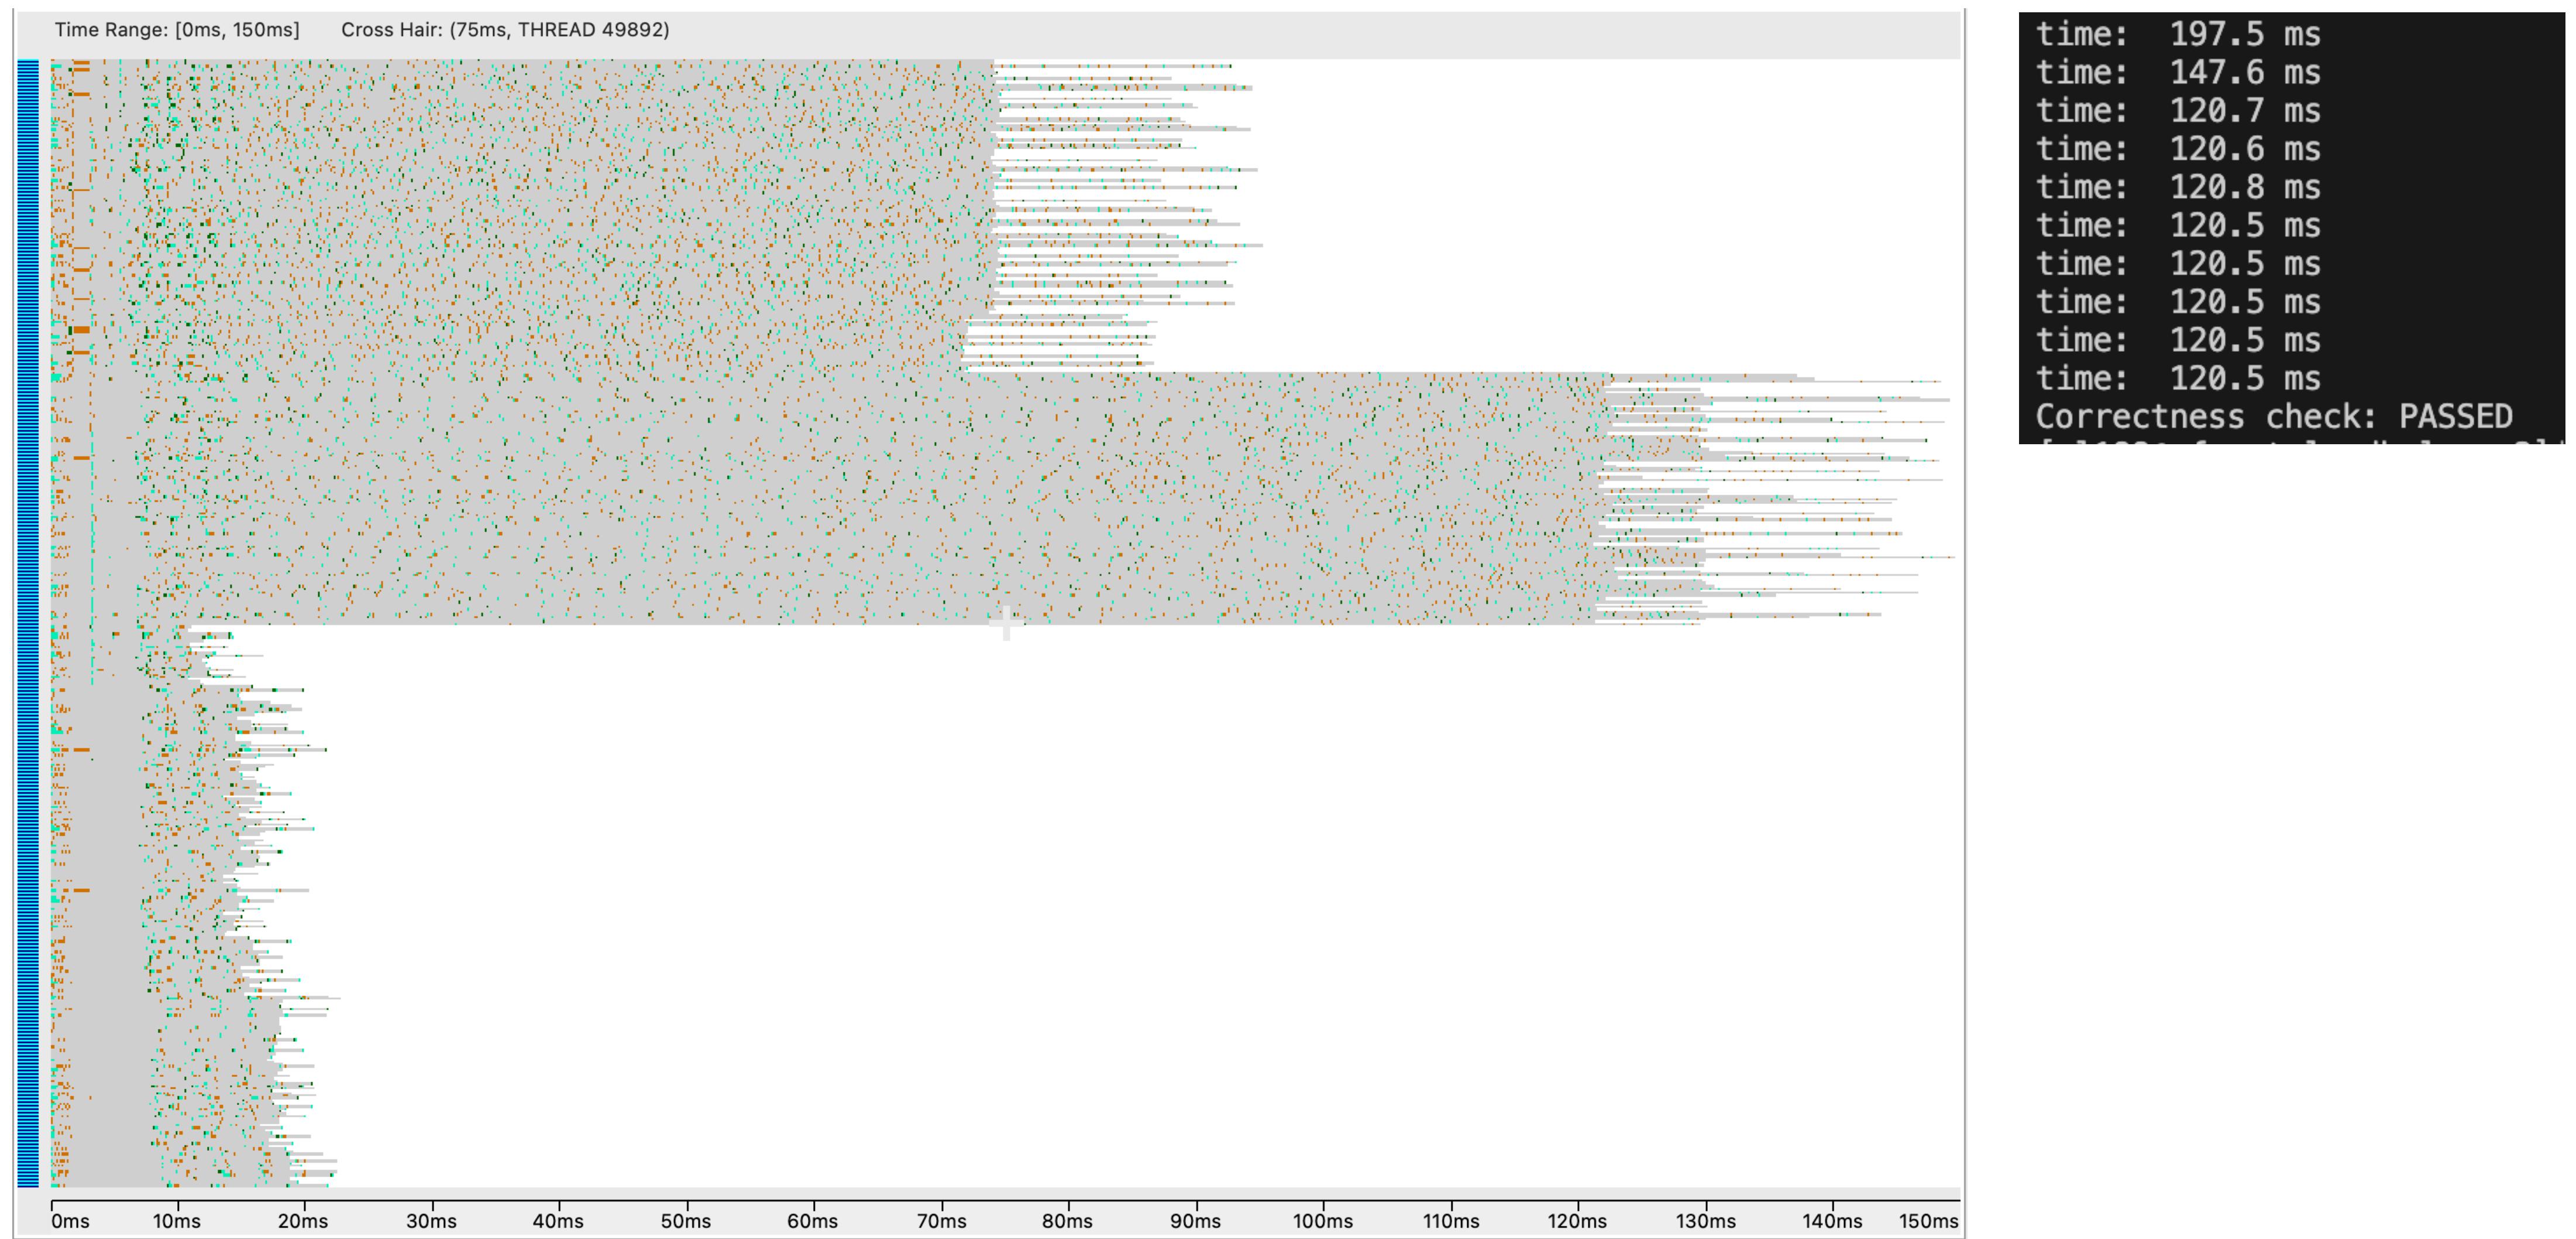
- Same amount of total workload, different threads assignment (who get what task)

Version	LONG task	SHORT task
odd-even	even threads	odd threads
long-short	first half	second half
short-long	second half	first half

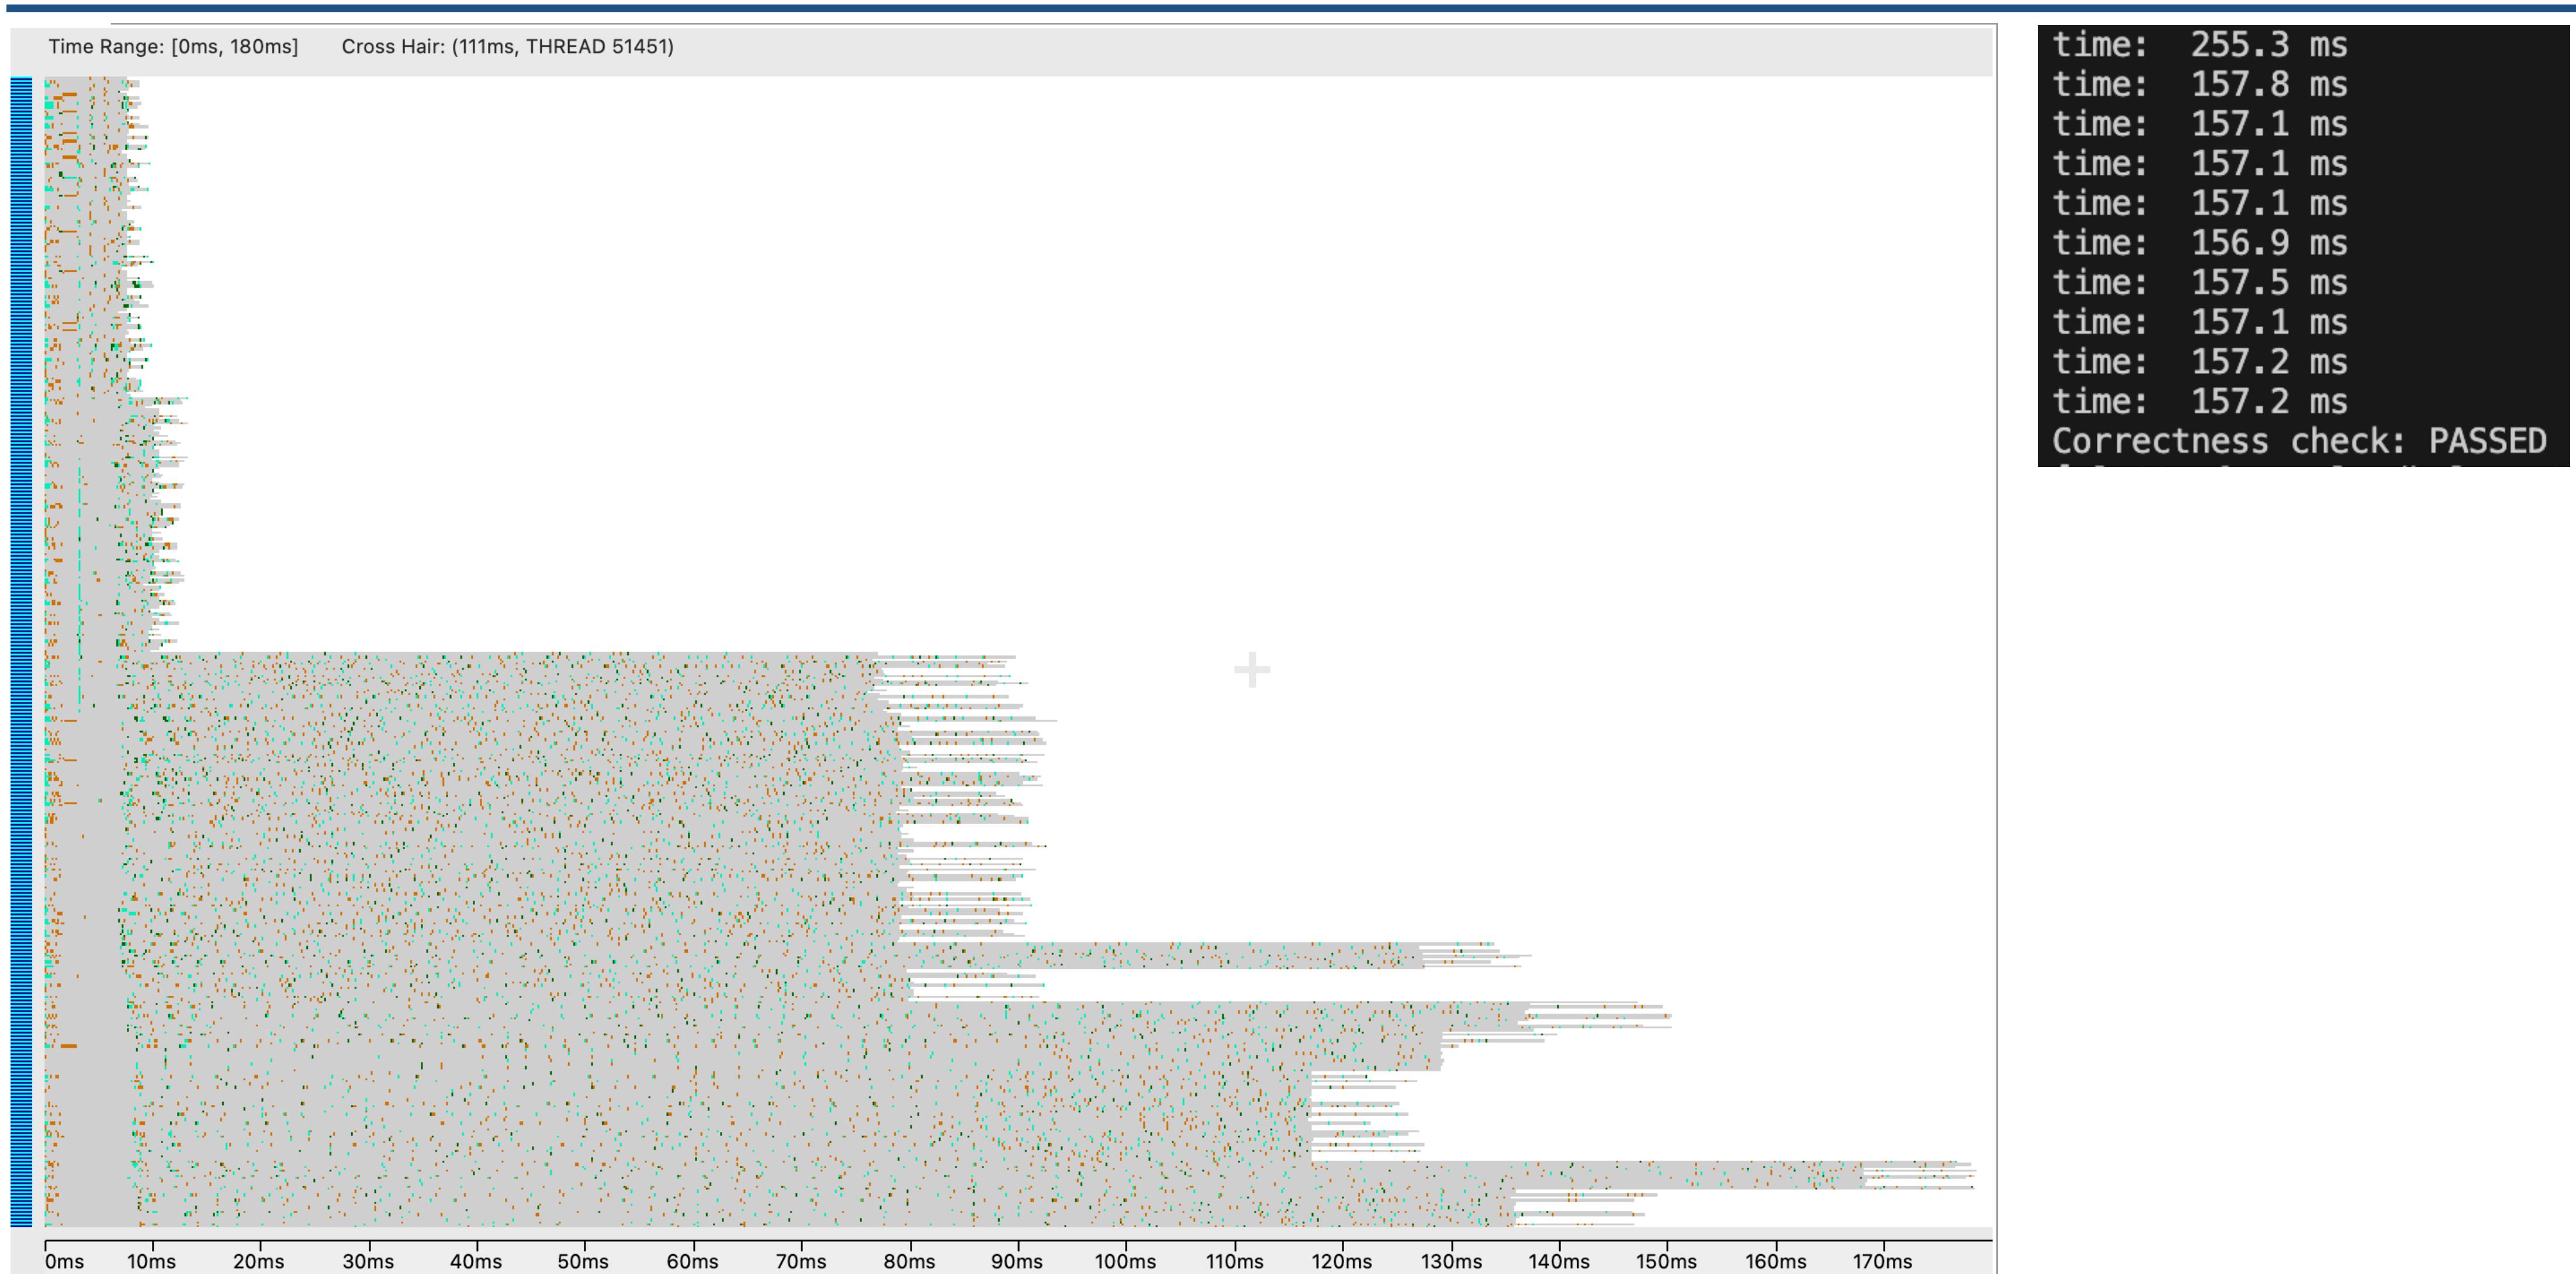
Mini Proxy App - Version odd-even



Mini Proxy App - Version long-short



Mini Proxy App - Version short-long



Compact Work is Faster

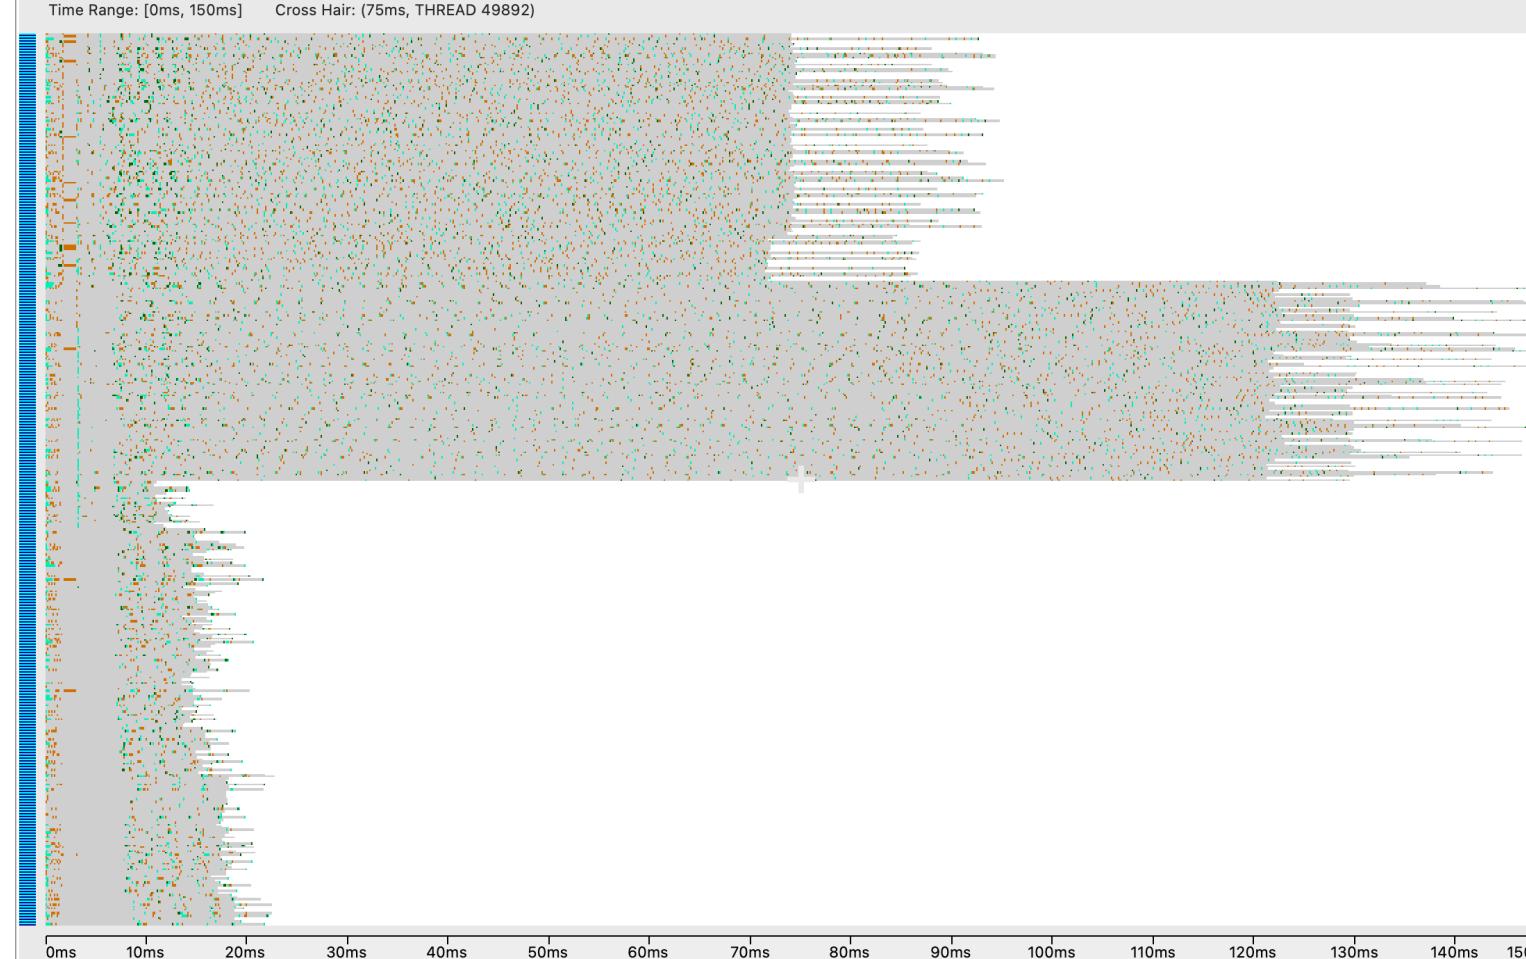
odd-even



```
time: 336.0 ms
time: 234.2 ms
time: 233.5 ms
time: 233.2 ms
time: 233.5 ms
Correctness check: PASSED
```

Slow

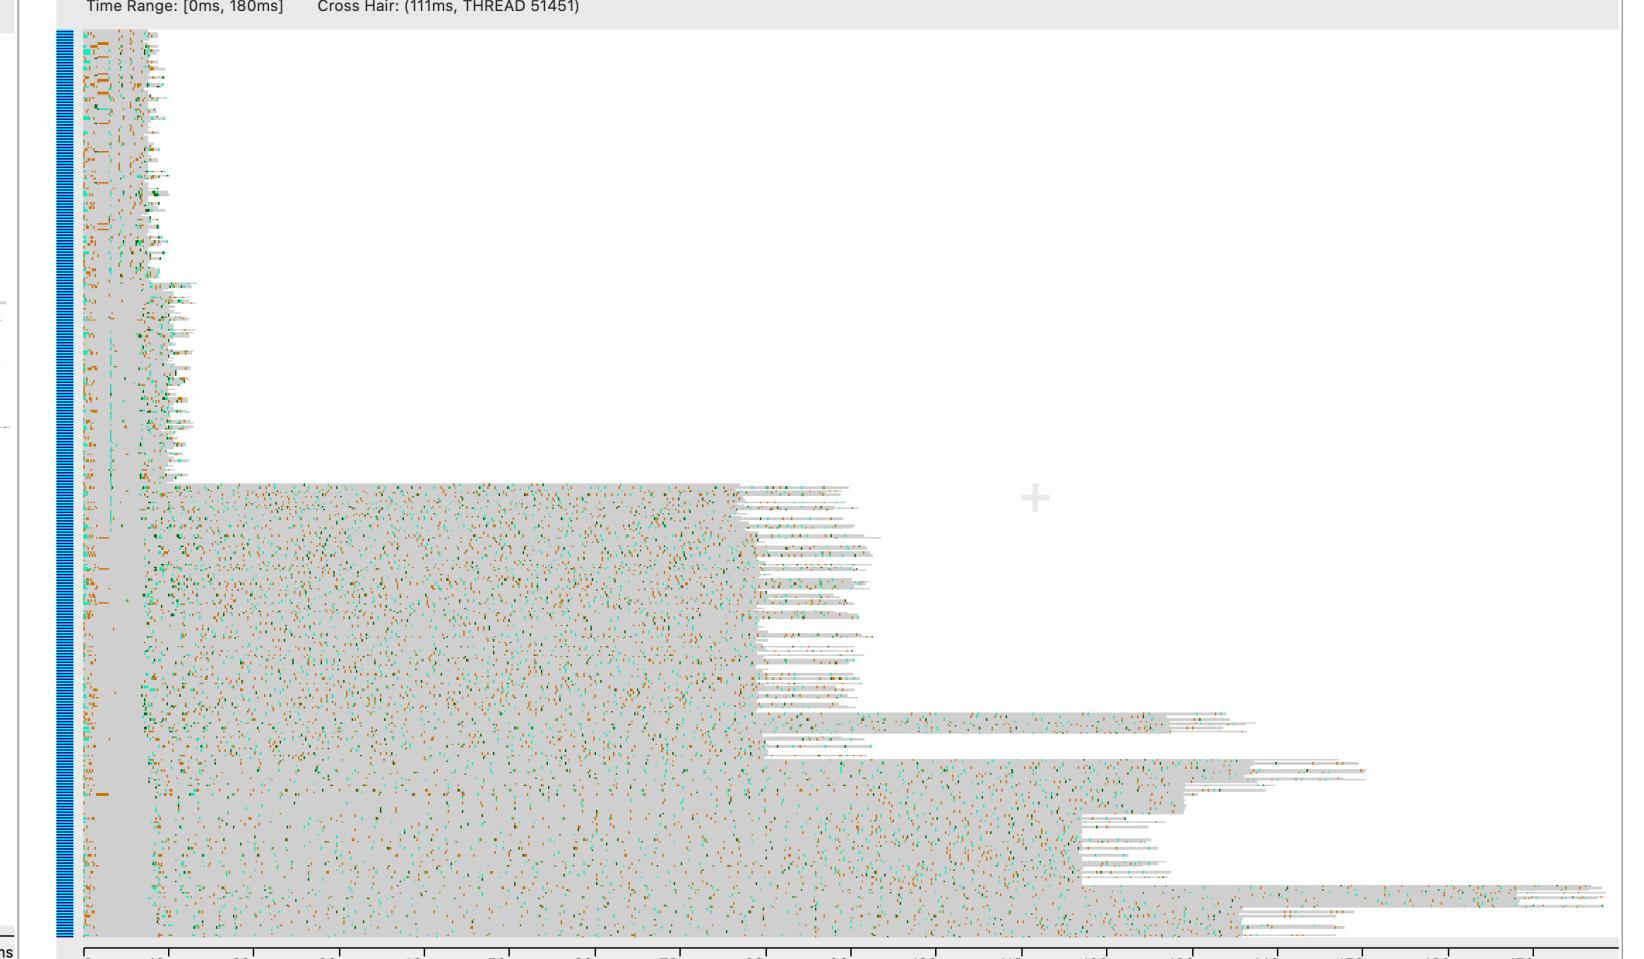
long-short



```
time: 197.5 ms
time: 147.6 ms
time: 120.7 ms
time: 120.6 ms
time: 120.8 ms
time: 120.5 ms
time: 120.5 ms
time: 120.5 ms
time: 120.5 ms
Correctness check: PASSED
```

1.94x Faster

short-long



```
time: 255.3 ms
time: 157.8 ms
time: 157.1 ms
time: 157.1 ms
time: 157.1 ms
time: 156.9 ms
time: 157.5 ms
time: 157.1 ms
time: 157.2 ms
time: 157.2 ms
Correctness check: PASSED
```

1.49x Faster

Outline

- Our trace tool
- Potential optimization opportunities of Quicksilver found with the trace tool
- Exploration of an optimization opportunity
- **Work stealing algorithm among GPU threads**
- Summary

Work Stealing Among GPU Threads

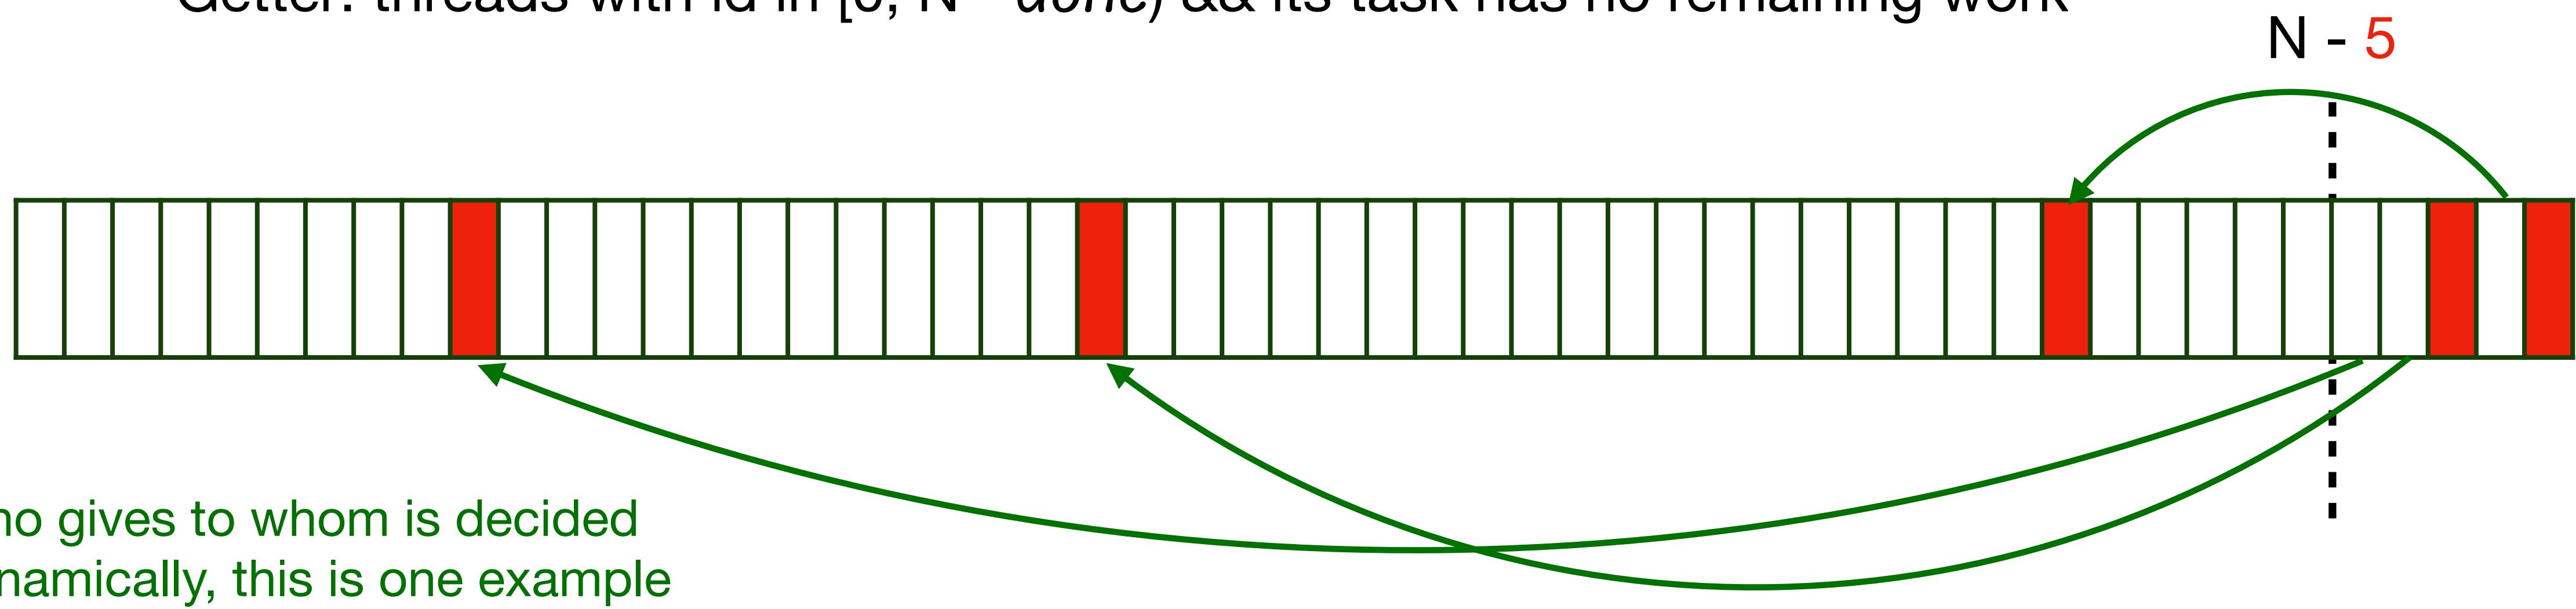
- In Quicksilver, unknown workload imbalance between threads motivates dynamic work stealing
- Pay attention
 - Avoid thread-level spin loop
 - Otherwise, one thread spinning may block the rest in its warp
 - Use `__syncwarp()`
 - Consider the scenario that not all the threads start at the same time

Work Stealing Among GPU Threads

- Each thread starts with its own task, and in each iteration
 - Step 1: works on its task
 - Step 2: decides if it should
 - give another thread its task
 - get another thread's task
 - exit
 - nothing
 - Step 3: behaves according to the decision from Step 2
 - Step 4: `__syncwarp()`

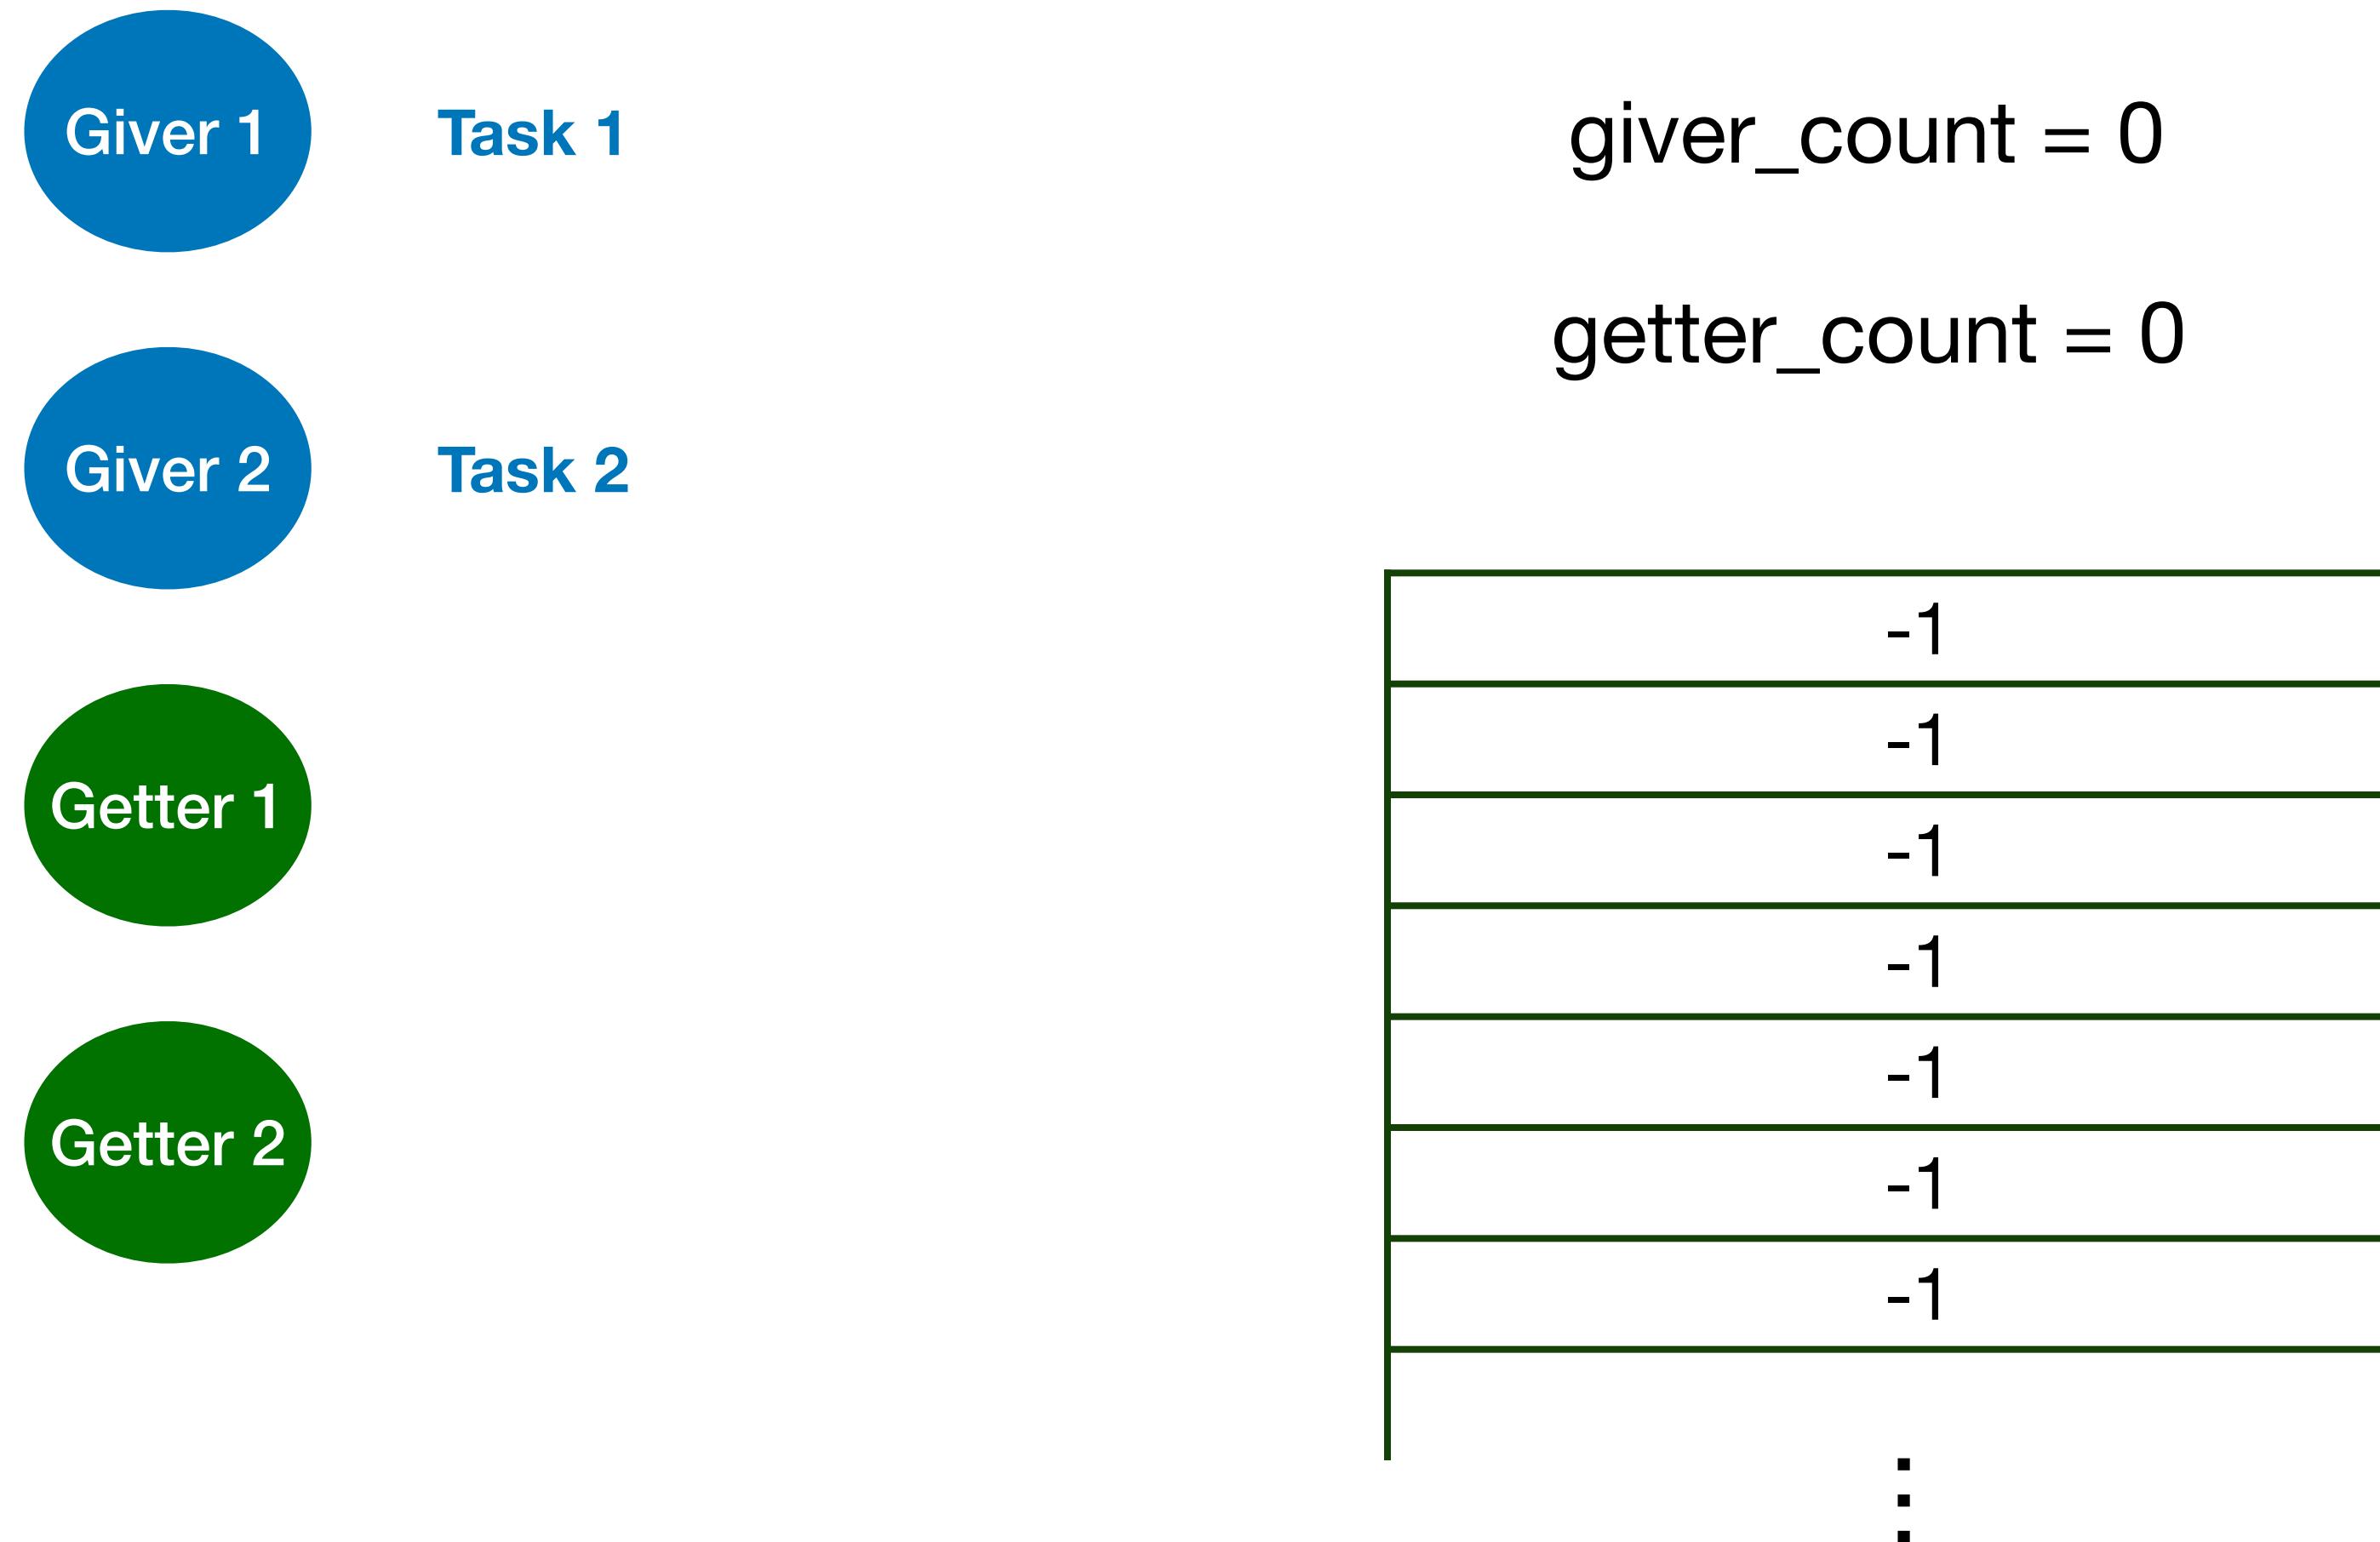
Work Stealing Among GPU Threads

- How to decide if the thread is a giver or getter or none?
 - Global variable *done*: number of finished tasks
 - Giver: threads with id in $[N - done, N)$ $\&\&$ its task still has remaining work
 - Getter: threads with id in $[0, N - done)$ $\&\&$ its task has no remaining work



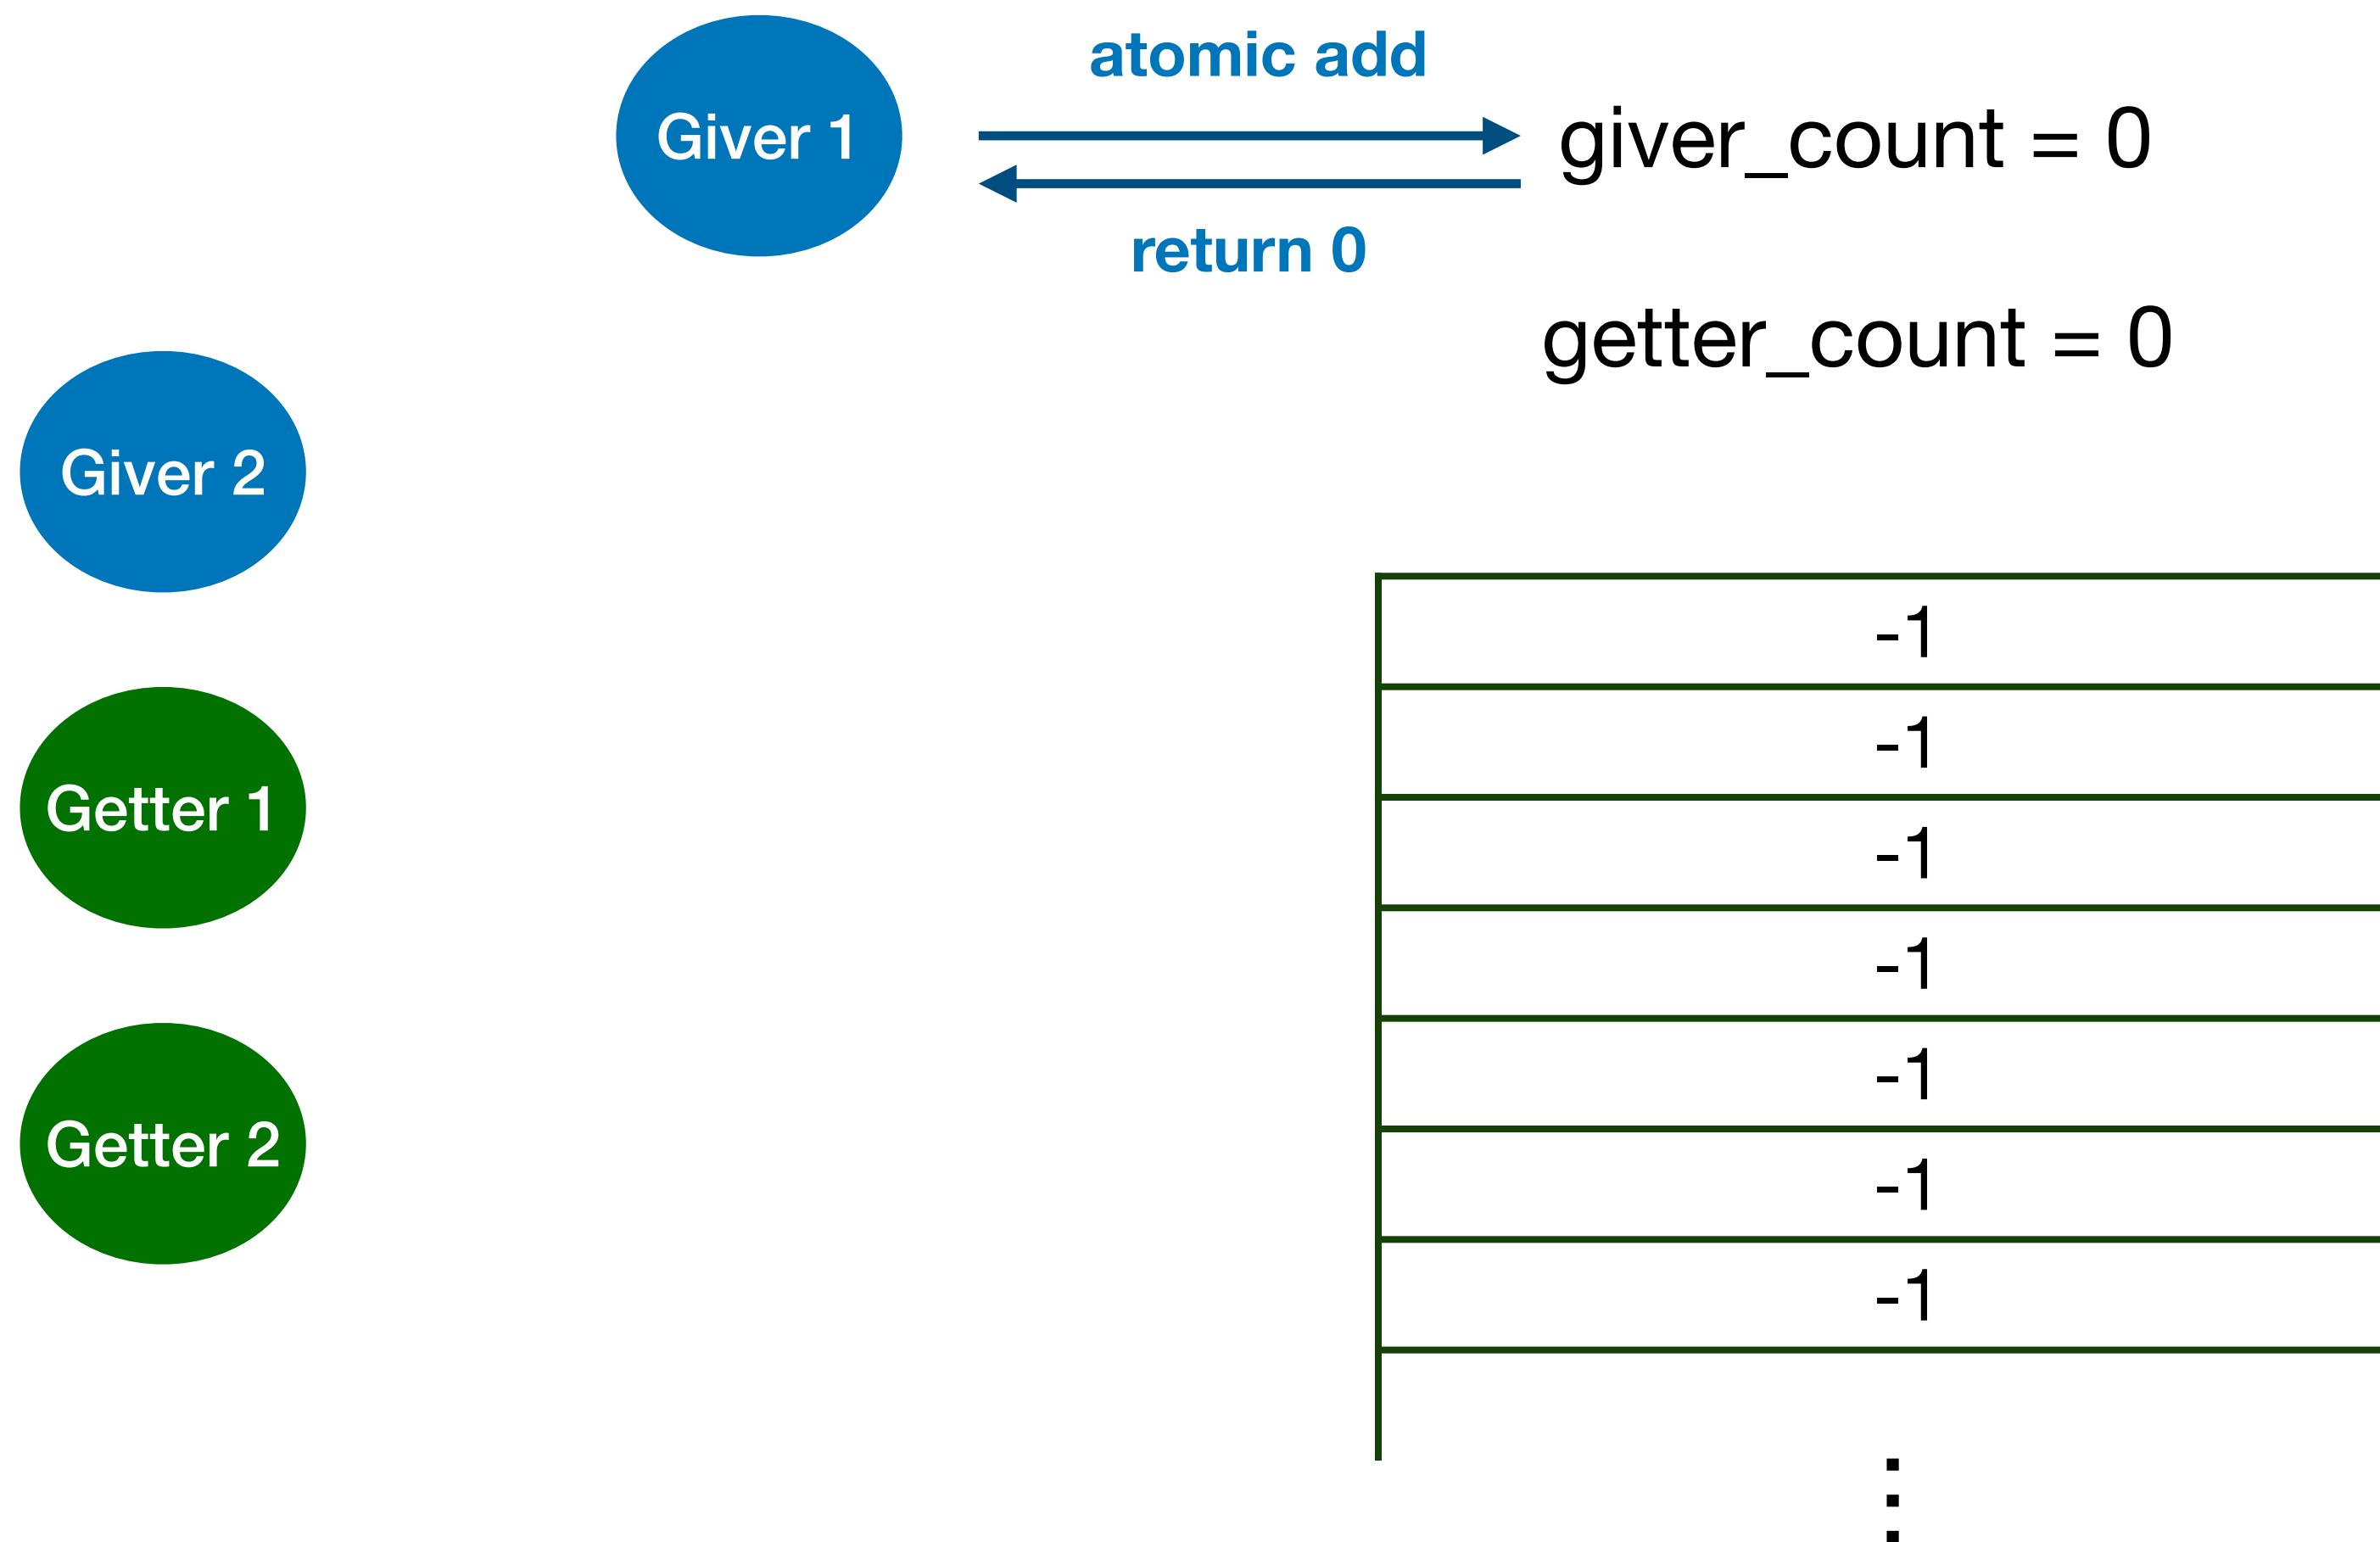
Work Stealing Among GPU Threads

- How to give or get? **Use atomic swap**



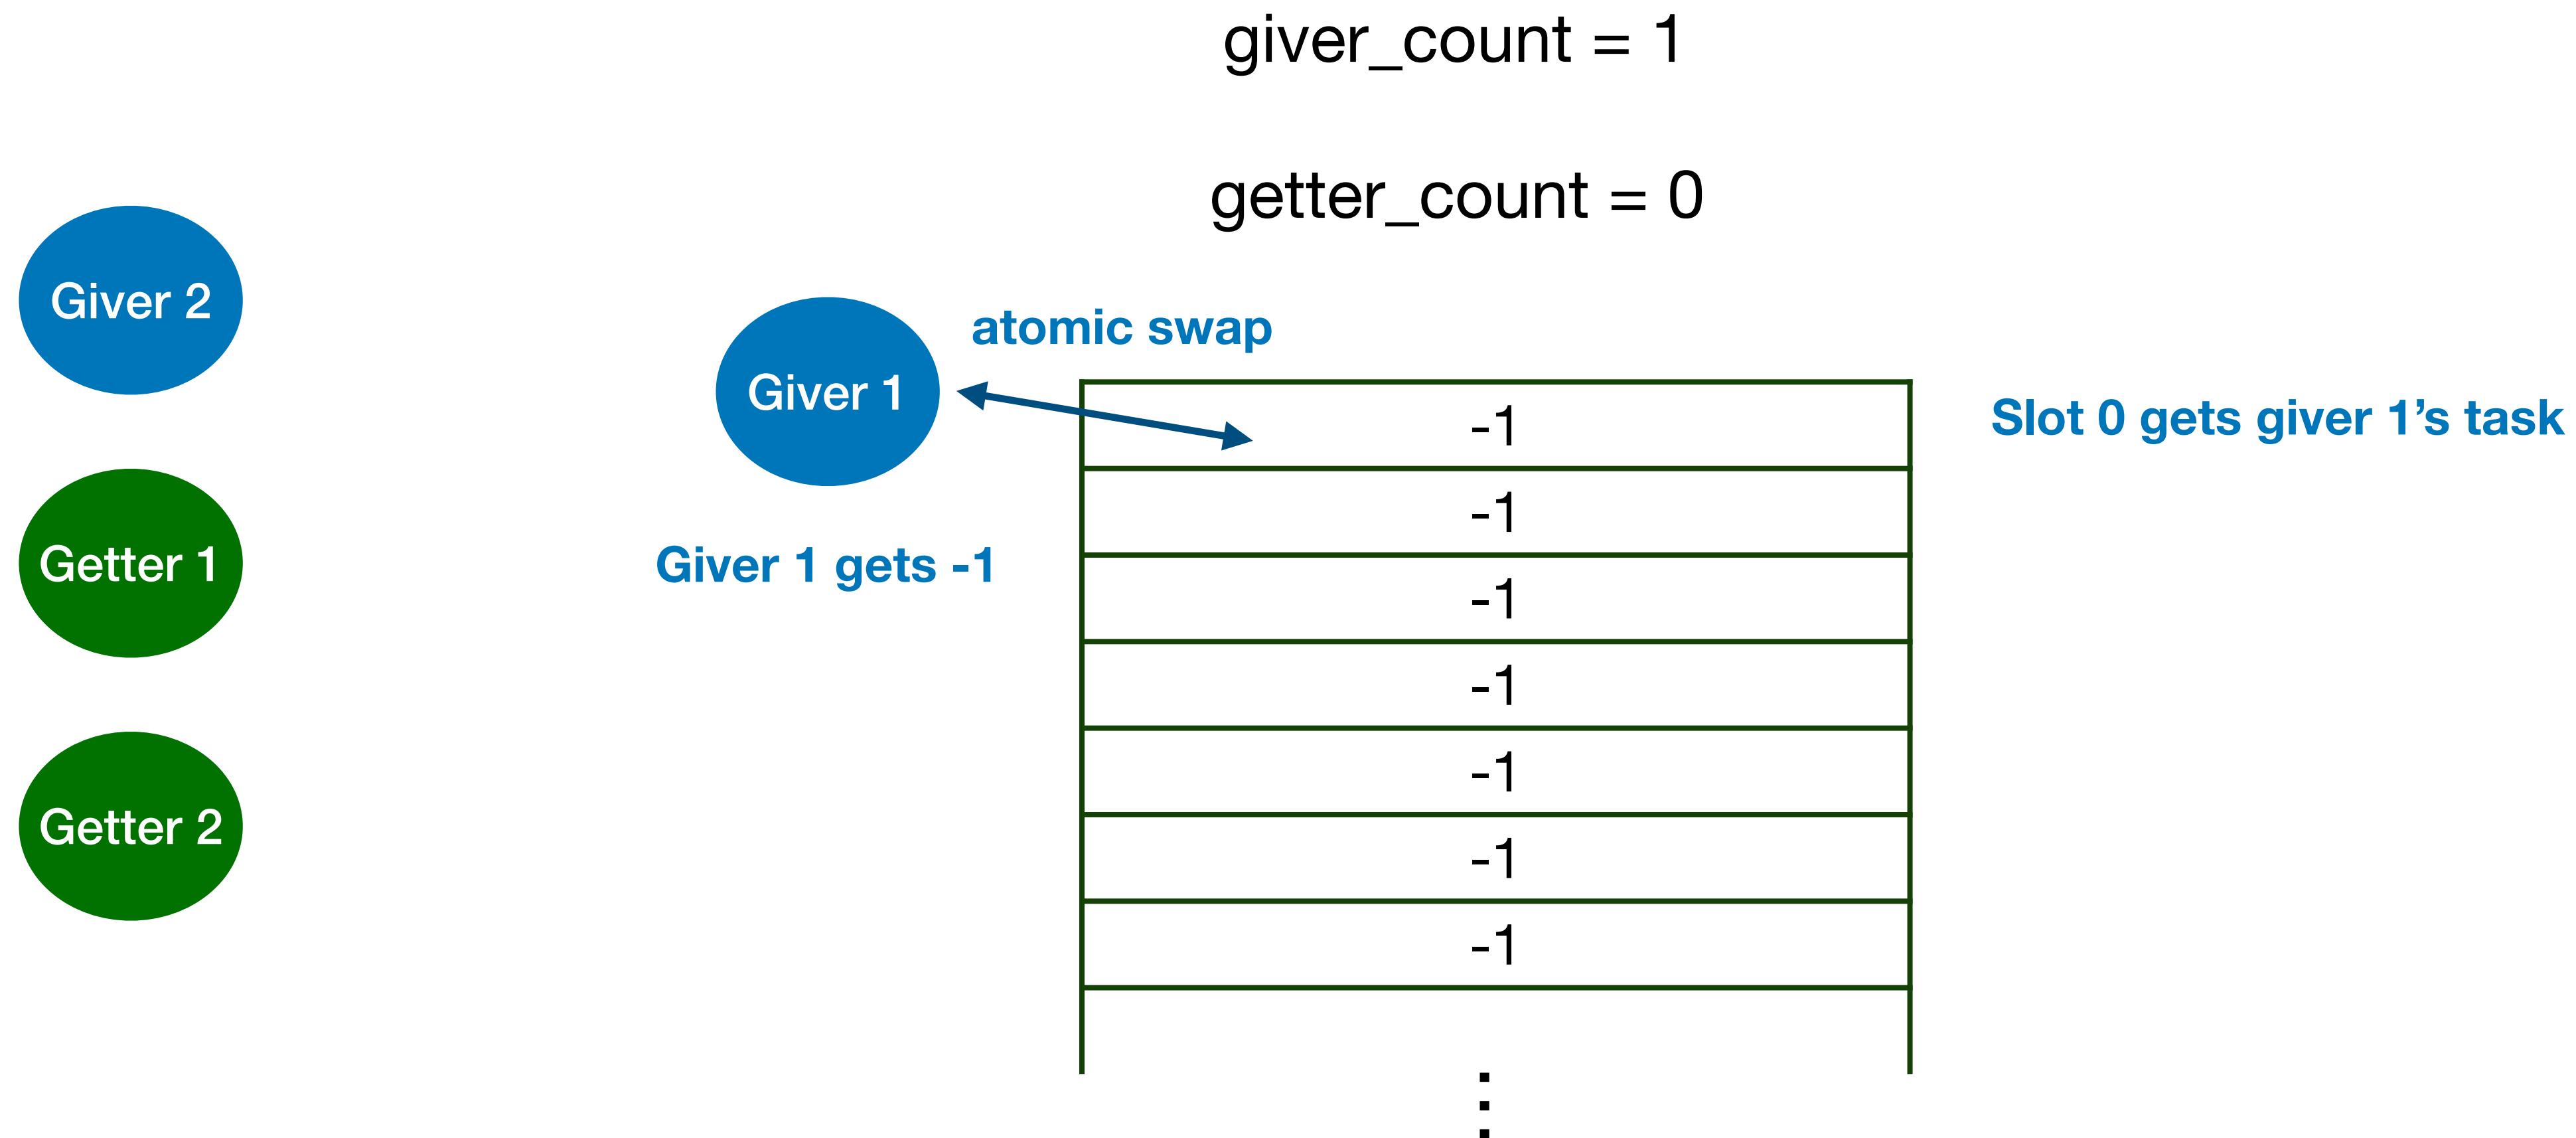
Work Stealing Among GPU Threads

- How to give or get? **Use atomic swap**



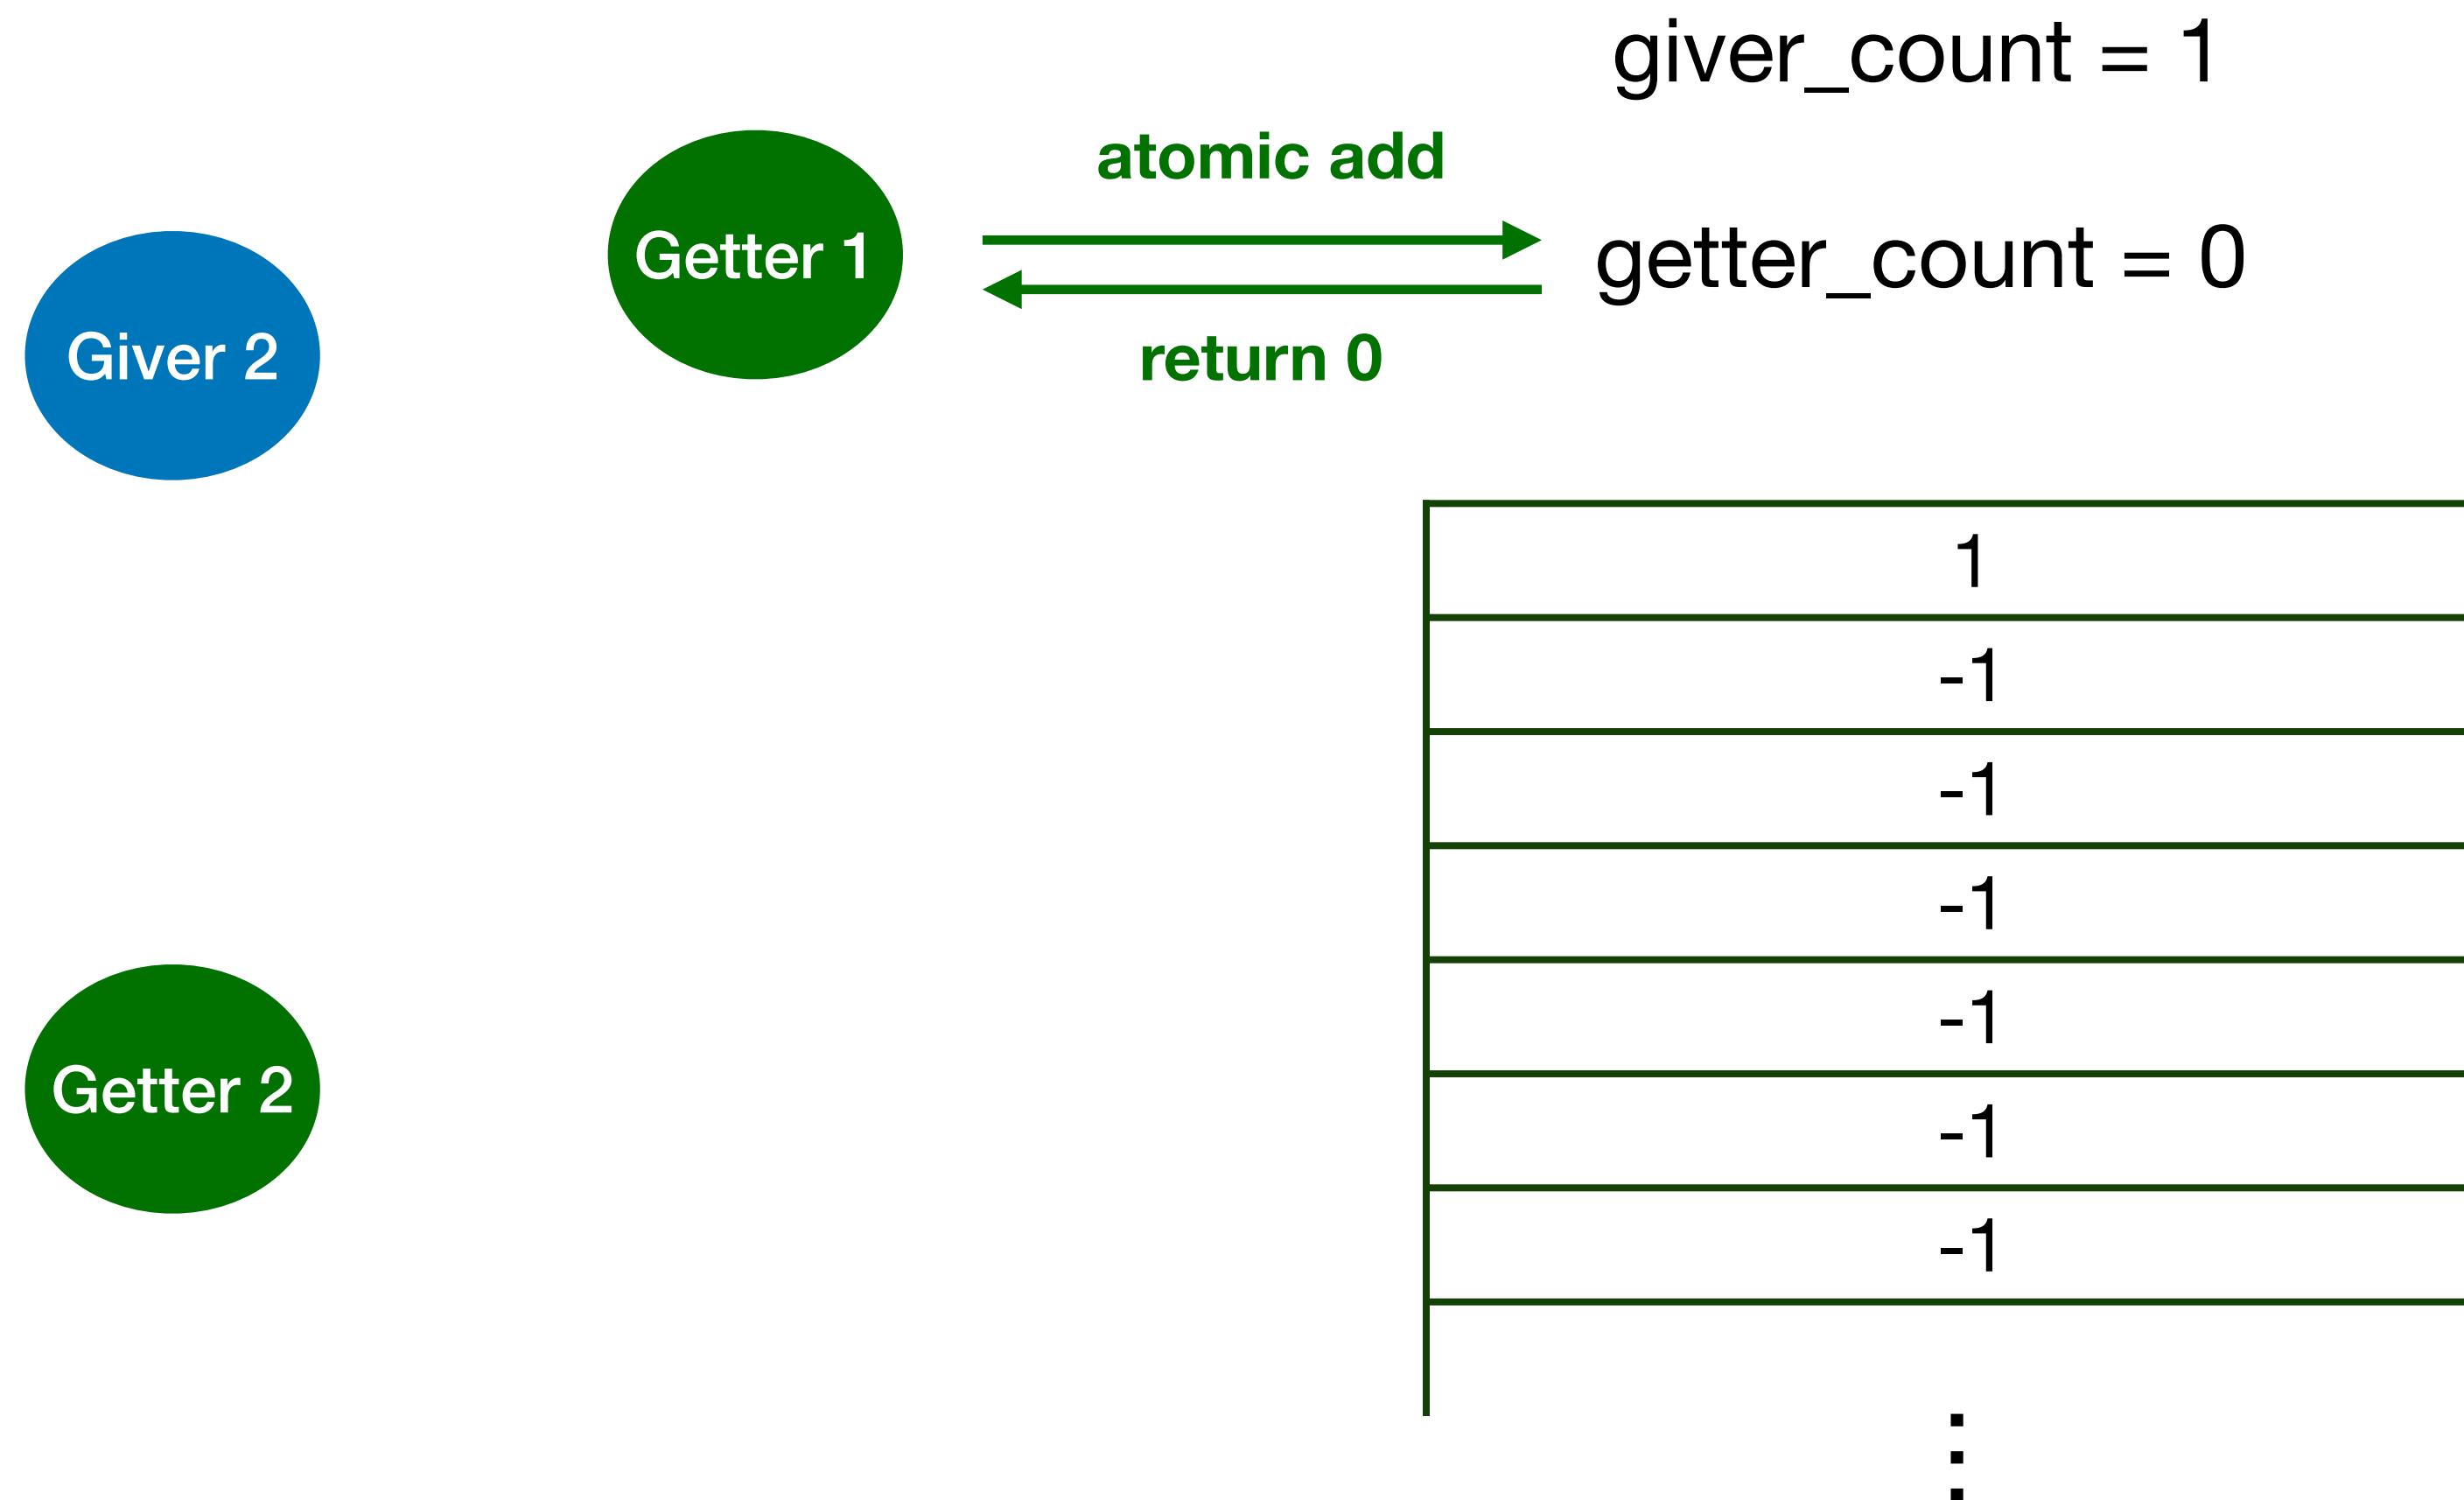
Work Stealing Among GPU Threads

- How to give or get? **Use atomic swap**



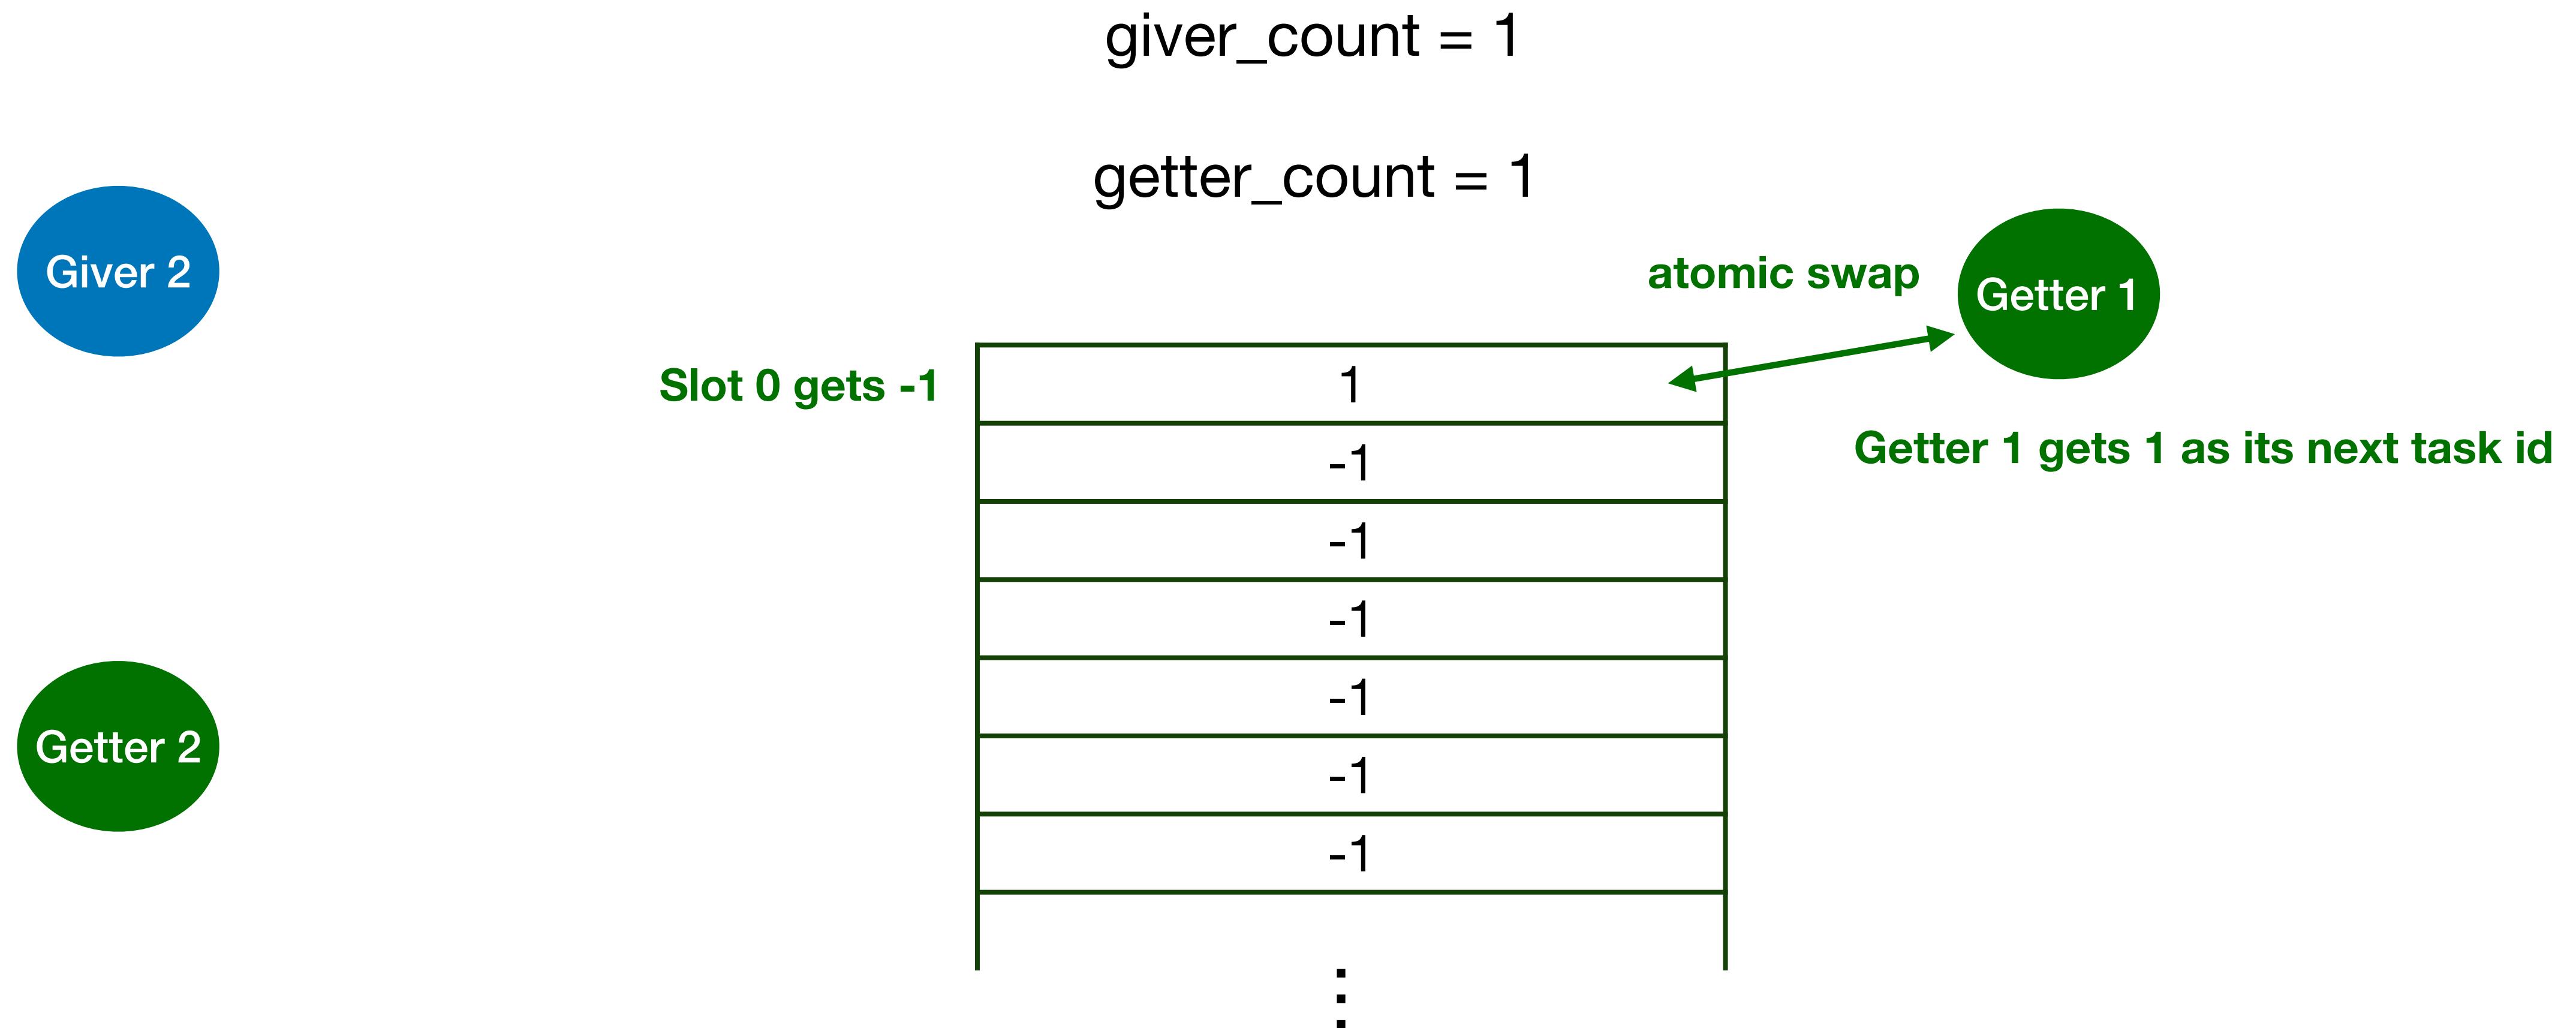
Work Stealing Among GPU Threads

- How to give or get? **Use atomic swap**



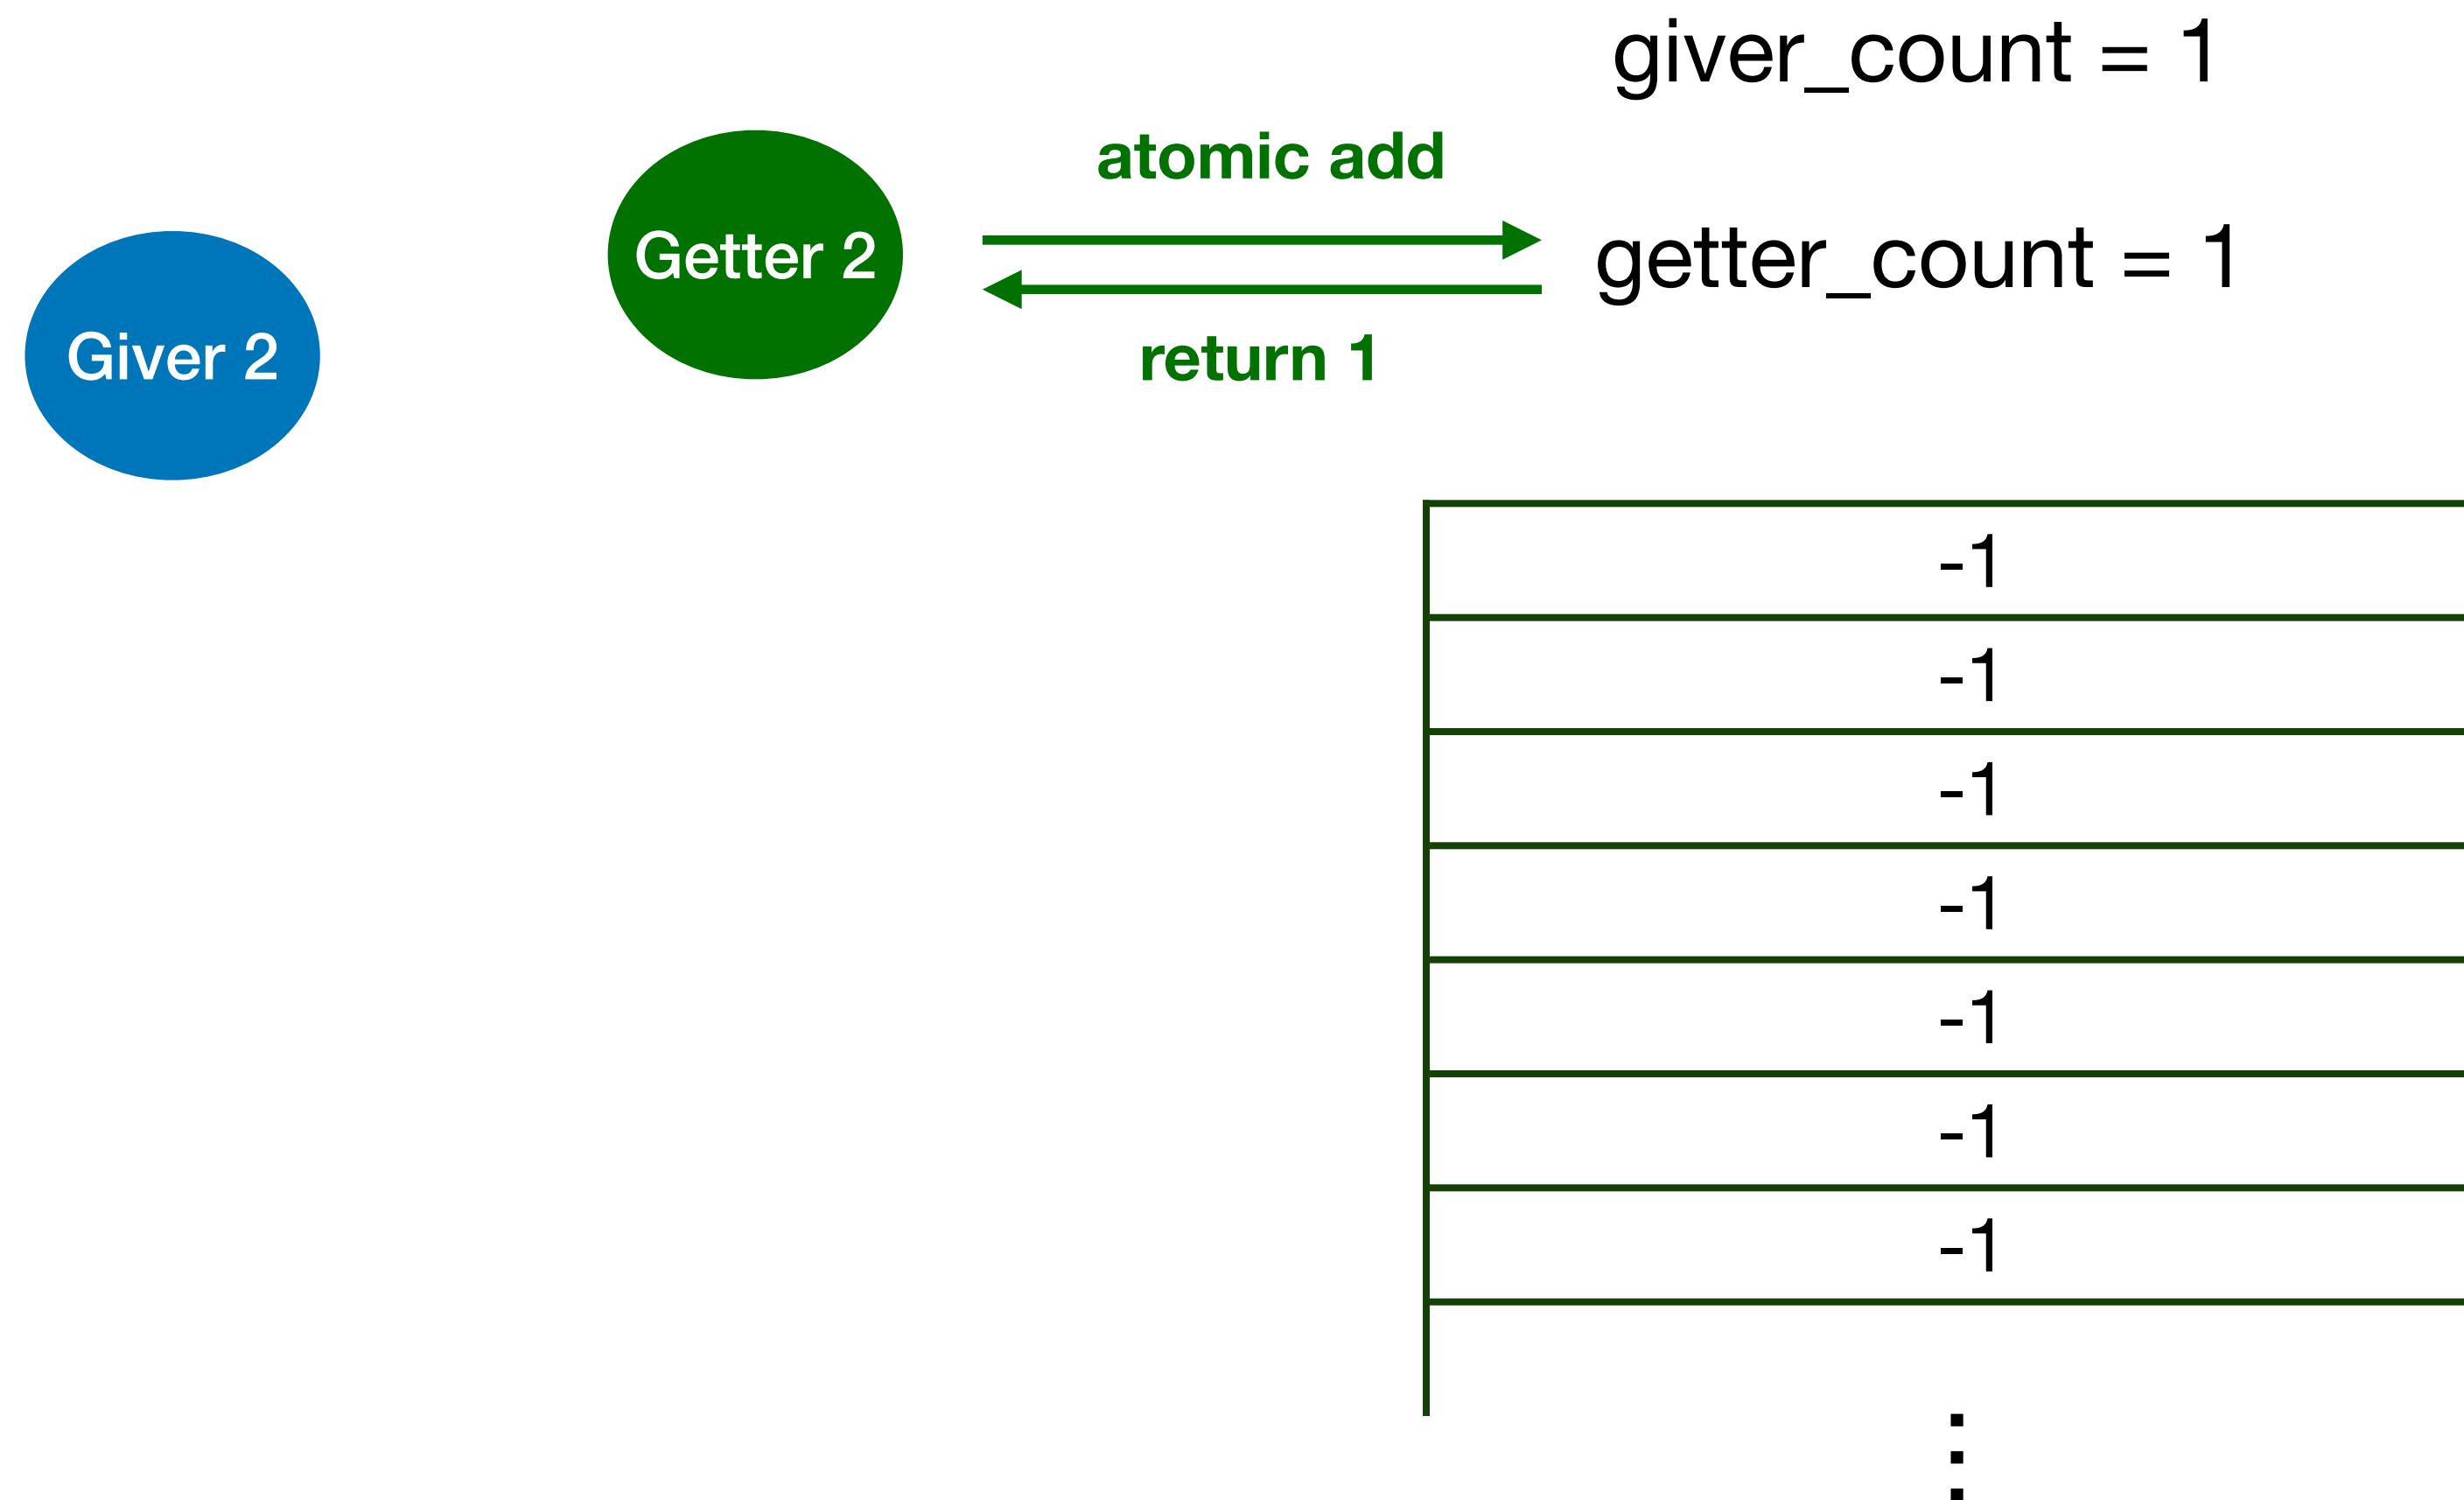
Work Stealing Among GPU Threads

- How to give or get? **Use atomic swap**



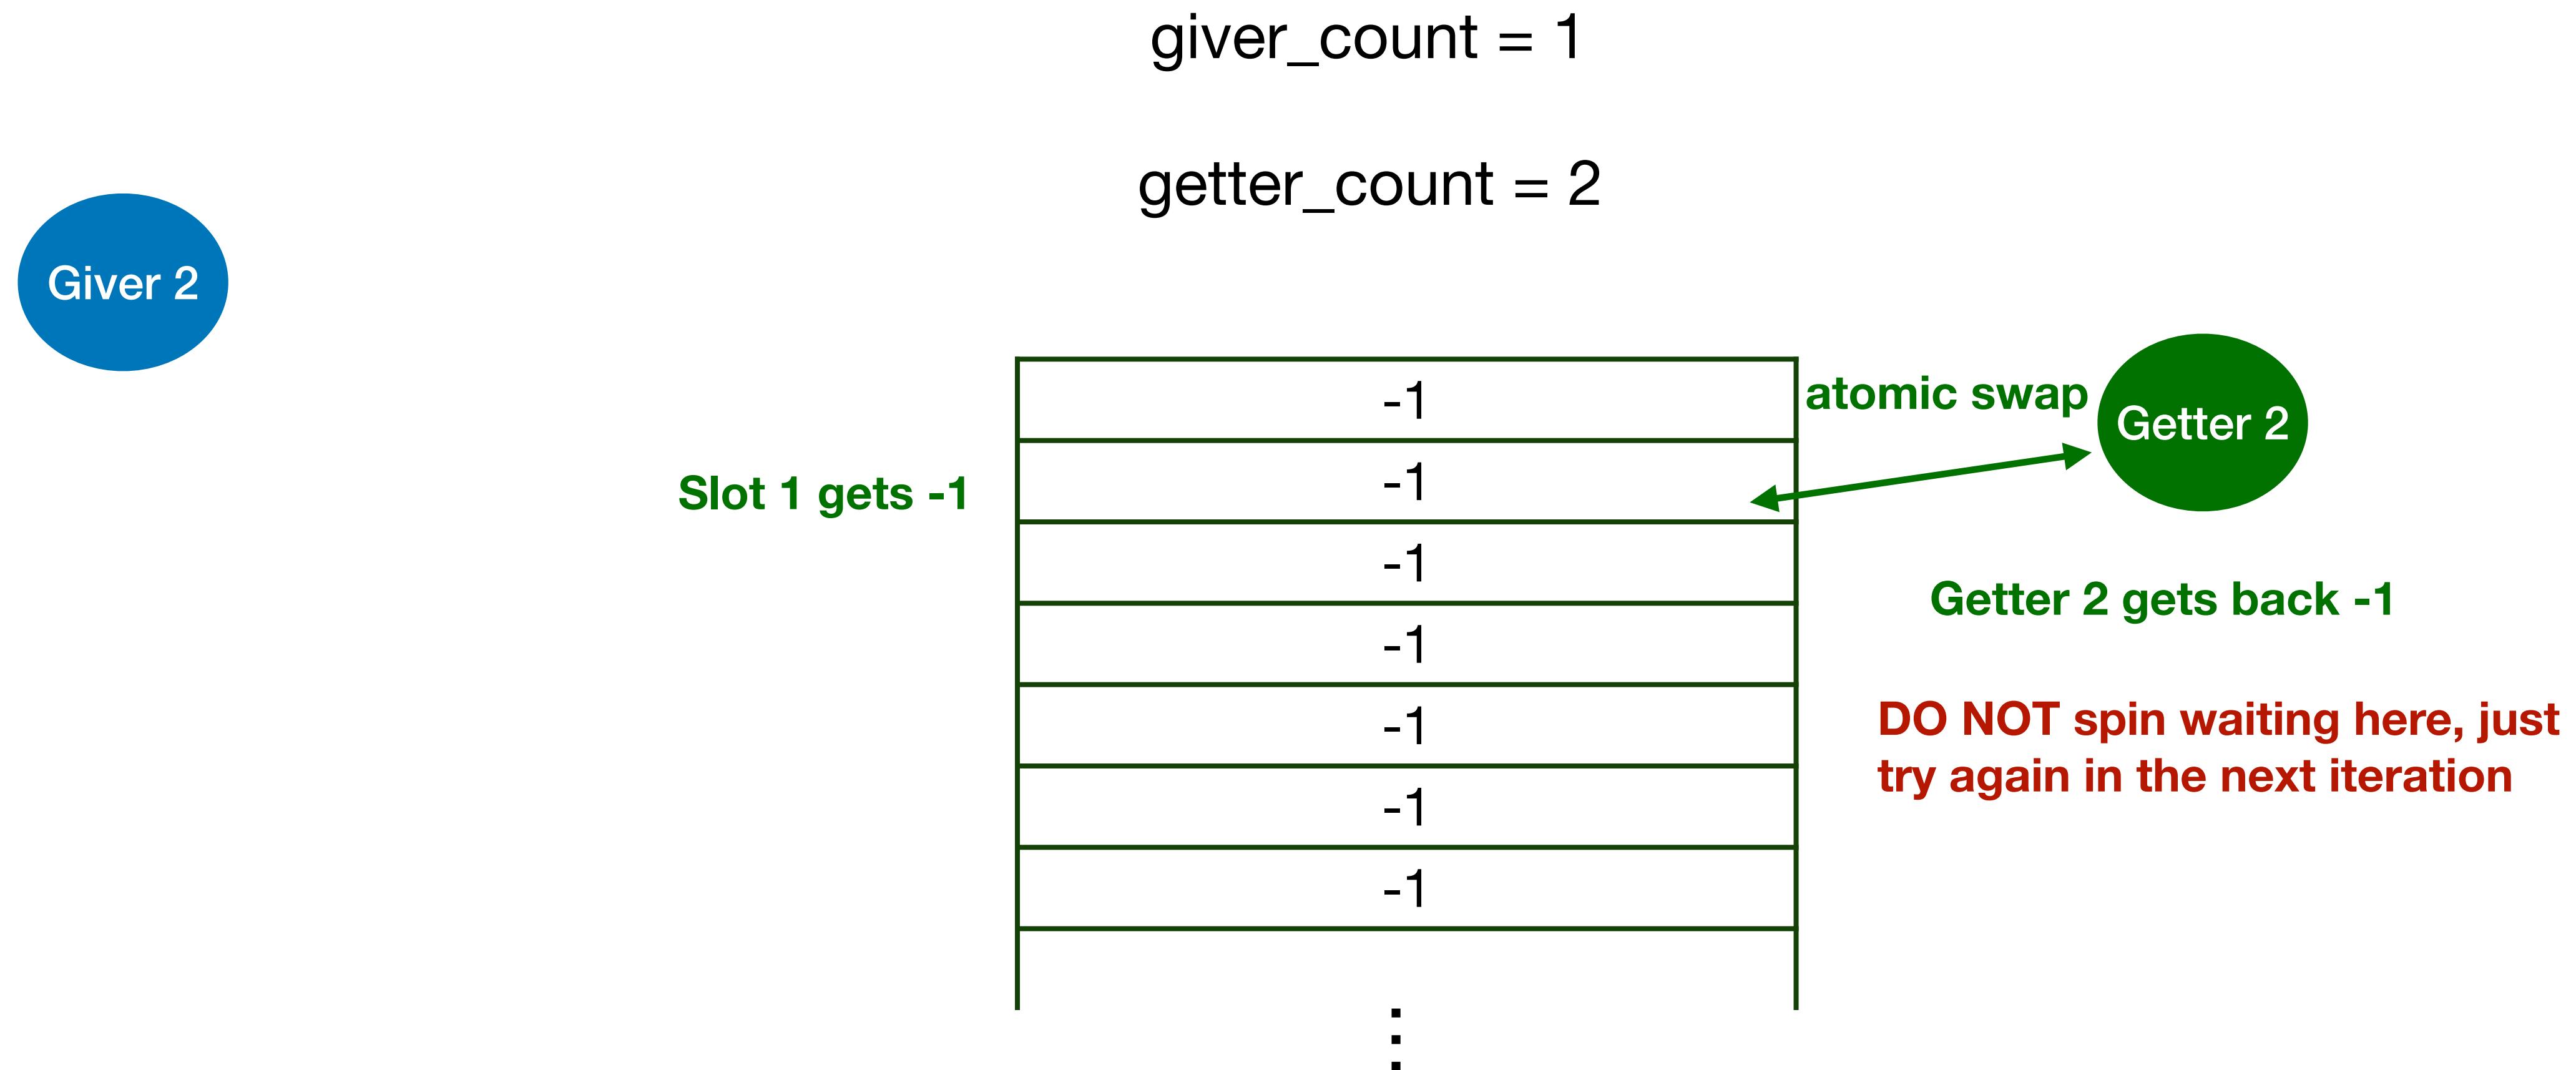
Work Stealing Among GPU Threads

- How to give or get? **Use atomic swap**



Work Stealing Among GPU Threads

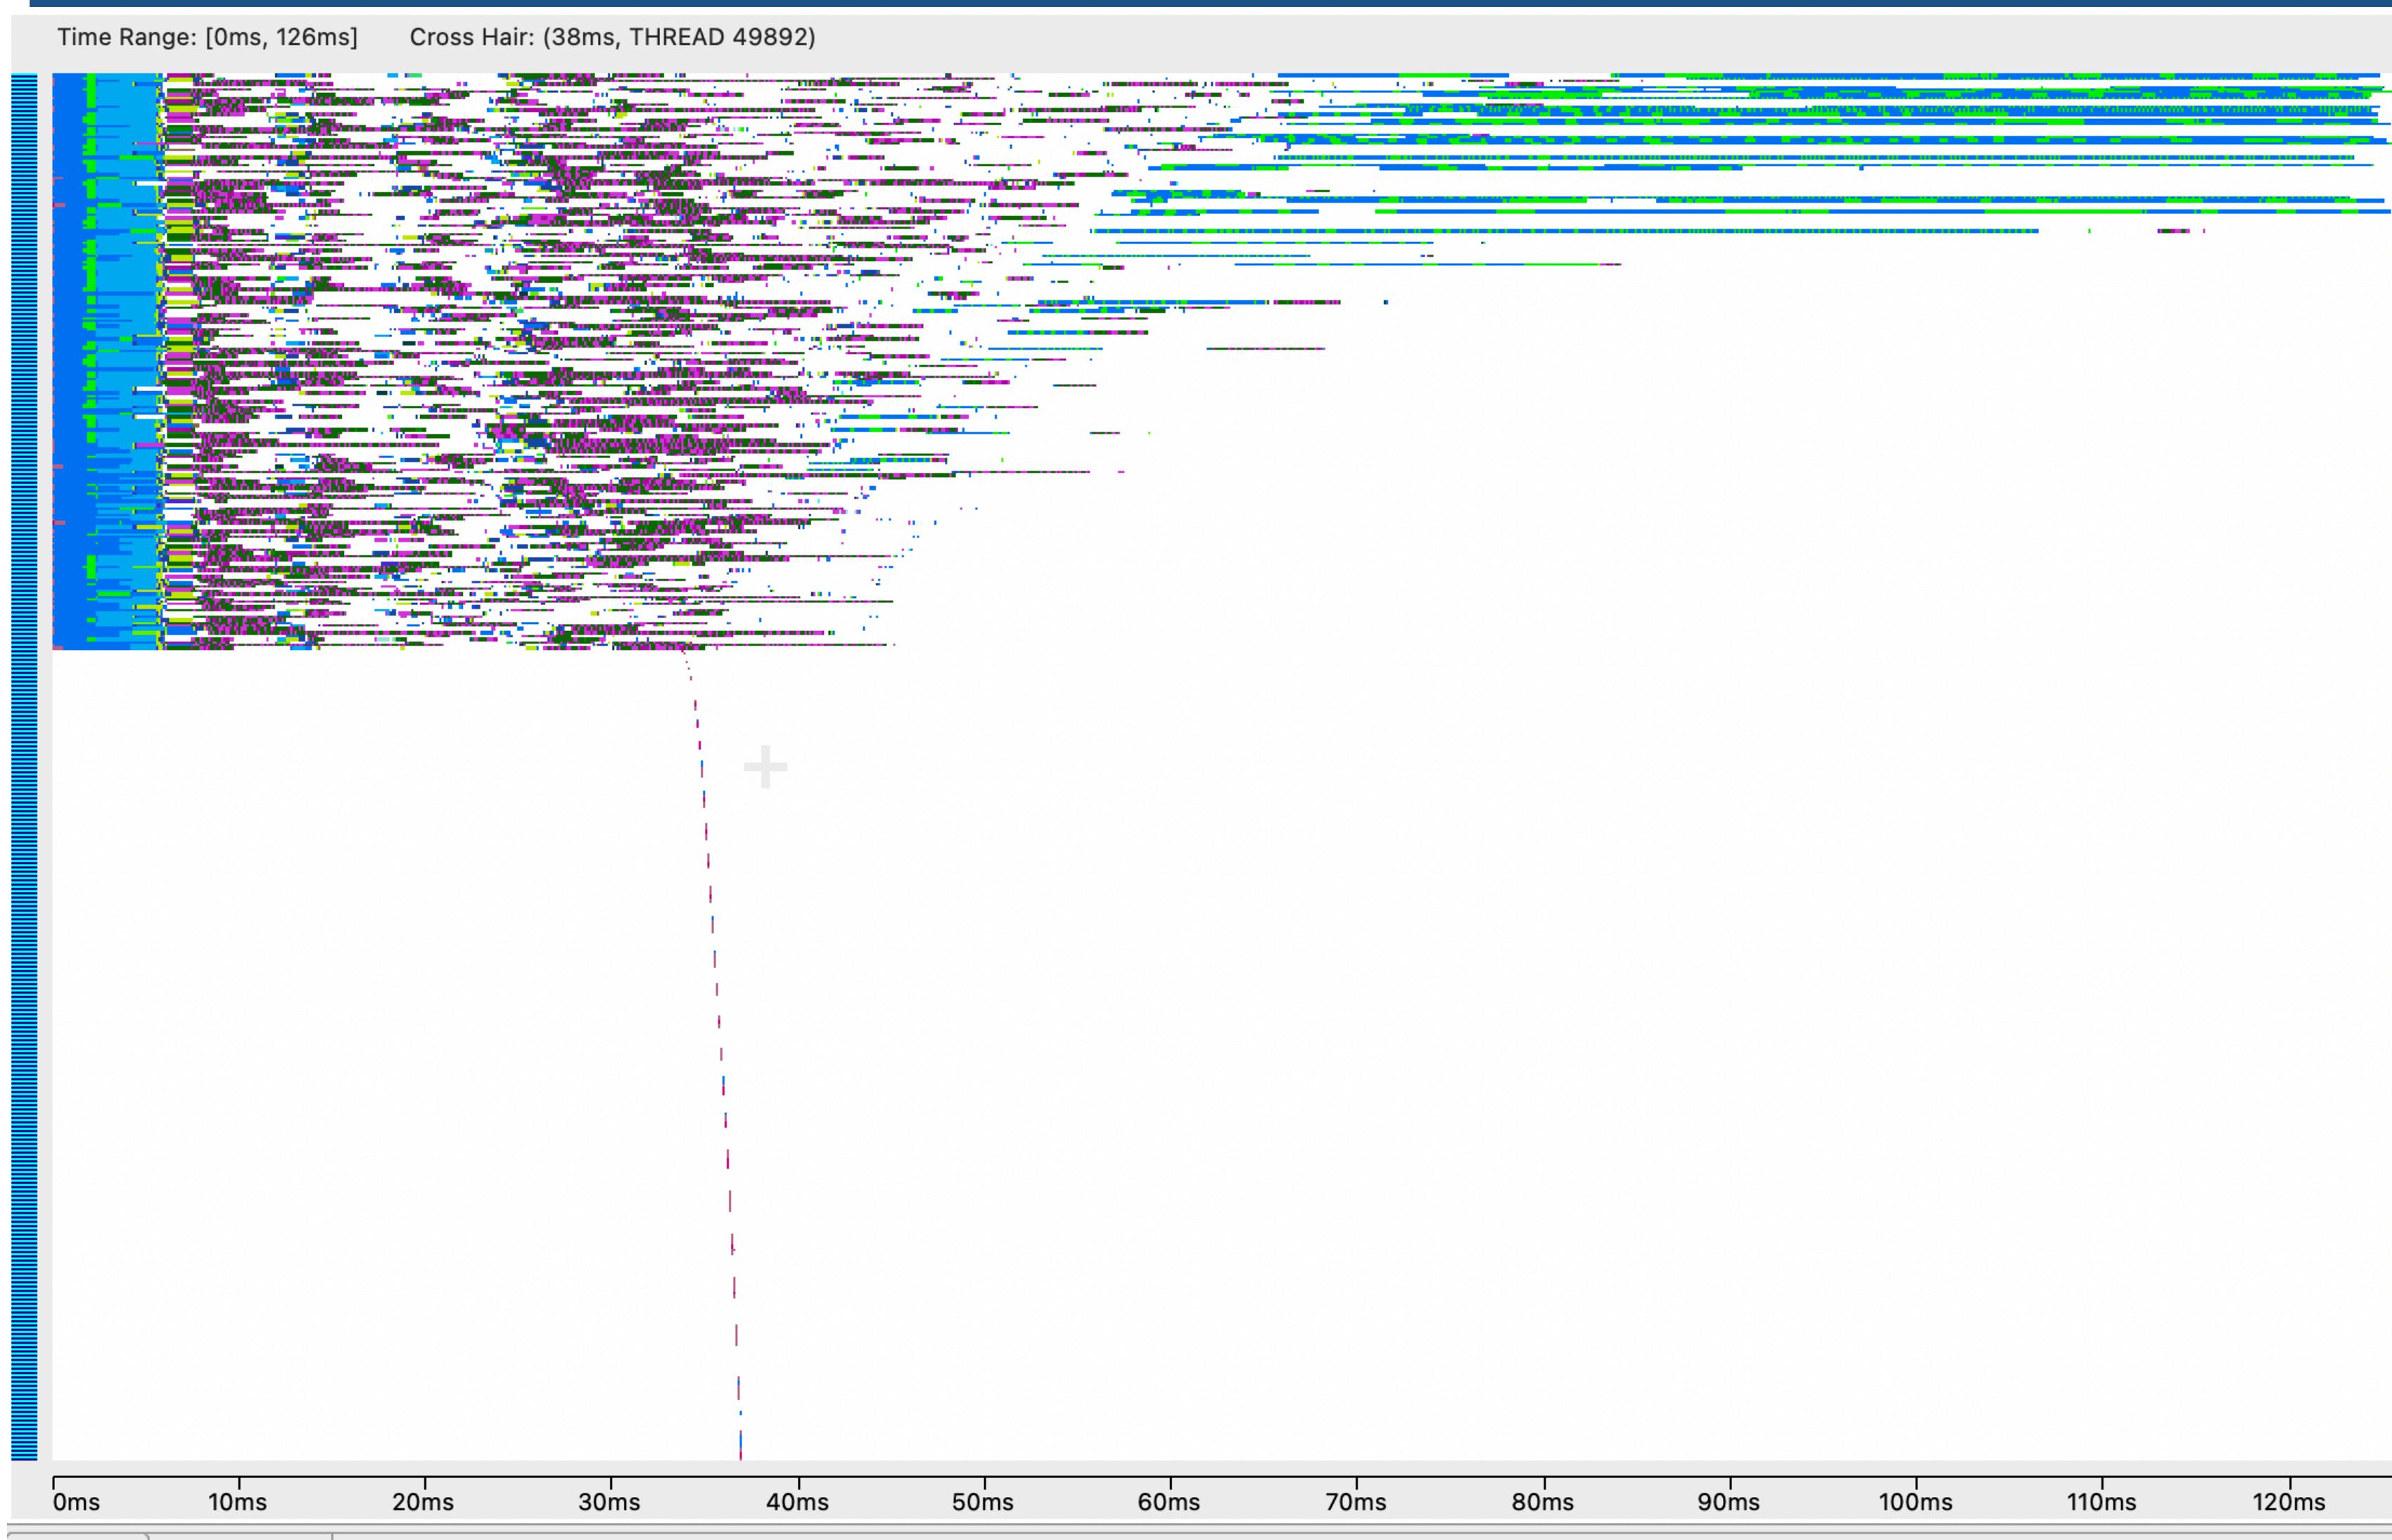
- How to give or get? **Use atomic swap**



Work Stealing Among GPU Threads

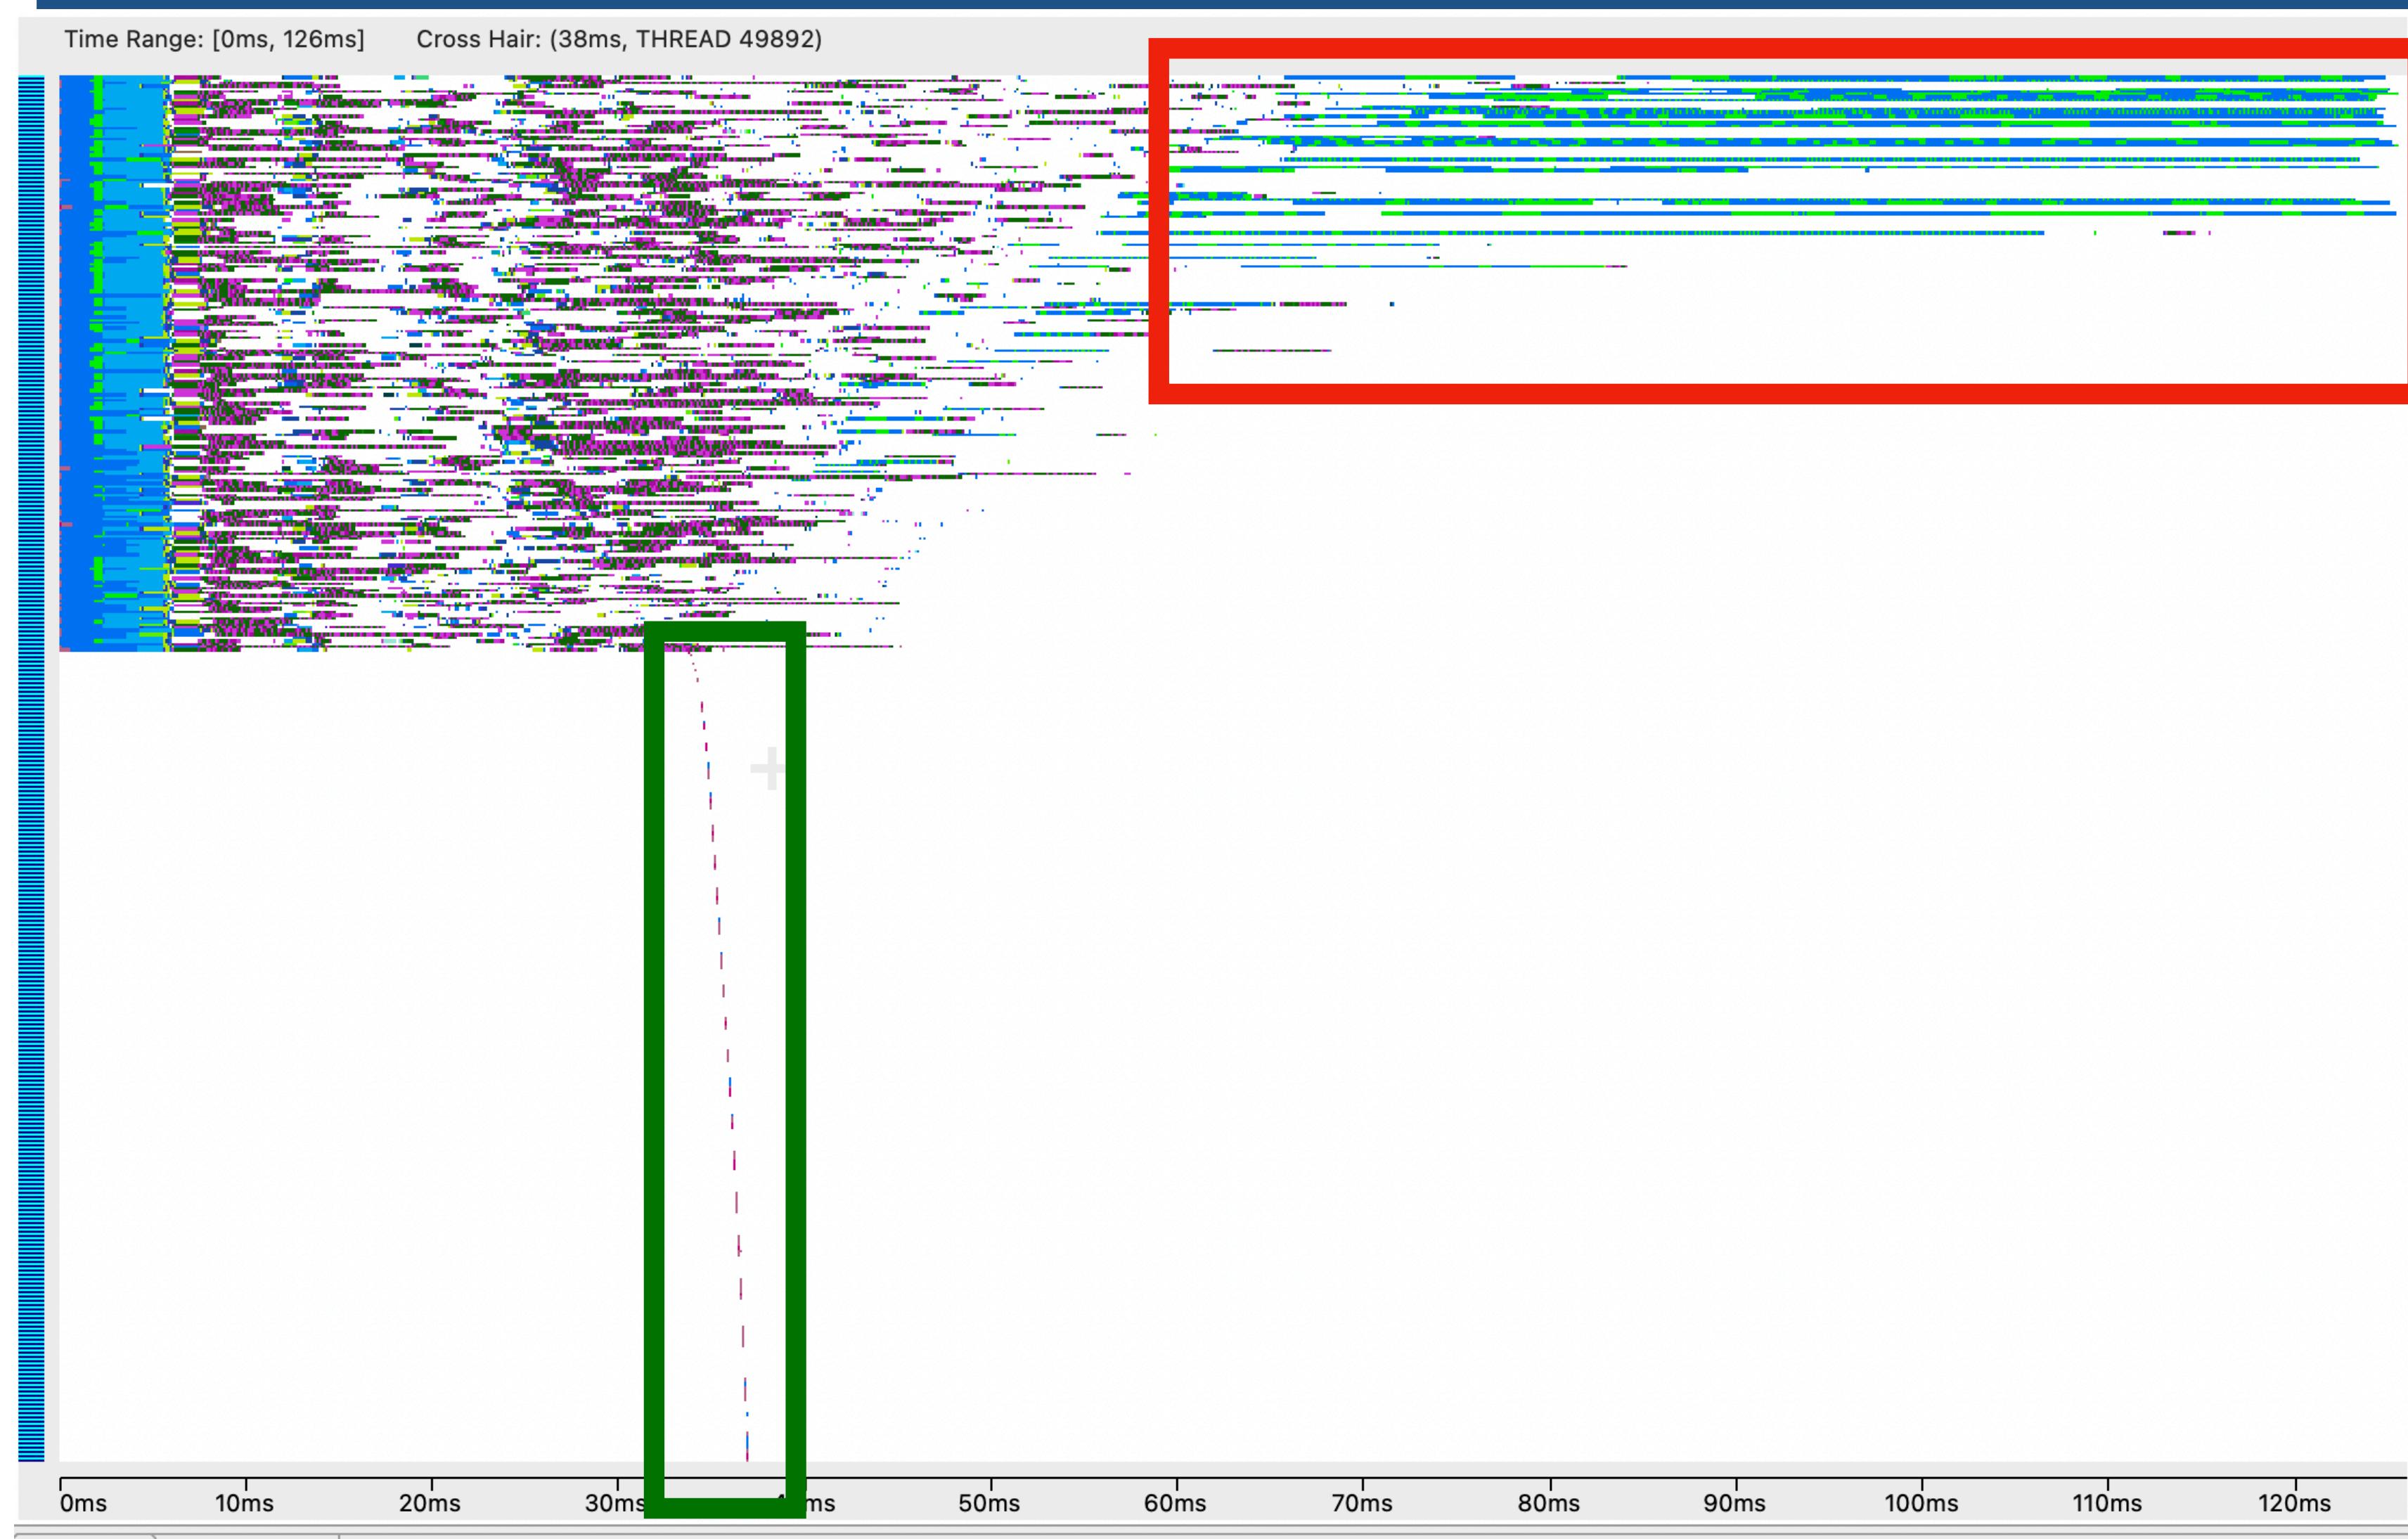
- Each thread starts with its own task, and in each iteration
 - Step 1: works on its task
 - Step 2: decides if it should
 - give another thread its task
 - get another thread's task
 - exit
 - nothing
 - Step 3: behaves according to the decision from Step 2
 - Step 4: `__syncwarp()`
- Only if it is NOT trying to get or give**
- Only try one time of get() or give() in each iteration**

Work Stealing Among GPU Threads - Quicksilver



- Successfully compacted tasks

Work Stealing Among GPU Threads - Quicksilver



- Successfully compacted tasks
- Threads started late find no work to do and just exit
- Threads keep trying to get() when only small amount of task remained

Outline

- Our trace tool
- Potential optimization opportunities of Quicksilver found with the trace tool
- Exploration of an optimization opportunity
- Work stealing algorithm among GPU threads
- **Summary**

Summary

- Binary instrumentation of GPU device functions within complex kernels reveals potential optimization opportunities
- Compacting threads with work into small amount of active blocks improves the performance
- A study of Quicksilver revealed two opportunities for optimizations
 - Imbalance between tasks
 - Imbalance between cases within each iteration
- Work stealing algorithm on GPU threads seems promising to accelerate the kernel