

# Improving the Tooling Interface Ecosystem

Working Group Outbrief

Scalable Tools Workshop 2025

# Needs

Must support use of multiple tools

Appropriate sequencing between runtime and tools.

Static constructors are a bit of a problem – everybody wants to be first

Hardware needs to be initialized/TLS needs to be initialized/runtimes

Would like a gestalt view of linking of the application

Where am I in the initialization - no way to see where we are and what is already loaded.

Challenge: Sometimes tools need to attach to a process (GDB/ptrace)

Need a way to get the current state of the process

Lots of Linux kernel difficulties, might be made easier for ROCm if GPUs were properly integrated into proc space

Many many fundamental differences vs. first-party interfaces

# Common Requirements

- Multiple tools must be able to monitor at once
- Appropriate sequencing between runtime and tools.
  - Static constructors are a bit of a problem – everybody wants to be first
  - Hardware/TLS/runtimes/etc. need to be initialized at the right time
  - Could possibly be achieved through a “gestalt” view of application initialization
    - So tools that run at odd times can check what is or isn’t loaded yet
- What about tools late-attaching to a process? (Think GDB)
  - Need an interface to get the current state of the process
  - ROCm looking into this for their interface, serious ordering challenges in the Linux kernel
    - Ben: This might be easier if GPUs were actual Linux processes
  - Interface design is vastly different compared to first-party interfaces
- Performance impact when “off” must be near zero

## Tooling for orchestration/resource manager

RM: flux/slurm

Orchestration/workflow: kubernetes, merlin

Highly multitool - occurs at multiple levels task level data vs. workflow level

Things are dynamic

We see things at task level not at the higher level

Detect conflicts

See how things fit together

# Target Audience

- We want tooling in GPU runtimes
  - HPCToolkit: overall successful collaboration with ROCm, requested Intel copy from them
- We want tooling in CPU runtimes
  - E.g. OpenMP
- We want tooling in workload orchestrators/resource managers
  - Examples: K8s, Merlin, Flux, Slurm
  - Workload performance ≠ task performance
    - Multiple levels of performance to consider, highly multitool
- We want tooling in Glibc
  - Allocator (malloc/free/...) hooks
  - Pthread operations (create/mutex/...) hooks
  - Filesystem I/O (open/close/stat/read/write/realpath/...) hooks
  - ...and others

# Runtime initialization

We have solved problems with ROCm

Can Intel copy the overall design? (current solution from Intel is not acceptable)

Pre-and-post thread creation callbacks. Pre-create callback indicates role of the forthcoming thread

Runtime initialization - tool do your stuff now

Tool issues upcall saying I'm initializing

API that enables one to connect GPU operation invocation context to performance information for a GPU event

PC Sampling with correlation id enables a tool to correlate back to invocation context

Correlation ids for relating asynchronous activities and the lifetime management

Make a base standard for tooling and runtime.

Multiple tools need to be able to attach

# Action Items

- Collect lessons from ROCm interaction, for GPU tooling interfaces
  - Role markers for runtime-created threads
  - Tool boot trigger on runtime initialization
  - Runtime boot trigger on tool initialization
  - Correlation between GPU performance events and CPU calling context
    - Especially for instruction-level performance (PC sampling)
  - Correlation between related asynchronous events
  - Lifetime management for correlation ids
- Start standardizing the tooling interfaces we want
  - No one likes making standards... but only way to combat API divergence
  - Start with the runtimes we care about and expand from there

# The Pitch

- Symbol wrapping is not a sustainable solution
  - Optimized libraries (e.g. CUDA) don't support wrapped symbols, internal calls avoid PLT
  - "Shadowing" the upstream APIs takes persistent development effort from tool developers
  - Wrong abstraction level, missing mapping information to make sense of performance
- Tooling interfaces in libraries/applications lead to more robust tools
  - Maintained by the actual developers (who understand their APIs!) instead of a third party
  - Less effort for tools to adopt new tech, more man-time can be spent on useful tasks (& interesting problems)
- We want widespread high-quality tooling interfaces
  - What can we do to make this happen?

# Resources

- OpenMP / OMPT: <https://www.openmp.org/spec-html/5.1/openmpch4.html>
- ROCProfiler-SDK:  
<https://github.com/ROCM/rocprofiler-sdk/blob/30e239d1b7bb74df8d255f01956e4631780723d8/source/include/rocprofiler-sdk/registration.h#L222>
- PerfStubs:  
<https://github.com/UO-OACISS/perfstubs>
- “The Case for a Common Instrumentation Interface for HPC Codes”  
<https://ieeexplore.ieee.org/document/8955675>

# Random

We need a generic demangler that avoids the baggage of libiberty

Libc++ doesn't do Rust, doesn't have all the features

LLVM libdemangle is missing features (also LLVM)

# Off-topic Discussion: Symbol Demangling

- Background: Binutils (libbfd, libiberty, etc.) is missing shared libraries
  - Binutils doesn't maintain a stable ABI, so major distros don't ship shared libraries
  - Some past difficulties getting -fPIC static libraries, but fixed in recent Fedora/RHEL
- Fedora/RHEL: static linking is **STRONGLY** discouraged
- Most usage of libbfd can be replaced by Elfutils
  - HPCToolkit made the jump successfully a few years ago
  - Spindle currently using a mix, but could consider moving to pure Elfutils
- Challenge: libiberty's symbol demangler has no direct replacement
  - C++ ABI: Only for C++ symbols (not e.g. Rust), missing options, not usable from pure C
  - LLVM libdemangle: Missing options, not usable from pure C, requires building LLVM
  - Need a new “libdemangle” to replace libiberty

# Attendees

- Jonathon Anderson
- John Mellor-Crummey
- Ben Woodard
- Matt LeGendre
- Jonathan Madeson
- Add your name here