

# Efficient Instrumentation and Tracepoint Insertion for GPU Compute Kernels

Low-overhead Trace Collection on GPU

---

Sébastien Darche <[sebastien.darche at polymtl.ca](mailto:sebastien.darche@polymtl.ca)>

July 8th, 2025

Dorsal - Polytechnique Montréal

# Current projects at DORSAL lab

- The Distributed Open Reliable Systems Analysis Lab
- Strong focus on trace collection and performance analysis
- LTTng, Trace Compass



# Tooling for HPC

- Score-P traces support through CTF conversion, ROCm runtime instrumentation
- Multiple analyses available
  - Critical path for linux kernel traces
  - Hardware performance counters through Score-P
  - Call stack among ranks, statistics
  - Flame graph
  - Communicators, bandwidth
  - Critical path for MPI
  - ...
- Scalability of Trace Compass through distributed analyses (ongoing work)
- Current work on kernel instrumentation

# GPU Tracing with hip-analyzer

- Few tools for tracing on GPUs, and often at the cost of very high performance impact (at minima 10× and up to 120×) [1] [2]
- GPU Tracing is unwieldy : clumsy memory management, massive parallelism (concurrency control, high throughput)
- Separate buffer allocation and event collection using two kernel runs
- "Online" tracing methods
- LLVM IR (static) instrumentation

## What's new ?

- Introduced tracing methods that do not require two executions, but has its own set of challenges
- Requires specific tuning for the hardware
  - Memory locality
  - Allocation granularity
  - Implementation choices
- Last project focuses on reducing instrumentation in the kernel

# Results

- Instrumentation tested on the HeCBench [3] benchmark. Overhead is reported as the slowdown factor between the traced kernel execution time and the original, uninstrumented kernel.

	mean	median
hip-trace	2.07×	1.50×
4 × padded hip-trace	2.18×	1.58×
hip-global-mem	3.73×	1.96×
hip-cu-mem	2.47×	1.60×
hip-chunk-allocator	1.79×	1.33×
hip-cu-chunk-allocator	1.77×	1.32×

# Reducing the number of tracepoints

- For most kernels, the number of tracepoints could be reduced
  - Reduction in trace size
  - Reduction in run time overhead
- Intuitively, if half of the threads go through an *if* statement, we can deduce the other half goes to the *else* statement
- Can be generalized to switch statements and more complex control flow (more than two outgoing edges)

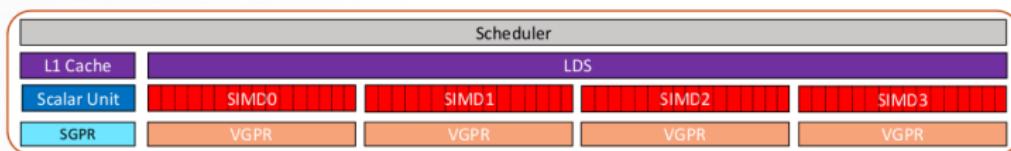


Figure 1: AMD GCN Compute unit <sup>1</sup>

<sup>1</sup> Reproduced from *AMD GPU Hardware Basics*, 2019 Frontier Application Readiness Kick-off Workshop

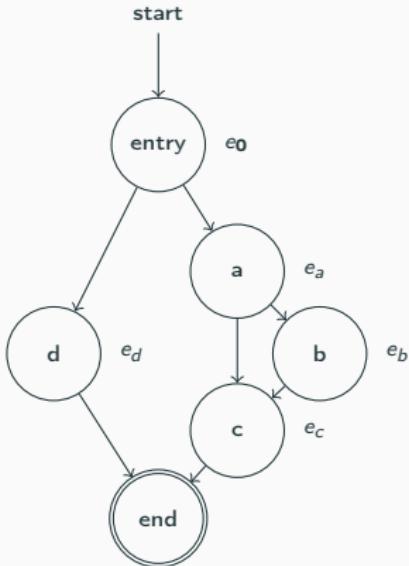
## Static analysis

- In an acyclic CFG (simple case), the control flow can be completely computed by instrumenting  $n - 1$  outgoing edges
- Vertices in the CFG are processed using a variant of Kahn's algorithm

$$\bigvee_{e_i \in \text{incoming}} T(e_i) = \bigvee_{e_i \in \text{outgoing}} T(e_i) \quad (1)$$

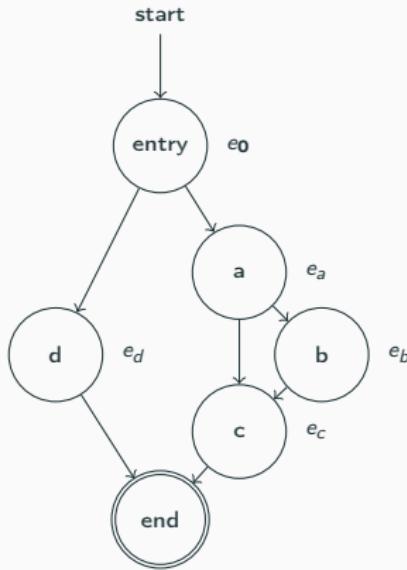
- The algorithm does not terminate for CFGs containing a cycle
- We identify back edges using a depth-first search (DFS), and run the algorithm on the CFG stripped of its cycles. Back edges must be instrumented.

## Resulting trace



```
--global--
void simple_kernel() {
    entry();
    if(c0()) {
        a();
        if(c1()) {
            b();
        }
        c();
    } else {
        d();
    }
    end();
    return;
}
```

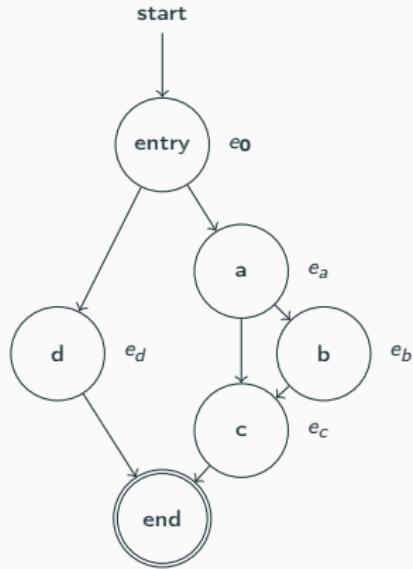
# Static analysis



**Figure 2:** Thread-centric CFG Example

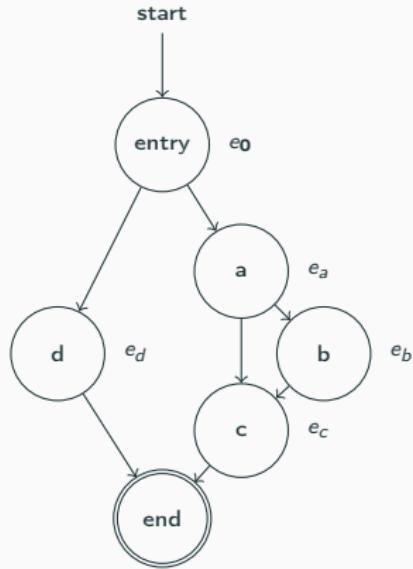
- $e_a = \overline{e_d} \cdot e_0$
- $e_c = e_a + \overline{e_b} \cdot e_a = e_a$
- $e_{end} = e_d + e_c = e_d + e_a = e_d + \overline{e_d} \cdot e_0 = e_0$

# Resulting trace



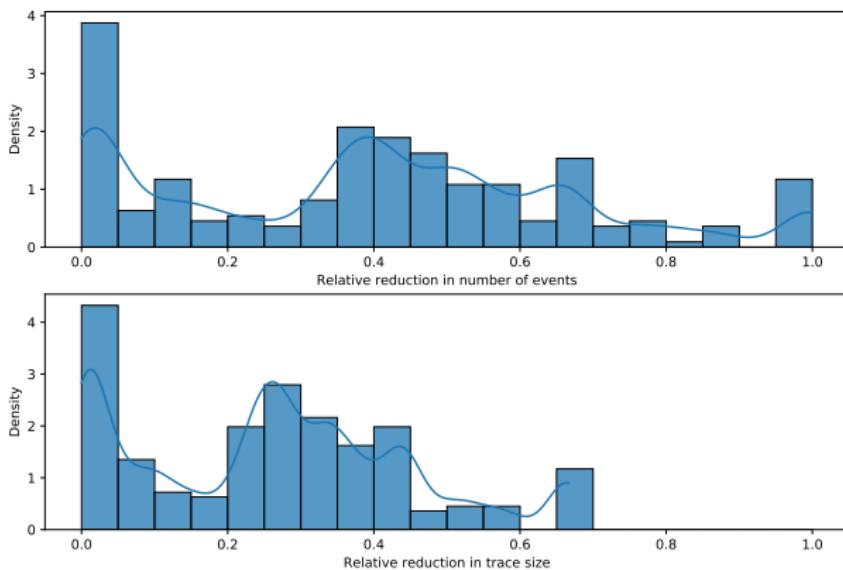
Timestamp	Basic block	Execution mask
$t_0$	entry	$1111_2$
$t_1$	a	$0111_2$
$t_2$	b	$0010_2$
$t_3$	c	$0111_2$
$t_4$	d	$1000_2$
$t_4$	end	$1111_2$

# Resulting trace



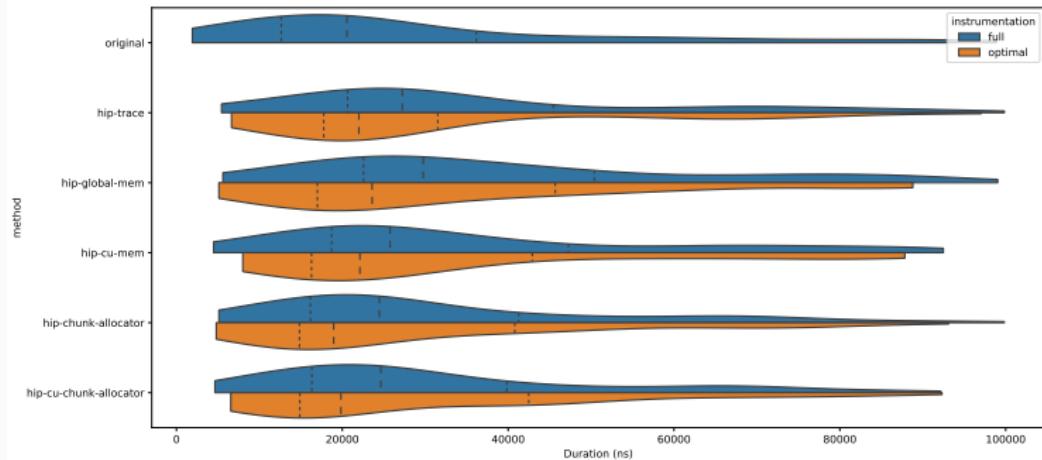
Timestamp	Basic block	Execution mask
$t_0$	entry	$1111_2$
$t_1$	entry → a	$0111_2$
$t_2$	a → b	$0010_2$
$t_3$	(c)	$(0111_2)$
$t_4$	(d)	$(1000_2)$
$t_4$	(end)	$(1111_2)$

# Trace size reduction



**Figure 3:** Relative reduction in number of events and total trace size

# Run time overhead



**Figure 4:** Distribution of kernel run time as a function of collection method and instrumentation

## Trade-offs in GPU Tracing – Instrumentation

- Instrumentation methods are intrusive and *will* modify how the kernel runs. Tradeoff between :
  - Increased register pressure (may affect occupancy)
  - Reusing registers (*scavenging*) will probably mean spills
- Tracepoints will incur a runtime overhead

## Trade-offs in GPU Tracing – Memory management

- Trace management is a major concern
- Uncertain trace size – may exceed memory
- Synchronization inside kernel bounds is not defined by the memory model
- "Smarter" trace management methods are more costly (*cf* instrumentation)

## Trade-offs in GPU Tracing – Trace analysis

- What data are we presenting to the end user?
- Thread-centric (programming language) vs. Vector representation (ISA)
- Large (!) trace files

# Tradeoffs in GPU Tracing – Future works

- Better compiler support
  - Scalar / vector registers specifications
  - Scalar / vector instructions
  - Intrinsics
  - Backend plug-ins?
- Better hardware support
  - CU-wide registers and memory access
  - Host / Device interaction
- Finer memory model

# Conclusion and future work

- PhD project is nearing its end
- Explored instrumentation methods for tracing compute kernels
- Studied the performance impact of data structures for online tracing
- Improved baseline results by reducing the number of tracepoints
- Interest for the project from partners
- Available freely on Github, feedback and/or use cases are more than welcome

 [dorsal-lab/hip-analyzer](#)

Compiler plugin for performance analysis of HIP applications

 C++  2  1

 [dorsal-lab/TraceCompassGpu](#)

Trace Compass GPU plugins

 Java

## Q&A

---

## References

---

- [1] D. Shen, S. L. Song, A. Li, and X. Liu, “**Cudaadvisor: LLVM-based runtime profiling for modern GPUs**,” in *Proceedings of the 2018 International Symposium on Code Generation and Optimization*, 2018.
- [2] Y. Arafa, A.-H. Badawy, A. ElWazir, et al., “**Hybrid, scalable, trace-driven performance modeling of gpgpus**,” in *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, 2021, pp. 1–15.
- [3] Z. Jin and J. S. Vetter, “**A benchmark suite for improving performance portability of the SYCL programming model**,” in *2023 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, IEEE, 2023, pp. 325–327.