

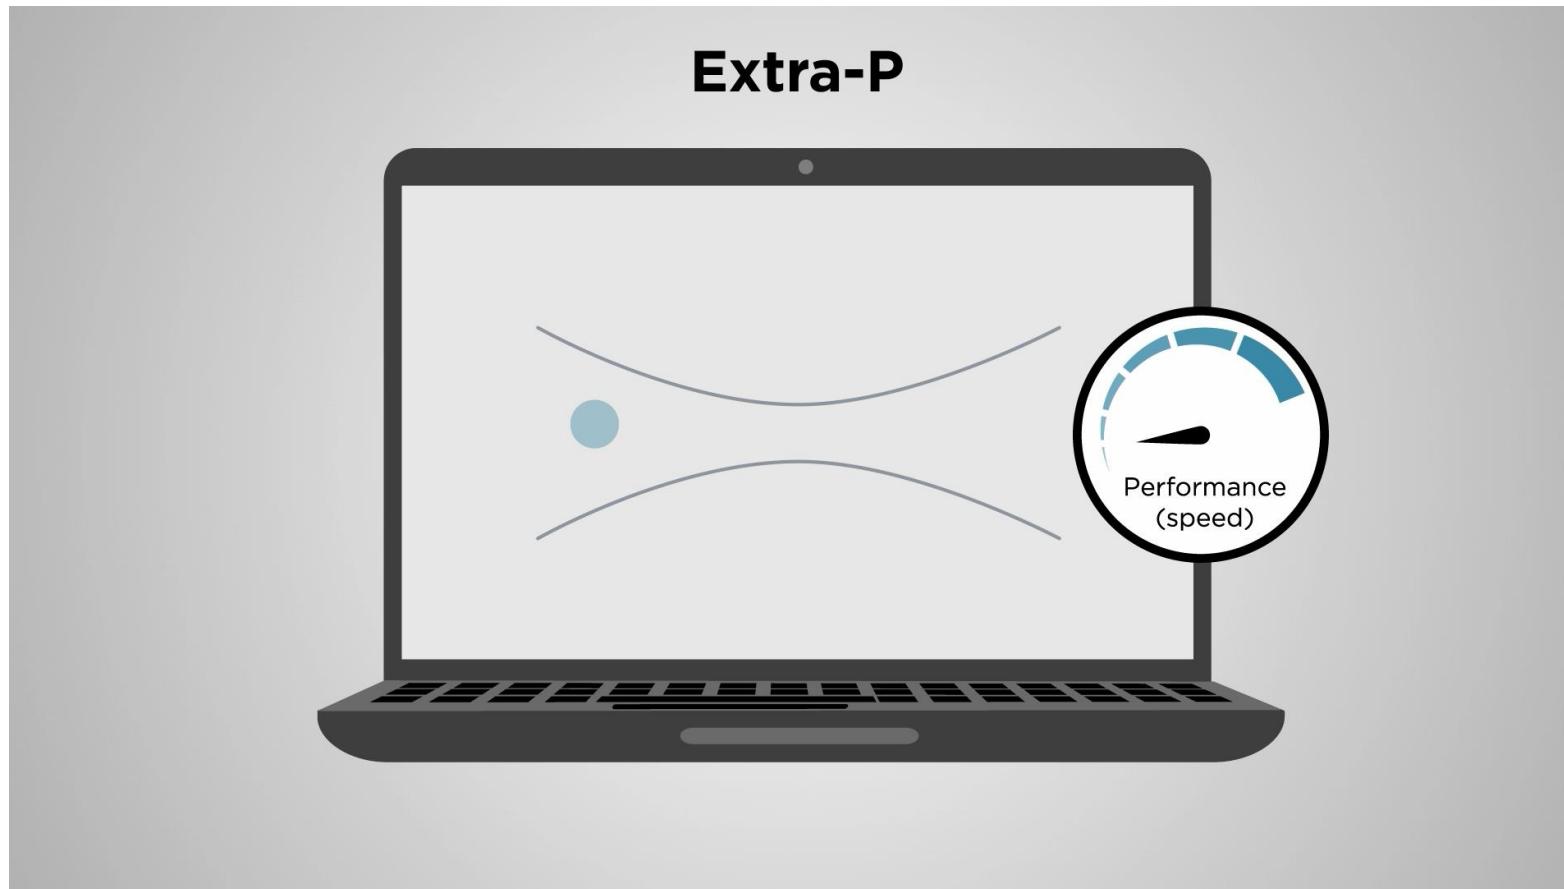


TECHNISCHE
UNIVERSITÄT
DARMSTADT

Parallel
Programming

Denoising Application Performance Models with Noise-Resilient Priors

Motivation





Watch Extra-P
overview video




<https://www.youtube.com/watch?v=Cv2YRCMWqBM>

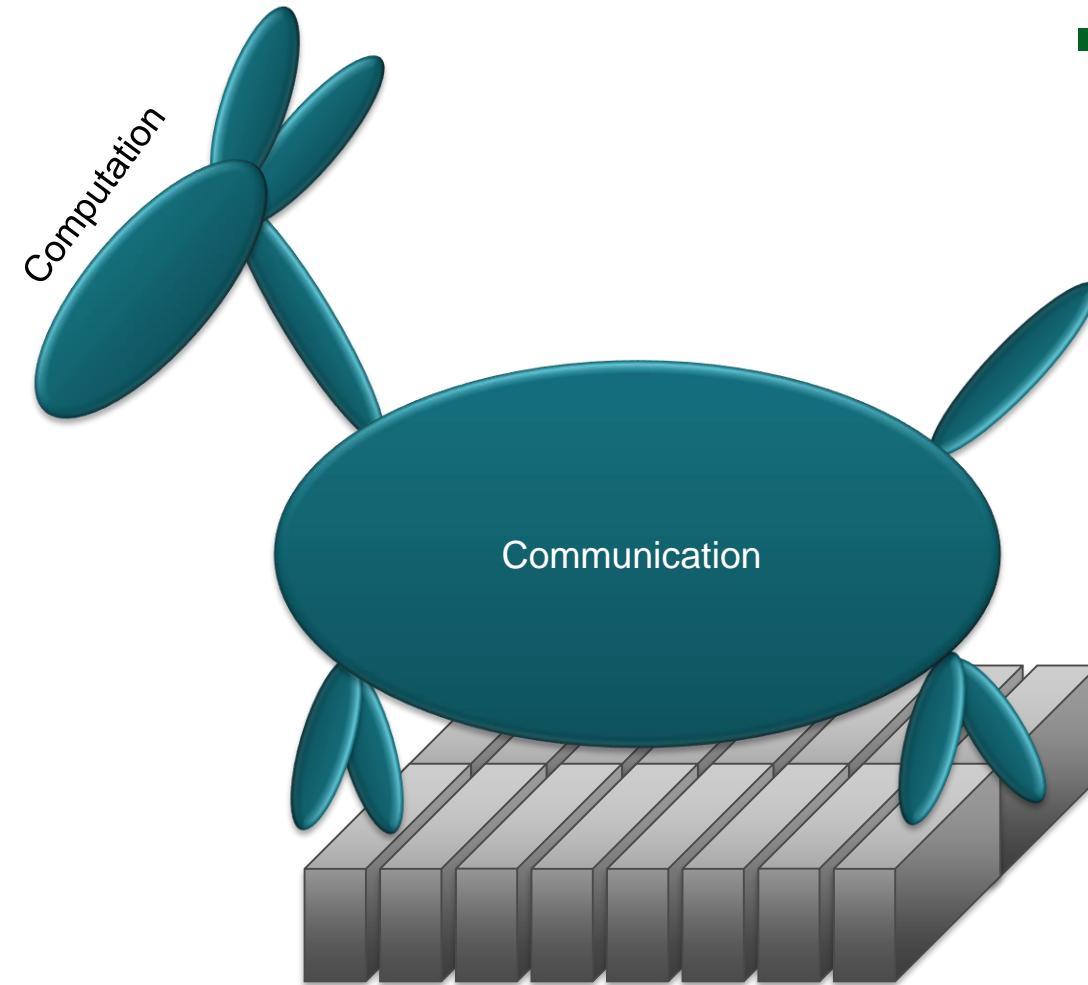

Motivation



Motivation

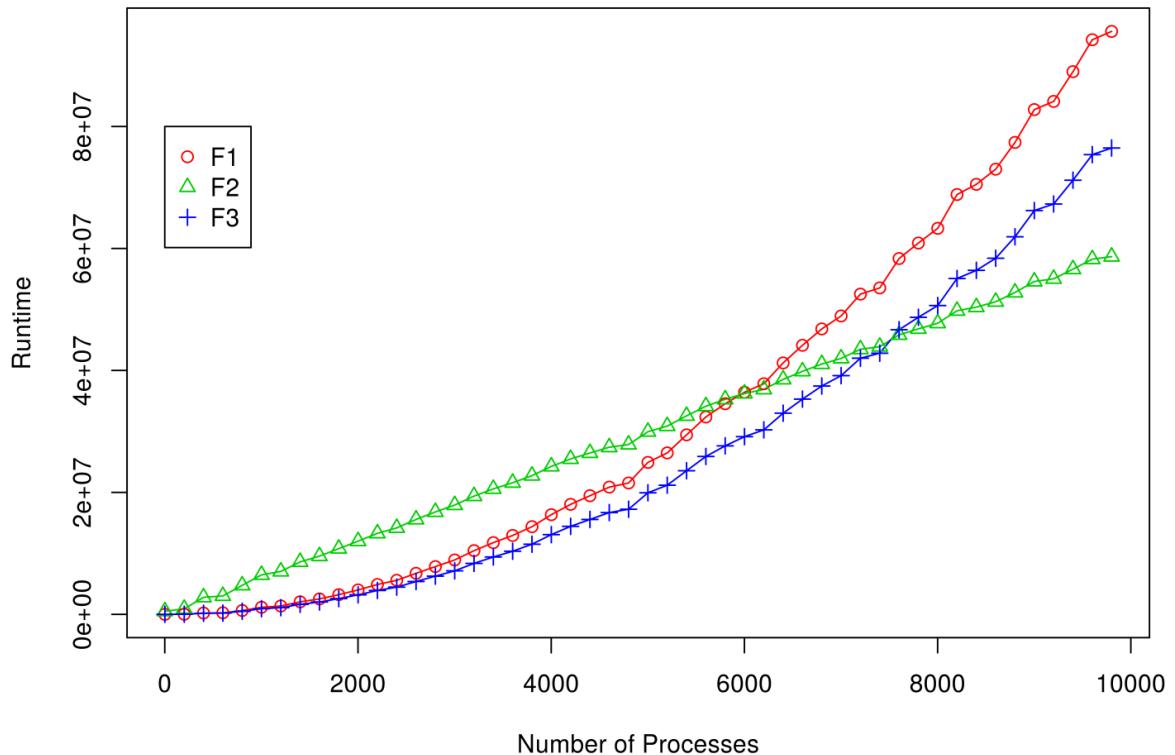


Motivation




Motivation

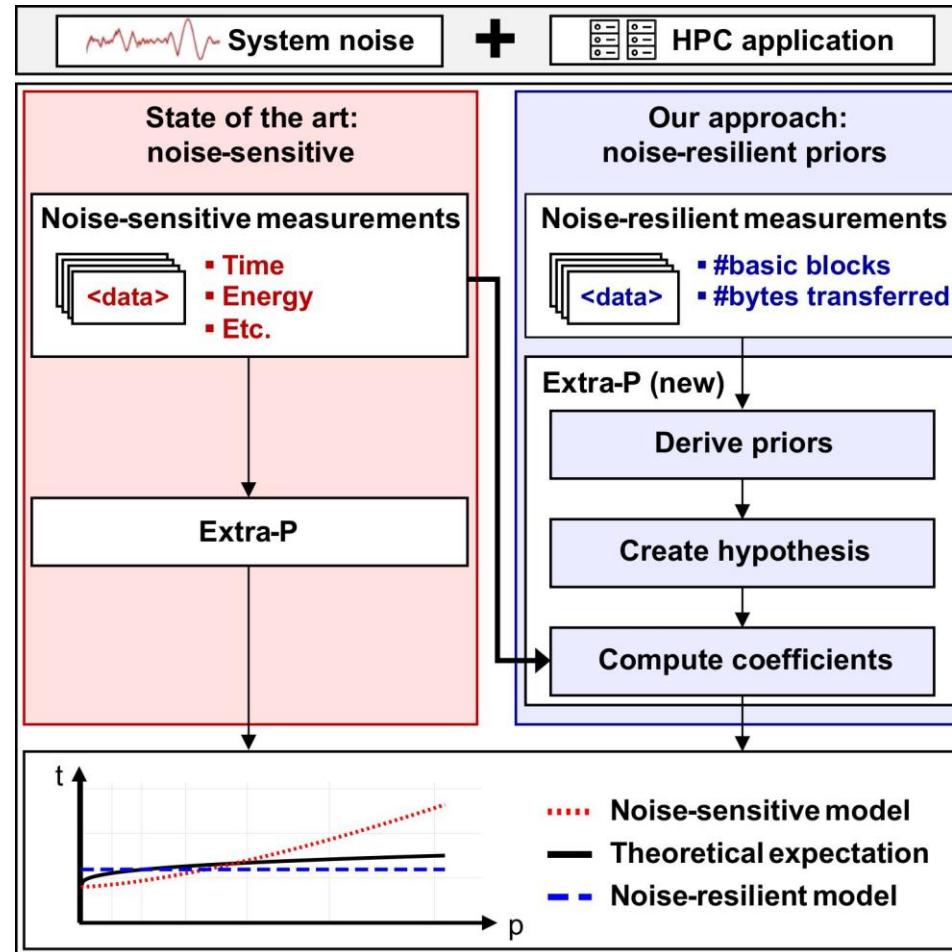
Common performance analysis chart in a paper



Production reality



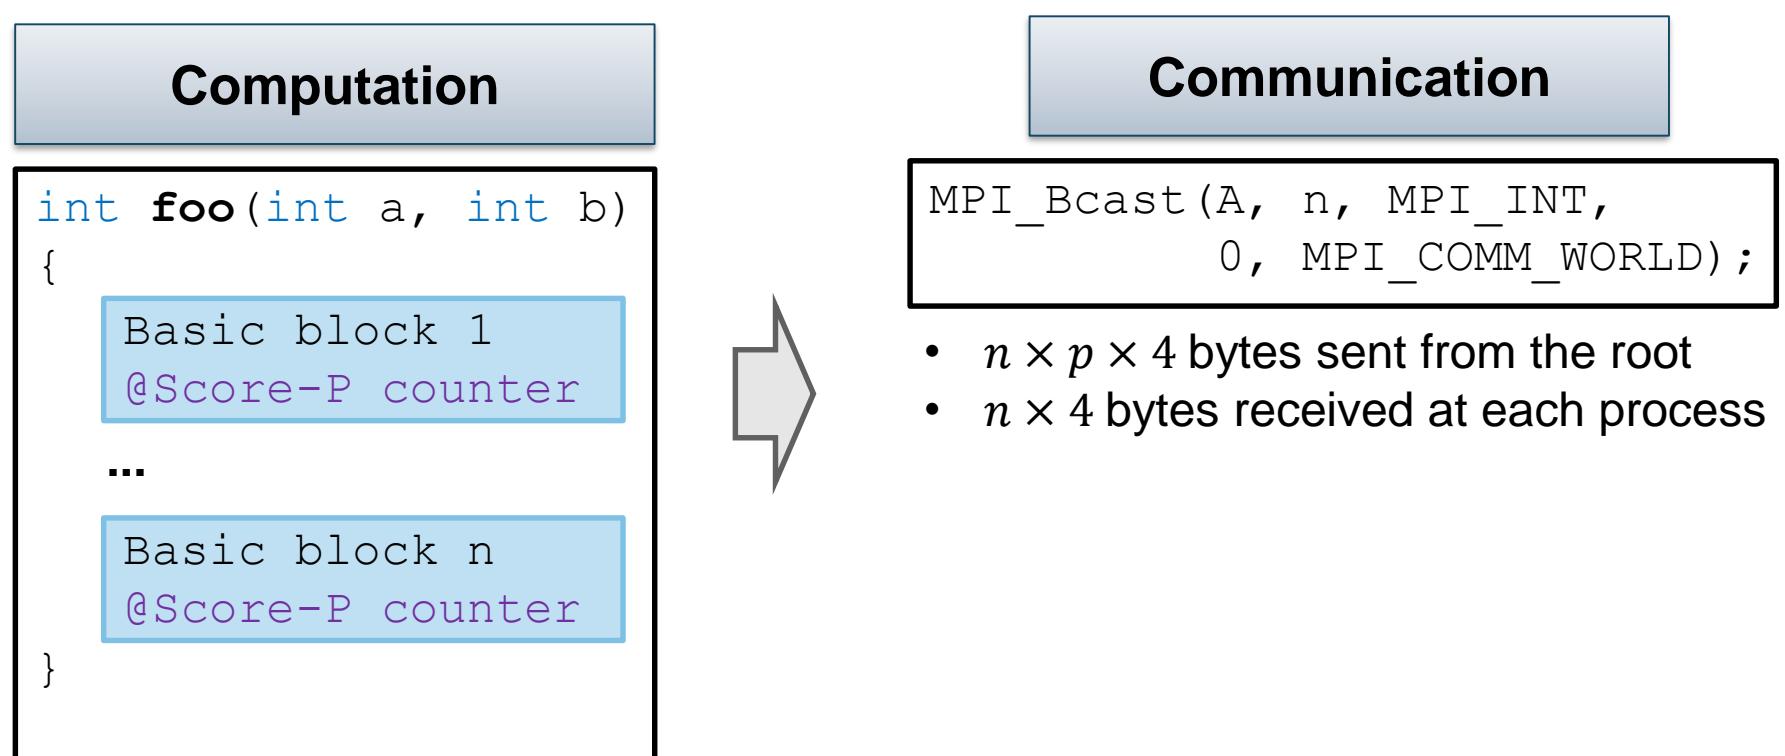
Motivation



The noise-resilient model aligns with the theoretical expectation more closely

Noise-resilient measurements

- LLVM-IR [2] plug-in into Score-P [3] framework



Multi-parameter performance modeling

- Performance Model Normal Format (PMNF) [1]


$$f(x_1, \dots, x_m) = \sum_{k=1}^n c_k \prod_{l=1}^m x_l^{i_{kl}} \cdot \log_2^{j_{kl}}(x_l)$$

$$\begin{aligned} m, n &\in \mathbb{N} \\ i_k &\in I \\ j_k &\in J \\ I, J &\subset \mathbb{Q} \end{aligned}$$

Model candidates

▪ Constant	c_1
▪ Single parameter	$c_1 + c_2 \cdot x_1$
▪ Multiple parameters	\dots
▪ Additive	$c_1 + c_2 \cdot x_1 + c_3 \cdot x_2$
▪ Multiplicative	$c_1 + c_2 \cdot x_1 \cdot x_2$
▪ Complex	$c_1 + c_2 \cdot x_1 \cdot x_2 + c_3 \cdot \log x_2 \cdot x_2^3$

Creating models from priors




Creating models from priors

▪ Communication

MPI function	Expected runtime	Ref.
Send	$f(x, p) = \alpha + B(x, p) \cdot \beta$	[5]
Receive		
Broadcast	$f(x, p) = \log_2(p) \cdot \alpha + B(x, p) \cdot \beta$	[6]
Scatter	$f(x, p) = \log_2(p) \cdot \alpha + B(x, p) \cdot \frac{p - 1}{p} \cdot \beta$	[6]
Gather		
Allgather		
Reduce	$f(x, p) = \log_2(p) \cdot \alpha + \left(\beta + \frac{p - 1}{p} \cdot \gamma \right) \cdot B(x, p)$	[6]
Allreduce		

	Summary
$f(x, p)$	Prior model
$B(x, p)$	Bytes model
x	Input parameters
p	MPI ranks
α	Latency
β	Bandwidth
γ	Computation cost

Creating models from priors



Creating models from priors

■ Example

```
void F0 (int* V, int n, int p) {
    // int V[n*p];
    MPI_Bcast(V, n*p, ...);
    for(int i = 0; i < n/100; i++) {
        // calculate something
        for(int j = 0; j < p; j++) {
            // calculate something
    }}}
```

Communication
Computation

Communication

Time: $13.0 \cdot 10^3 + 0.506n + 0.185np$
 Bytes transferred: $0.1 \cdot 10^{-3} + 4.0np$
 Theoretical model: $\log_2(p)\alpha + B(p,n)\beta$
 Prior: $c_0 + c_1 \log_2(p) + c_2 np$
 Noise-resilient: $7366 + 0 \log_2(p) + 0.189np$

Determine prior
Create model



Computation

Time: $-7782 + 0.755n + 0.002 np^{\frac{3}{4}} \log_2^2(p)$
 Basic blocks: $275 + 1.6n + 0.48np$
 Prior: $c_0 + c_1 n + c_2 np$
 Noise-resilient: $-7071 + 0.142n + 0.037np$

Determine prior
Create model

Benefits

Accuracy

Robustness to noise

Experimental cost

Evaluation

Test systems

Lichtenberg II Cluster

Jureca-DC Cluster

Comparison baseline

SWC-based

Classic

DNN-based [6]

Accuracy metrics

Exponent deviation (ED)


$$ED(f_1(x_i), f_2(x_i)) = |n_1 - n_2|$$

Relative error (RE)

$$RE(f_1(x_i)) = \frac{|y_i - f_1(x_i)|}{y} \cdot 100\%$$

Evaluation with synthetic benchmarks

- Benchmark Generator for parallel codes
- Allows flexibility on the performance behavior
- Functions with known theoretical analytical complexity



```
for (int i = 0; i < xy; i++) {
    sum = v1[i] + v2[i]}
```

$O(x^y)$

```
for (int i = 0; i < x; i++) {
    for (int j = 0; j < x; j++) {
        mul = v1[i] * v2[j]}}
```


$O(x^2)$

```
for (int i = 0; i < xy; i++) {
    a = std :: max(a, v1[i]);
    v2[i] = a;}
```

$O(x^y)$

Evaluation with synthetic benchmarks

- Benchmark Generator



Evaluation with synthetic benchmarks





- 200 synthetic functions
- Accuracy
- Exponent deviation: comparing performance models with their theoretical expectation

Models	Computation		Communication	
	MPI ranks	Message size		
SWC-based	0	0	0	
Classic	0.44	1.14	0.57	

- Relative error
 - SWC-based: 35% (computation) and 60% (communication)
 - Classic: 45% (computation) and 91% (communication)

Evaluation with synthetic benchmarks

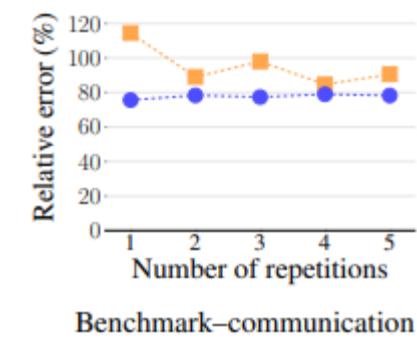
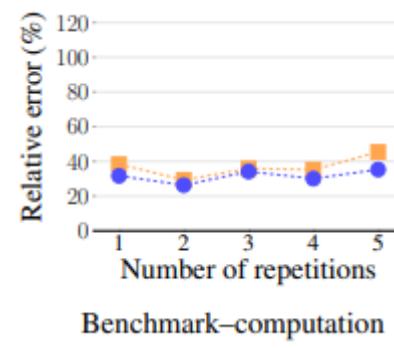
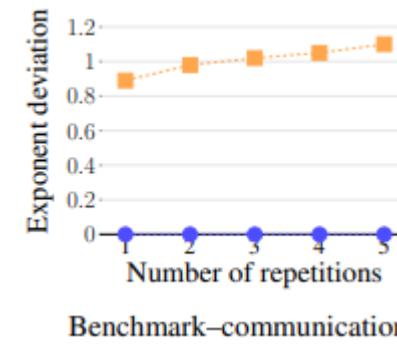
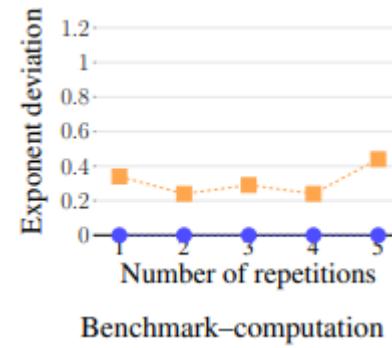
- Robustness to noise



Classic SWC-based

Evaluation with synthetic benchmarks

- Experimental costs



■ Classic ● SWC-based

Application case studies

- Kripke [7]
- Relearn [8]

Known theoretical performance

App/System	Training configurations	Test points
Kripke	$p \in \{512, 1000, 1728, 2744, 4096\}$	$(p, G, Z) = (5832, 160, 20^3)$
Lichtenberg II	$G \in \{32, 64, 96, 128, 160\}$ $Z \in \{4^3, 8^3, 12^3, 16^3, 20^3\}$	$(p, G, Z) = (4096, 192, 20^3)$ $(p, G, Z) = (4096, 160, 24^3)$
RELeARN	$p \in \{32, 64, 128, 256, 512\}$	$(p, n) = (1024, 450)$
Jureca-DC	$n \in \{250, 300, 350, 400, 450\}$	$(p, n) = (512, 500)$


Evaluation

- Application case studies
- Accuracy

Model	Asymptotic complexity	ED		
		Δp	ΔG	ΔZ
Kripke-computation				
Theoretical	$\mathcal{O}(G \cdot Z)$			
Classic	$\mathcal{O}(p \cdot \log_2^2(p) \cdot G^{\frac{3}{4}} \cdot \log_2(G) \cdot Z^{\frac{4}{5}})$	1	0.25	0.20
DNN-based	$\mathcal{O}(p \cdot G^{\frac{5}{4}} \cdot Z^{\frac{2}{3}})$	1	0.25	0.33
SWC-based	$\mathcal{O}(G \cdot \log_2(G) \cdot Z^{\frac{3}{4}})$	0	0	0.25
Kripke-communication				
Theoretical	$\mathcal{O}(p^{\frac{1}{3}} + G \cdot Z^{\frac{2}{3}})$			
Classic	$\mathcal{O}(p^{\frac{4}{3}} \cdot \log_2(p) \cdot G^{\frac{3}{4}} \cdot \log_2(G) \cdot Z^{\frac{1}{3}} \cdot \log_2^2(Z))$	1	0.25	0.33
DNN-based	$\mathcal{O}(G^{\frac{5}{4}} \cdot Z^{\frac{1}{2}})$	0.33	0.25	0.16
SWC-based	$\mathcal{O}(G \cdot Z^{\frac{2}{3}})$	0.33	0	0
RELeARN				
Theoretical	$\mathcal{O}(p + n \cdot \log_2(n \cdot p))$			
Classic	$\mathcal{O}(p^{\frac{2}{3}} \cdot n^{\frac{3}{4}} \cdot \log_2(n))$	0.33	0.25	
DNN-based	$\mathcal{O}(p^{\frac{2}{3}} \cdot \log_2(p) \cdot n^{\frac{1}{4}})$	0.33	0.75	
SWC-based	$\mathcal{O}(p + n^{\frac{5}{4}} \cdot \log_2(n) \cdot p^{\frac{1}{4}})$	0	0.25	

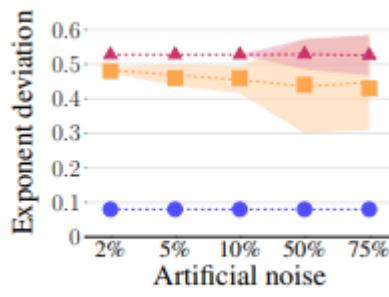
Evaluation

- Application case studies
- Accuracy
- Better in 6/8 cases




Application case studies


- Robustness to noise




Kripke-computation




Kripke-communication




RELeARN



Kripke-computation



Kripke-communication



RELeARN







Classic

DNN-based

SWC-based

Application case studies

- Experimental costs



■ Classic ▲ DNN-based ● SWC-based

Conclusion

- Our method accurately captures the computational effort of an application in close alignment with its theoretical performance model
- We reduce, if not eliminate, the need for multiple time measurements
- Under artificial noise, our models maintained stable error rates

References

- [1] A. Calotoiu, D. Beckinsale, C. W. Earl, T. Hoefler, I. Karlin, M. Schulz, and F. Wolf, “Fast multi-parameter performance modeling,” in 2016 IEEE International Conference on Cluster Computing (CLUSTER). IEEE, 2016, pp. 172–181.
- [2] LLVM admin team. (2023) LLVM website. Accessed 2023/08/21.[Online]. Available: <https://llvm.org/>
- [3] Score-P developer community. (2023) Scalable performance measurement infrastructure for parallel codes (Score-P). Accessed 2023/08/21. [Online]. Available: <https://www.vi-hps.org/projects/score-p>
- [4] W. Zhang, M. Hao, and M. Snir, “Predicting hpc parallel program performance based on llvm compiler,” Cluster Computing, vol. 20, pp.1179–1192, 2017.
- [5] E. Chan, M. Heimlich, A. Purkayastha, and R. Van De Geijn, “Collective communication: theory, practice, and experience,” Concurrency and Computation: Practice and Experience, vol. 19, no. 13, pp. 1749–1783, 2007.

References

- [6] M. Ritter, A. Geiß, J. Wehrstein, A. Calotoiu, T. Reimann, T. Hoefer, and F. Wolf, “Noise-resilient empirical performance modeling with deep neural networks,” in 2021 IEEE International Parallel and Distributed Processing Symposium (IPDPS). IEEE, 2021, pp. 23–34
- [7] A. J. Kunen, T. S. Bailey, and P. N. Brown, “Kripke-a massively parallel transport mini-app,” Lawrence Livermore National Lab.(LLNL), Livermore, CA (United States), Tech. Rep., 2015
- [8] S. Rinke, M. Butz-Ostendorf, M.-A. Hermanns, M. Naveau, and F. Wolf, “A scalable algorithm for simulating the structural plasticity of the brain,” Journal of Parallel and Distributed Computing, vol. 120, pp. 251–266, 2018

Thank you!

- You can contact us via email: extra-p-support@lists.parallel.informatik.tu-darmstadt.de
- Or on GitHub using the issues tool: <https://github.com/extra-p/extrap>

