

# Binary Analysis and Instrumentation for AMDGPU in Dyninst

Hsuan-Heng Wu and Ronak Chauhan

Computer Sciences Department  
University of Wisconsin-Madison

**Scalable Tools Workshop**

July 2025



*Dyn*<sub>inst</sub> *Para*<sub>dyn</sub>

# High-Level Goals

## Boring Goal

Normal Dyninst binary analysis and instrumentation should work for AMDGPU

## Interesting Goal

Adapt analysis and instrumentation for SIMD/SIMT architectures:

Abstract SIMD/SIMT control flow operations to a conventional CFG

Represent SIMD/SIMT data flow accurately and compactly

Efficient Fine Grained Instrumentation

Fine grained:

Record information at wave or thread level granularity

Efficient:

Synchronization-free instrumentation

Use control and data flow analysis to guide instrumentation

# Outline

## Background

- SIMT, wavefronts and execution mask

- Predicated execution and predicated control flow

## Binary Analysis of AMDGPU Kernels

- Control flow analysis of GPU binaries

- Dataflow analysis of GPU binaries

## Design Efficient Instrumentation for GPU Applications

- Leverage GPU architecture for efficient GPU instrumentation

# Background

# Background – SIMT, Wavefronts and Execution Mask

Threads grouped in **wavefronts (warps)** to execute in Single Instruction Multiple Data (SIMD)

AMDGPU Kernels uses 64 threads in a wavefront

Each thread in a wavefront (lane)


Shares **the same** program counter

By default operates on their corresponding lane of the vector data






Can be disabled by clearing a bit in a execution mask register

Single Instruction Multiple Thread - Multiple wavefronts executing the same kernel

## SIMD Processor



# Multiple SIMD processors result in SIMT



**Scalar (Wavefront)  
Processor**

**Scalar  
Registers**

# Background – Vector Registers & Vector Instructions

Execution Mask

1	...	1	1	1
---	-----	---	---	---

v\_add\_u32 v0, v1, v2

v1

v1[63]	...	v1[2]	v1[1]	v1[0]
--------	-----	-------	-------	-------

v2

v2[63]	...	v2[2]	v2[1]	v2[0]
--------	-----	-------	-------	-------



v0

v1[63]+v2[63]	...	v1[2]+v2[2]	v1[1]+v2[1]	v1[0]+v2[0]
---------------	-----	-------------	-------------	-------------

# Cross Lane Operations – Lane Shifting

Execution Mask

1	...	1	1	1
---	-----	---	---	---

v\_add\_u32 v0, v1, v2 wave\_shr:1

v1

v1[63]	...	v1[2]	v1[1]	v1[0]
--------	-----	-------	-------	-------

v2

v2[63]	...	v2[2]	v2[1]	v2[0]
--------	-----	-------	-------	-------



v0

unmodified	v1[63]+v2[62]	...	v1[2]+v2[1]	v1[1]+v2[0]
------------	---------------	-----	-------------	-------------


# Background – Predicated Execution

Execution Mask




v\_add\_u32 v0, v1, v2

v1



v2



v0[63]




...

v1[2]+v2[2]




v1[1]+v2[1]




v0[0]

```
IF (threadIdx.x % 2)
{
    A; // odd threads
} ELSE {
    B; // even threads
}
```

## Wavefront View CFG



## Per Thread View CFG



# **Binary Analysis for AMDGPU Kernels**

# Control Flow Analysis for AMDGPU Kernels

Analyzing control flow at the **wavefront level** is similar to that of CPU binary

Threads in a wavefront share the same PC

Control flow instructions are scalar

Analyzing per thread control flow is the interesting topic here

**Control flow analysis enables data flow analysis**


Under predicated control flow, threads can have different control flow based on EXEC

**IF-THEN-ELSE, DO-WHILE, SWITCH-CASE, SIMT Jump Table**


Map these code constructs to conventional code constructs for programmer's to understand the semantics of the binaries

```
IF (threadIdx.x % 2)
{
    A; // odd threads
} ELSE {
    B; // even threads
}
```

## Wavefront View



## Per Thread View



# Control Flow Analysis for AMDGPU Kernels

Capture and represent per-thread control flow compactly

**Threads with the same execution mask shares the same CFG**

The exact value of an execution mask *might be* only known at runtime

Rely on symbolic analysis to determine the set of equivalence classes of the symbolic execution mask value (when possible)

**For each predicated control flow construct**

Generate a per-thread control flow graph per equivalence classes of the symbolic execution mask

In the cases where all threads have the same mask, this looks like a traditional CFG

# Data Flow Analysis for AMDGPU Kernels

Data flow analysis needs to capture SIMD/SIMT data dependencies accurately

- Threads can be disabled by the execution mask

- For each instruction


- Only active threads contribute to the data flow

- Vector registers can be partially updated by active threads

- The DEF and USE set of instructions need to expand to individual lanes of a vector register

```
int index = threadIdx.x;
if (index%2) {
    v[index] = index + 10;
} else {
    v[index] = index * 20;
}
```

```
vcc = if(index%2)
s_and_saveexec_b64 s[0:1], vcc
v_add_u32_e32 v1, 10, v0
s_xor_b64 exec, exec, s[0:1]
v_mul_u32_u24_e32 v1, 20, v0
...
```



# Data Flow Analysis for AMDGPU Kernels

Capture In-Lane dataflow compactly

In-lane operations have the same data flow for all threads

Capture Cross-Lane & Predicated dataflow accurately

Vector instructions under predicated control flow

Permutation instructions can introduce dependencies between any two lanes of two vector registers

Need to track the dependencies between all lanes

# **Design and Implementation of Efficient AMDGPU Binary Instrumentation**

# Instrumentation - Accessing Instrumentation Variables

Insert instructions to observe the behavior of the kernel

Insert instructions to bookkeep execution in instrumentation variables  
Counters / Timestamps / Tracing

Data eventually needs to be available on the host


Simplest: at the end of kernel execution. Effective for counters and timers.

More complex: periodically at initiative of the device. Allows counters and times to be viewed over time on the host. Allows trace buffers to be flushed when filled.

Perhaps more complex: periodically at initiative of the host. Allows interactive tools to poll current state of the instrumentation results.

# Instrumentation variables: per kernel launch


Two regular scalar variables A and B



Count cumulative information  
across waves

Requires synchronization or atomic  
operations as all waves write to the  
same locations

# Instrumentation variables: per waveform



Each **wave** has own instance of the scalar variable

No synchronization need

Measure per-wave information on device and aggregate if necessary

# Instrumentation variables: per thread


Wave 1		Wave 2		Wave 3	
$A_{1,1}$	$B_{1,1}$	$A_{2,1}$	$B_{2,1}$	$A_{3,1}$	$B_{3,1}$
$A_{1,2}$	$B_{1,2}$	$A_{2,1}$	$B_{2,2}$	$A_{3,2}$	$B_{3,2}$
$A_{1,3}$ ...	$B_{1,3}$ ...	$A_{2,3}$ ...	$B_{2,3}$ ...	$A_{2,3}$ ...	$B_{3,3}$ ...
$A_{1,64}$	$B_{1,64}$	$A_{2,64}$	$B_{2,64}$	$A_{2,64}$	$B_{3,64}$

Each **thread** has its own instance of the scalar variable (so, a vector per wave)


Computations on thread-level variables use vector operations

Only active threads read/write to their instances

# Mechanics of offloading on the device (GPU)




# Host-side implementation for Instrumentation Variables



We modify kernel signature to take extra argument  
The extra argument is the address for instrumentation variables

# Device-side implementation for Instrumentation Variables

Example for accessing wave-level instrumentation variables




For each wave, calculate and store base address for its variables once in a register

Offsets for instances are same across waves

Therefore base + offset addressing will work at granularity of individual instances

# Instrumenting AMDGPU code objects



# Current Implementation Status

## Control Flow Analysis

[Wavefront-level control flow analysis](#)

## Dataflow analysis

Experimental support for instruction semantics and symbolic execution based on AMD's machine readable instruction semantic specification

## Instrumentation

Initial support for instrumentation at the granularity of a kernel launch

## Support tools

Tools that work alongside the Dyninst mutator for AMDGPU

Alpha level of readiness

# Next Steps

Develop per-thread control flow analyses

Develop per-thread data flow analyses

Testing infrastructure for the instrumentation tools to make it ready for general use

Complete support for wave-level instrumentation

Extend instrumentation to thread level granularity

# Questions ?

hwu337@wisc.edu

ronak@cs.wisc.edu