The background of the slide features a scenic landscape of a river flowing through a forested area. The sky is filled with dramatic, colorful clouds reflecting off the water. In the distance, mountains are visible under a bright, setting sun.

SIMON: A Simple Monitoring Framework for Heterogeneous Application Observability

*Michael Dushkoff, Allen D. Malony, Grace McLewee, Kevin Huck
University of Oregon*

Scalable Tools Workshop 2025
Granlibakken Resort, Lake Tahoe



O

HPC Heterogeneity


- Rapidly changing HPC landscape
- Heterogeneous architectures and system hardware
 - Processor types / diversity: CPU, GPU, NPU, ... from multiple companies
 - Memory: cache hierarchy, NVRAM, CXL, unified memory, ...
 - Network: interconnection, smarNICs, DPUs, ...
- Drives software technology for dealing with heterogeneity
 - Higher-level programming languages and frameworks
 - Support for cross-platform performance portability
- Heterogeneity in HPC applications (parallelism)
 - Traditional “faster is better” orthodoxy of (*hierarchical (vertical) parallelism*)
 - New paradigms of concurrent heterogeneous tasks (*horizontal parallelism*)
 - Broader computing continuum and variety in types of users/applications
 - Different execution models (SPMD, workflow, streaming, AI-coupled, ...)

Observability

- In general, *observability* of a system is the ability to make measurements of its operation in order to understand how and how well the system performs when used by application
 - Ability to understand the internal state of a system by analyzing its external outputs to gain insights into its health and performance.
- Concept of observability has gained considerable attention in scaled-out distributed enterprise system environments (cloud)
 - Challenges in developing/deploying applications in diverse ecosystems
 - Complexity of workflows of many heterogeneous tasks
 - Various runtime stacks in use and scaling out performance challenges

Cloud-native Observability

- Visibility into how applications behave in cloud-based systems
 - Identify execution anomalies and evaluate performance issues
- Observation telemetry data produced from measurements
 - *Event logs*: record details of an event
 - *Metrics*: data to quantify performance and health
 - *Traces*: track end-to-end services
 - Collectively use the term *monitoring*
 - Emphasize online telemetry accessibility
- Observability is more about making sense of it
 - Reasoning and improving in real time
- Commercial importance of cloud-based applications
 - Motivated development of sophisticated observability frameworks



HPC Performance Observability

- Older history of observability in HPC (see my Ph.D. thesis 😊)
- How to use HPC resources effectively to achieve optimal performance
- Observability focus on parallel performance measurement and analysis
 - Identify performance inefficiencies and limitations
 - Detailed understanding of processing, memory, and communications hardware
 - Inform application tuning decisions
 - Observability methods should minimally impact the system
- HPC research community has developed robust performance tools
 - Address HPC application observability challenges
 - Different in design and operation from cloud-based environments
- Heterogeneity in HPC raises issues about tools for observability
- Reimagine observability in the context of HPC application monitoring

My Obsession with Application-level Monitoring

- Interested HPC application-level monitoring for a long time:
 - Hypermon (1990)
 - VNG (2003)
 - TAUg (2006)
 - TAUoverSupermon (2007)
 - TAU snapshots (2007)
 - TAUoverMRNet (2008)
 - Monitoring sweetspots (2008)
 - TAUMon (2010)
 - POWMon (2015)
 - APEX (2015)
 - WOWMON (2016)
 - SOS (2016)
 - Artemis (2021)
 - Symbiomon (2021)
 - ZeroSum (2023)
 - SOMA (2024)
- Motivated by belief of value in online access to performance data


SIMON Objectives



- Support heterogeneous HPC application observability
 - Simple monitoring architecture and framework
 - Based on an asynchronous time-series observational model
 - Enable in situ analytics and visualization
- Offer (simple) functionality and development
 - Broadly available and easy to use
 - Straightforward to integrate with HPC applications and tools
 - Can operate out-of-the-box in standard HPC environments
 - Programmable, extensible, and configurable
 - Addresses modern observability requirements
 - Utilize modern observability technology
- Easily construct new monitoring solutions (however simple)


SIMON Design Architecture for Application Monitoring

- Minimalist approach
 - Separation of concerns
 - Focus on telemetry data transport
- Components
 - **Providers** (measurement, telemetry)
 - Process-level (embedded in process)
 - Application node-level (daemon process)
 - **Collectors** (local)
 - Application node-level (SIMON process)
 - **Aggregator** (global)
 - SIMON resource (SIMON process)
 - Analytics helpers
 - **Messaging** (local, global)
 - Dashboard
- SIMON launches / terminates with the application



SIMON Prototype

- SIMON started as a graduate course project
 - *Advanced Parallel Computing*, Winter 2025
 - Two 1st-year graduate students
 - 2-3 months worth of effort thus far
- Development decisions
 - Python for all SIMON processes
 - ZeroMQ for messaging
- Leverage existing measurement providers
 - HPC performance: TAU, APEX, ZeroSum, ...
 - System: /proc, Linux top program, ..
 - Other profilers: PyTorch Profiler, ...
- Work with Slurm to launch
- Python-based visualization
 - Seaborn, Matplotlib



- Application process
- Process-level measurement
- Node-level measurement
- SIMON local process
- SIMON global process
- SIMON analysis process


Why Python and ZeroMQ?

- Python
 - Ubiquitous and portable
 - Highly functional
 - Many libraries and tools
 - Used in other observability tools and frameworks
 - Used in other application environments of interest
 - Easy to use for prototyping
- ZeroMQ
 - Open source messaging library
 - Robust and universally available
 - Asynchronous message processing
 - Multiple communication patterns
 - Fast enough for clustered systems
 - Several language APIs and OS
 - Inter-process communication transport (`zmq_ipc`, `zmq_tcp`)
 - NOT MPI
- Why not?




Launch, Configuration, and Measurement Provider

- SIMON collector processes and aggregator starts first
 - Establish OMQ connections
- Application node measurement daemon starts
 - Discovers SIMON collector process
 - Establishes OMQ connections
- Application processes starts (w/ measurement provider)
 - Special magic to set up SIMON at init
 - Discovers SIMON collector process
 - Establishes OMQ connections
- Measurement data converted to structured text
 - Structure determined by type of data
 - Example: CSV, key-value, formatted trace records, ...
- Uses ZeroMQ API to send asynchronously to collector
- Each message is timestamped with provider ID



Test Case: Distributed “Top” Monitor


- Multi-node heterogeneous application
 - Complex fluid field simulation on image
 - 16 instances (processes) on 4 nodes
 - Ran for 300 seconds
- Experiment on Talapas cluster
 - Each node has 2x 32-core AMD EPYC
 - Script ran with other programs active
- Local SIMON daemon measurement
 - Ran `top -b -n 1` every second
- Deliver telemetry to SIMON aggregator
 - Allocated its own node
 - >1 MB / second total “top” telemetry
- Realtime analytics and visualization



- Certain application observation requires application and system measurement
 - Can not diagnose performance problems otherwise
 - Need for periodic sampling
- Need to associate application performance view with system behavior
 - Resource utilization monitors for applications
 - Check for misconfiguration and efficient utilization
 - Confirmation of expected hardware / operating system behavior
 - Aid to identify failure scenarios
- Monitor coverage (per process, all nodes, sampling)
 - Application threads (LWP)
 - CPU hardware (HWT)
 - Memory
 - GPU hardware for all processes, all nodes in the allocation
- Utilizes an extra thread for monitoring

Test Case: XGC Monitor

- SIMON integration in heterogeneous HPC application
 - XGC gyrokinetic PIC code
 - C++ / Fortran 90, MPI, GPU, Kokkos, Cabana
- Measurement providers for SIMON telemetry data
 - ZeroSum modified to write out sample windows
 - Each process samples every second
 - Writes to SIMON every 10 seconds
- Experiment on Frontier
 - 64 MPI ranks distributed on 8 nodes
 - Each with 64-core AMD EPYC and MI250X GPU
 - Each rank assigned graphic compute die (GCD)
- Deliver telemetry to SIMON aggregator
 - Node collectors gather telemetry for 8 ranks
 - Deliver to aggregator running on its own node
- Realtime analytics and visualization




PyTorch Profiler

https://pytorch.org/tutorials/recipes/recipes/profiler_recipe.html

- Built into PyTorch
 - Analyze deep learning models during training and inference
 - Identify performance bottlenecks, optimize operations, improve memory efficiency, and accelerate training and inference of PyTorch models
- Detailed insights into various aspects of model execution
 - Time spent on operations (different layers, CPU and GPU)
 - Memory usage (allocation/deallocation for tensors, CPU and GPU)
 - CUDA kernel activity
 - Tensor shapes to operations
 - Python tack traces
- Context manager to define code region to be profiled

Test Case: Distributed PyTorch Training (ResNet-50)

- Distributed PyTorch machine learning application
 - Training using ResNet-50 for image classification
- Measurement providers for telemetry data
 - Embedded in the application code (per process)
 - Daemon capturing node-level data
- PyTorch Profiler to capture model-level profiling information
 - Worked with profiling context manager
 - Capture profiling data for specific code regions
 - Wrapped epoch training code to see each layer and PyTorch routines used
 - Telemetry output for each epoch
- Daemon to observe GPU utilization
 - NVIDIA SMI monitor
 - GPU utilization, power consumption, and GPU memory usage
 - Telemetry output per second
- Experiment on 8 Talapas GPU nodes
 - Each with 24-core AMD EPYC and NVIDIA A100 GPU
- Deliver telemetry to SIMON aggregator
 - Allocated its own node
- Realtime analytics and visualization



Preliminary Takeaways and Possible Issues

- Took some trial-and-error to get things configured properly
- Relatively straightforward to hook up providers
- Mostly small-scale experiments
 - Real-time telemetry transport, analysis, and visualization
 - No noticeable performance effects
- Did not run experiments at moderate or large scale
- Obvious potential problem with overloading single aggregator

Related Work (and many more)

- Systems monitoring
 - Ganglia (<https://github.com/ganglia>)
 - Nagios (<https://www.nagios.org>)
- HPC systems monitoring
 - LDMS (<https://github.com/ovis-hpc/ldms>)
 - ClusterCockpit (<https://clustercockpit.org>)
 - Pika (<https://gitlab.hrz.tu-chemnitz.de/pika>)
 - Omnistat (<https://github.com/AMDRResearch/omnistat>)
- Enterprise-ready observability technologies
 - Grafana (<https://grafana.com>)
 - Prometheus (<https://prometheus.io>)
 - VictoriaMetrics (<https://victoriametrics.com>)
 - Dynatrace (<https://www.dynatrace.com>)
 - OpenTelemetry (<https://opentelemetry.io>)
 - InfluxDB (<https://www.influxdata.com>)




SIMON Next Steps

- Work on SIMON prototype was accepted for Heteropar 2025
- Interested in applying approach more use cases and scenarios
- Involve others in building out SIMON framework and tools
- Measurement providers
 - Develop library of telemetry data sources
 - Add other HPC performance data: Caliper, Score-P, HPCToolkit, ...
- In situ telemetry processing
 - Variety of actions at collectors: filtering, merging, sampling control, ...
 - Routing of telemetry types to specialized aggregators
 - Develop library of telemetry processing modules
- Aggregation
 - Address scaling concerns through replication, specialization, ...
 - Develop analytics and visualization modules

HPC Workflow Performance Observability

- Heterogeneous tasks dynamic heterogeneous platforms
 - Adaptive and dynamic response needed to break free of static execution
 - Heterogeneity and dynamism over many spatio-temporal scales
- Middleware architecture for decision-making and information-access
 - For task execution that provides performance-assurance
 - Optimal and robust decision-making depends on Quality of Information (QoI)
 - QoI only as good as performance monitoring and introspection
- Dynamism, heterogeneity, and state uncertainty in workflows and resources make performance monitoring and introspection challenging
 - Monitoring and introspection → modeling, prediction, and actuation

Canonical Ensemble Execution Environment



- Initial attempt with RADICAL Pilot and SOMA
- SOMA stack more complicated than SIMON
- Re-attempt with RP and SIMON

D. Yokelson, et al., “SOMA: Observability, Monitoring, and in situ Analytics for Exascale Applications,” CCPE, Vol. 36, Issue: 19, August 2024.

D. Yokelson, et al., “Enabling Performance Observability for Heterogeneous HPC Workflows with SOMA,” ICPP, August 2024.

Summary

- HPC applications are growing more diverse and heterogeneous
 - More asynchronous and dynamic
 - Observation needs will be different
- Post-mortem analysis is incomplete and limited
 - Ineffective in dynamic environments
 - Timely data needed to make dynamic decisions
- Interesting observability technology in cloud world
- Concern for HPC observability impact should not deter how a monitoring system is developed and used
 - Plenty of under-utilized resources
- Monitoring technology that is portable, robust, and easy-to-integrate will be more productively applied in many cases